

# **EECS 489**

# **Computer Networks**

**Fall 2020**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*


# Agenda

---

- Introductions
- Class policies, logistics, and roadmap
- Overview of the basics
  - How is the network shared?
  - How do we evaluate a network?
  - What is a network made of?

# GSI's

---



**Jie You**



**Joseph Buiteweg**

- Office hours: See course webpage

# Mosharaf Chowdhury

---

- At Michigan since 2016
- Research focus on application-infrastructure symbiosis in large-scale networked systems
- Office hours: Wednesday 2PM – 3:15PM in 4820 BBB
  - Queue: <https://officehours.it.umich.edu/queue/421>
  - Also, by appointment (pre-scheduled via email)
- Lectures will be recorded (but not discussions)
- Ask questions in chat

# 489 in EECS curriculum

---

- **EECS 281**

- High-level logic  $\Rightarrow$  Programs
- Coding skills learned in 281 are critical for 489 assignments

- **EECS 482**

- How do machines work?
- Execute programs, interact with users, etc.
- Prior 482 experience is not needed

# What is missing?


---

- How do we access *most* services?
  - Examples include search engines, social networks, video streaming, etc.
- How do two machines communicate?
  - When they are directly connected
  - When they are not directly connected
- Using a network

# What is a network?

---

- A system of “links” that interconnect “nodes” in order to move “information” between nodes



- We will focus primarily on the Internet

# What is EECS 489 about?

---

- To learn about (at a high level)
  - How the Internet works
  - Why it works the way it does
  - How to reason about complicated design problems
- What it's not about
  - How to write web services
  - How to design web pages
  - ...



# Class workload

---

- Four assignments
  - First one is an individual assignment
  - The rest are in groups of 3
- Exams:
  - Midterm: October 19
  - Final: December 16 8 AM – 10 AM

# Grading

---

	Allocation
Assignment 1	5%
Assignment 2	15%
Assignment 3	15%
Assignment 4	15%
Midterm	25%
Final	25%

# The ALL-NEW\* assignments

---

- **Assignment 1:** measure end-to-end throughput and delay of networks (i.e., simple speed test)
- **Assignment 2:** video streaming from CDNs (i.e., simple Netflix)
- **Assignment 3:** reliable transport (i.e., how to transfer data over an unreliable network)
- **Assignment 4:** router design (i.e., how do internal elements of the network work)

All on (emulated) realistic networks using *mininet*

# Bonus Quizzes

---

- 10 MCQ and solution key for each of the 20 lectures
- Made online sometime after the lecture; live for 48 hours
- Participation counts for 0.1 on top of your final grade
  - Max 2.0
- How well you do doesn't matter
- Prepared over the summer by Tejaswi



Tejaswi Worlikar

# Enrollment and wait list

---

- Wait-listed students will be admitted in the order of wait list
- If you're planning to drop, please do so soon!

# Communication protocol

---

- Course website: <http://mosharaf.com/eecs489/>
  - Assignments, lecture slides
- Confidential content on [canvas](#)
- Piazza for all communication
  - Sign up if you haven't already
  - <https://piazza.com/umich/fall2020/eecs489/>
- Assignment submission via Github
  - Start forming groups
  - Details will be sent out soon

# Policies on late submission, re-grade request, cheating ...

---

- Detailed description in the course webpage
- Don't cheat!

---

# LET'S TALK INTERNET



# The Internet consists of many end-systems

---

● ● ● car navigator

● heart pacemaker

smartphone ●

end-system



iPad



● Linux server


MAC laptop ●



Windows PC


# Connected by switches

---




# And links

---




# Managed by many parties

---



# Transfers data


---



# A federated system

---

- The Internet ties together different networks **by the IP protocol**
  - *A common interface binds them all together*




# Switched networks

---

- End-systems and networks connected by switches instead of directly connecting them
  - Why?
- Allows us to **scale**
  - For example, directly connecting  $N$  nodes to each other would require  $N^2$  links!

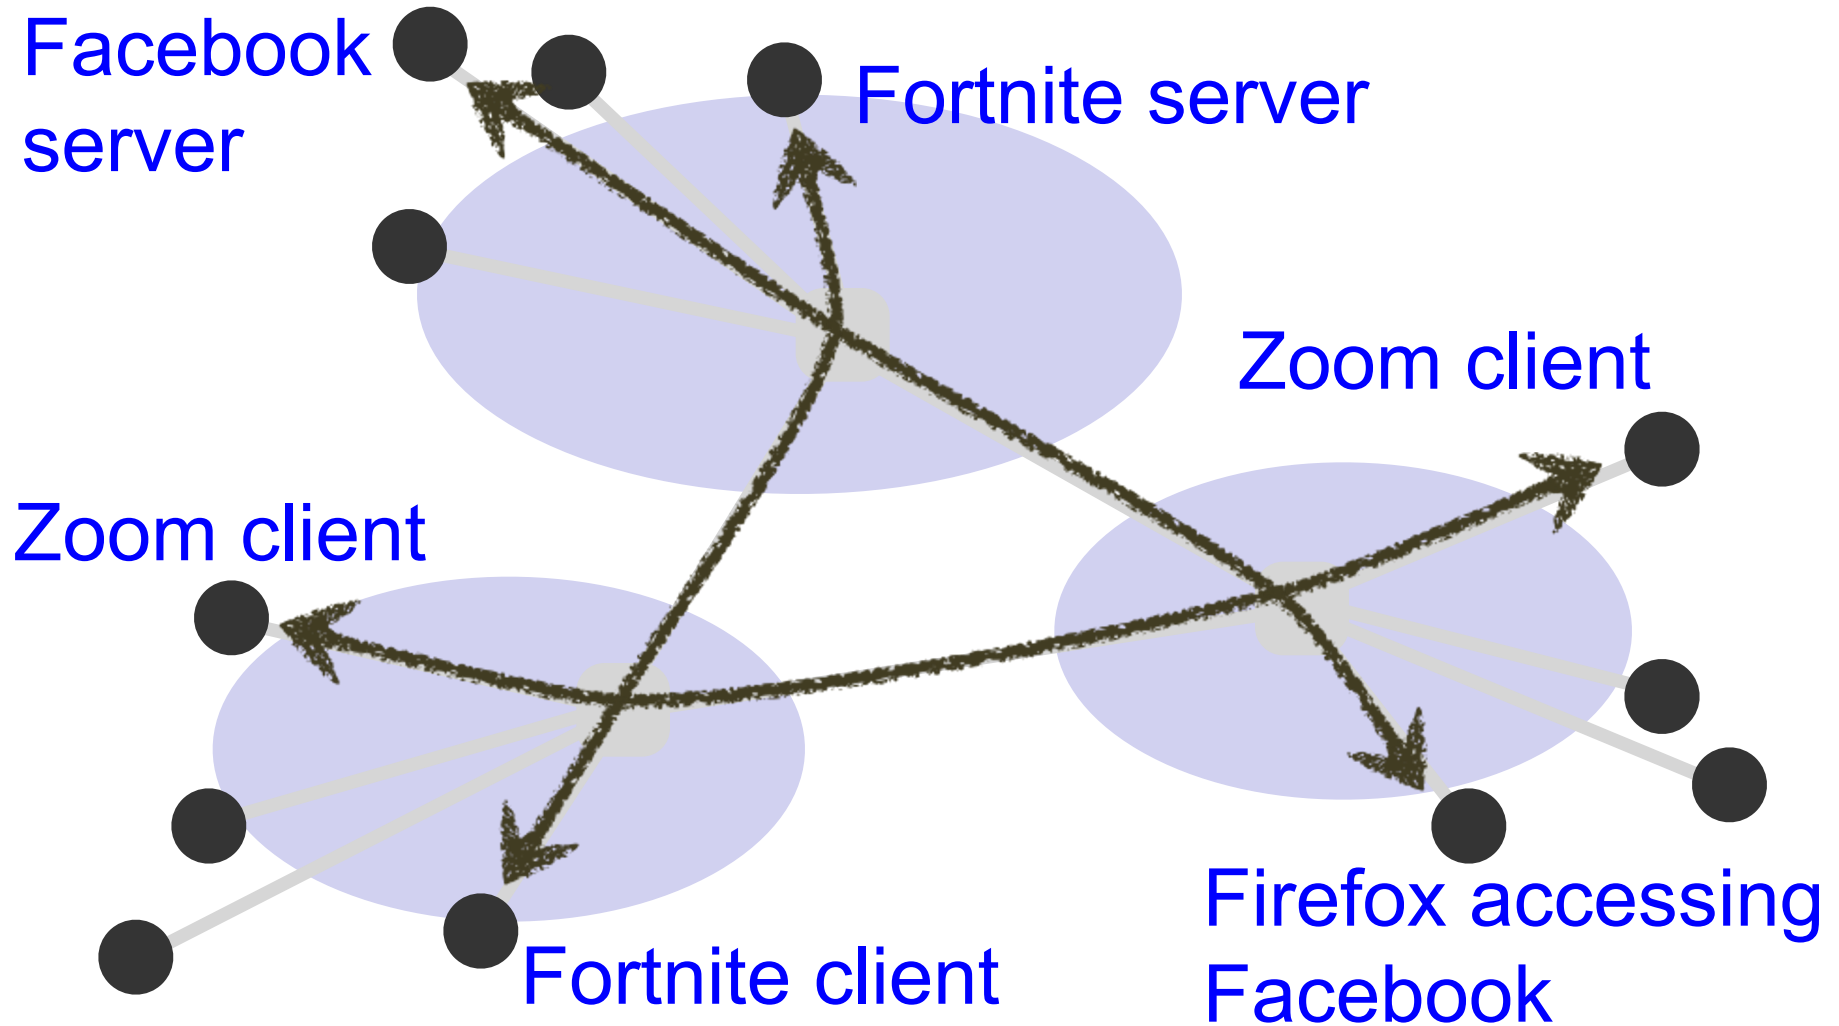
# When do we need to share the network?





# Shared among many services

---



# Two ways to share switched networks


---

- Circuit switching
  - Resource **reserved** per connection
  - Admission control: per connection
- Packet switching via statistical multiplexing
  - Packets treated independently, **on-demand**
  - Admission control: per packet

# Circuit switching


1. src sends reservation request to dst
2. Switches **create** circuit *after* admission control
3. src **sends** data
4. src sends **teardown** request

More details in backup



# Circuit switching

---



- Reservation establishes a “circuit” within a switch


# Circuit switching

---

- Pros
  - Predictable performance
  - Simple/fast switching (once circuit established)
- Cons
  - Complexity of circuit setup/teardown
  - Inefficient when traffic is bursty
  - Circuit setup adds delay
  - Switch fails → its circuit(s) fails

# Packet switching


---



- Each packet contains destination (**dst**)
- Each packet treated independently

# Packet switching

---



- Each packet contains destination (**dst**)
- Each packet treated independently
- **With buffers to absolve transient overloads**

# Packet switching

---

- Pros

- Efficient use of network resources
- Simpler to implement
- Robust: can “route around trouble”

- Cons

- Unpredictable performance
- Requires buffer management and congestion control



# Statistical multiplexing

---

- Allowing more demands than the network can handle
  - Hoping that not all demands are required at the same time
  - Results in unpredictability
  - Works well except for the extreme cases

---

**5-MINUTE BREAK!**

---

# HOW DO WE EVALUATE A NETWORK?

# Performance metrics

---

- Delay
- Loss
- Throughput

# Delay

---

- How long does it take to send a packet from its source to destination?

# Delay

---

- Consists of four components

- Transmission delay
- Propagation delay
- Queuing delay
- Processing delay

} due to link properties

} due to traffic mix and switch internals

# A network link

---



- Link bandwidth
  - ▢ Number of bits sent/received per unit time (bits/sec or bps)
- Propagation delay
  - ▢ Time for one bit to move through the link (seconds)

# 1. Transmission delay

---

- How long does it take to push all the bits of a packet into a link?
- Packet size / Transmission rate of the link
  - E.g., 1000 bits / 100 Mbits per sec =  $10^{-5}$  sec



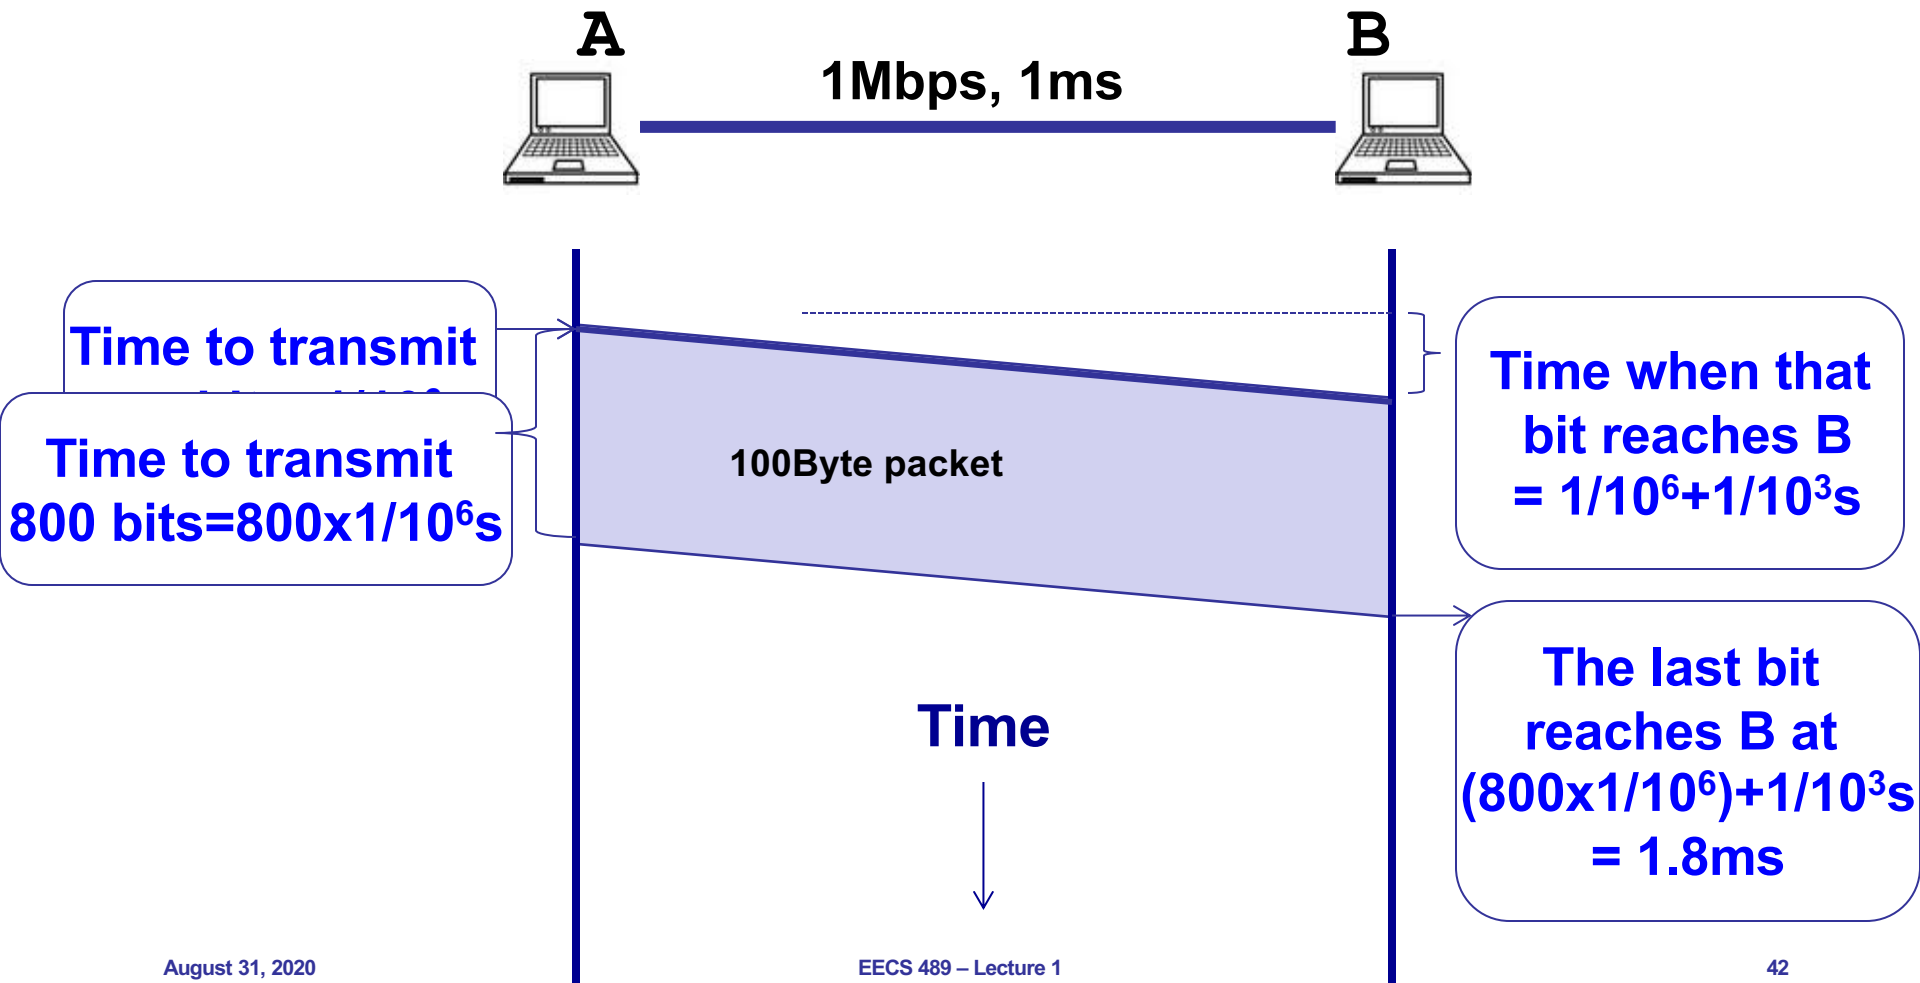
## 2. Propagation delay

---

- How long does it take to move one bit from one end of a link to the other?
- Link length / Propagation speed of link
  - E.g., 30 kilometers /  $3 \times 10^8$  meters per sec =  $10^{-4}$  sec

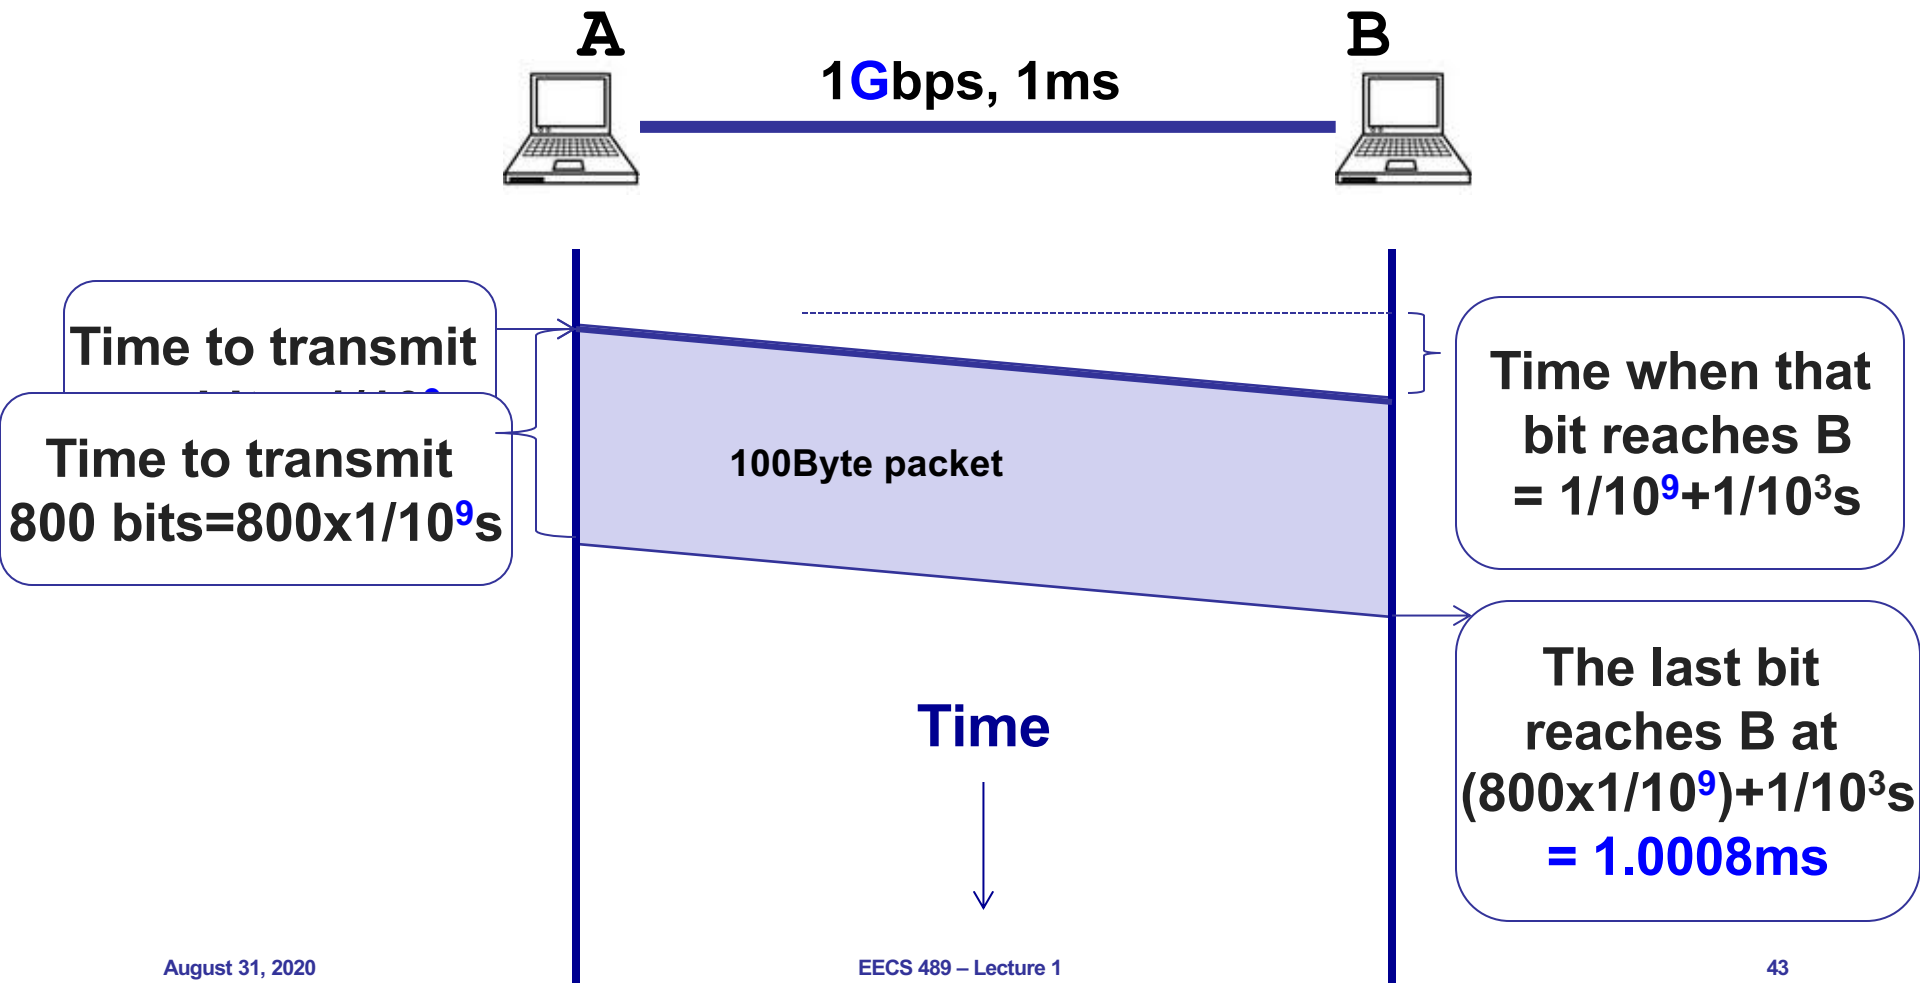
# Packet delay

## Sending a 100-byte packet




# Packet delay


## Sending a 100-byte packet



# Sending a large file using 100-byte packets



# Pipe view of a link




- Transmission delay decreases as bandwidth increases

# 3. Queuing delay

---


- How long does a packet have to sit in a buffer before it is processed?

# Queueing delay: “pipe” view




# Queueing delay: “pipe” view

---






# Queueing delay: “pipe” view




**Transient Overload**  
**Not a rare event!**


# Queueing delay: “pipe” view




# Queueing delay: “pipe” view



# Queueing delay: “pipe” view




# Queueing delay: “pipe” view



# Persistent overload leads to packet drop/loss

---



# Queueing delay

---

- How long does a packet have to sit in a buffer before it is processed?
- Depends on traffic pattern
  - Arrival rate at the queue
  - Nature of arriving traffic (bursty or not?)
  - Transmission rate of outgoing link

# Queueing delay

---

- How long does a packet have to sit in a buffer before it is processed?
- Characterized with statistical measures
  - Average queuing delay
  - Variance of queuing delay
  - Probability delay exceeds a threshold value



# Basic queueing theory terminology

---

- Arrival process: how packets arrive
  - Average rate  $A$
- $W$ : average time packets wait in the queue
  - $W$  for “waiting time”
- $L$ : average number of packets waiting in the queue
  - $L$  for “length of queue”

# Little's Law (1961)

---

- $L = A \times W$
- Compute L: count packets in queue every second
- Why do you care?
  - Easy to compute L, harder to compute W


# 4. Processing Delay

---

- How long does the switch take to process a packet?
  - Negligible

# End-to-end delay

---



# Round Trip Time (RTT)

---

- Time for a packet to go from a source to a destination and to come back
- Why do we care?
  - Measuring delay is hard from one end
- $RTT/2$  equals *average* end-to-end delay
  - Why not exact?

# Loss

---

- What fraction of the packets sent to a destination are dropped?

# Throughput

---

- At what rate is the destination receiving data from the source

# Throughput

---

Transmission rate  $R$  bits/sec



File of size  $F$  bits

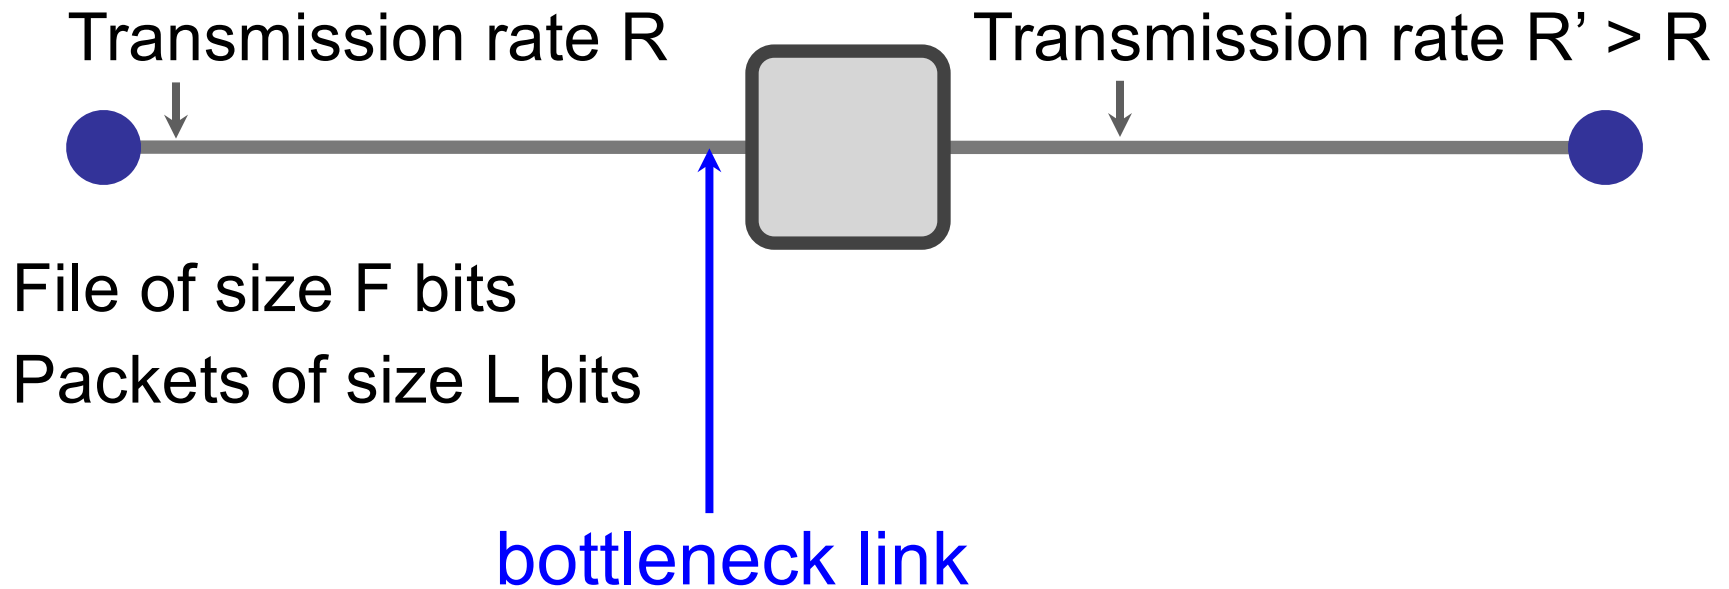
Packets of size  $L$  bits

Transfer time ( $T$ ) =  $F/R$  + propagation delay

Average throughput =  $F/T \approx R$



# End-to-end throughput



$$\text{Average throughput} = \min\{R, R'\} = R$$

# Summary

---

- How is the network shared?
  - On-demand or via reservation
- How do we evaluate a network?
  - Bandwidth, delay, loss, ...
- What is a network made of?
  - Whatever physical infrastructure exist
  - See backup slides




---

# **WHAT IS THE NETWORK MADE OF?**


# What is a network made of?

---




# What is a network made of?

---



# What is a network made of?

---



# The last hop


---





# How do we connect?

---




# Digital Subscriber Line (DSL)

---

- Twisted pair copper
- 3 separate channels
  - downstream data channel
  - upstream data channel
  - 2-way phone channel
- up to 25 Mbps downstream
- up to 2.5 Mbps upstream


# How about an cable provider as an ISP?

---



# Connecting via cable

---




# Cable

---

- Coaxial copper & fiber
- Up to 42.8 Mbps downstream
- Up to 30.7 Mbps upstream
- Shared broadcast medium

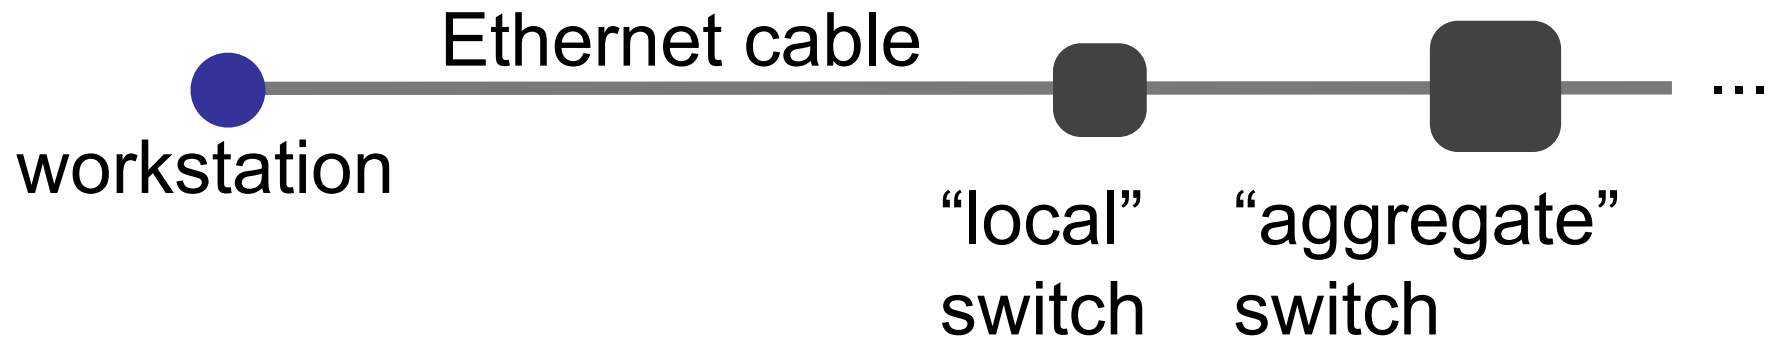
# Any other means?

---



# Ethernet

---



# Ethernet

---

- Twisted pair copper
- 100 Mbps, 1 Gbps, 10 Gbps (each direction)




# Many other ways

---

- Cellular (smart phones)
- Satellite (remote areas)
- Fiber to the Home (home)
- Optical carrier (Internet backbone)

# Where is WiFi?

---





# MASSIVE Scale

---

- 4.6 Billion users
- >1.8 Billion websites
- >200 Billion emails sent per day
- >2.5 Billion smartphones
- >2.7 Billion Facebook users
- >1 Billion hours of YouTube watched per day
- Routers that switch 10 Terabits/second
- Links that carry 100 Gigabits/second

# Have we found the right solution?

---

- We don't really know
- What we do know
  - The early Internet pioneers came up with a solution that was successful beyond all imagining
  - Several enduring architectural principles and practices emerged from their work
- Still, it is just one design with many questions

# The Internet is a lesson

---

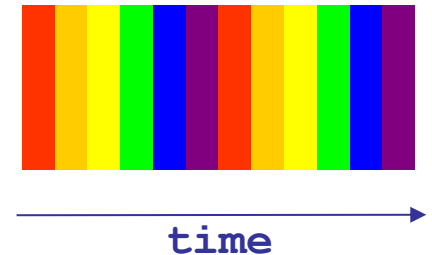
- In how to reason through the design of a very complex system
  - What are our goals and constraints?
  - What's the right prioritization of goals?
  - How do we decompose a problem?
  - Who does what? How?
  - What are the interfaces between components?
  - What are the tradeoffs between design options?

---

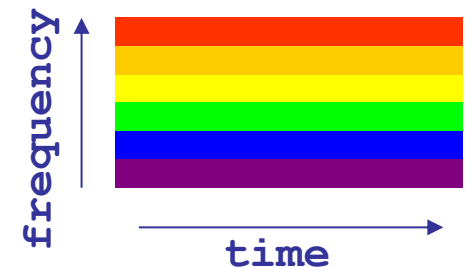
# **DETAILS ON CIRCUIT SWITCHING**

# Many kinds of circuits

- Time division multiplexing
  - divide time in time slots
  - separate time slot per circuit




- Frequency division multiplexing
  - divide frequency spectrum in frequency bands
  - separate frequency band per circuit



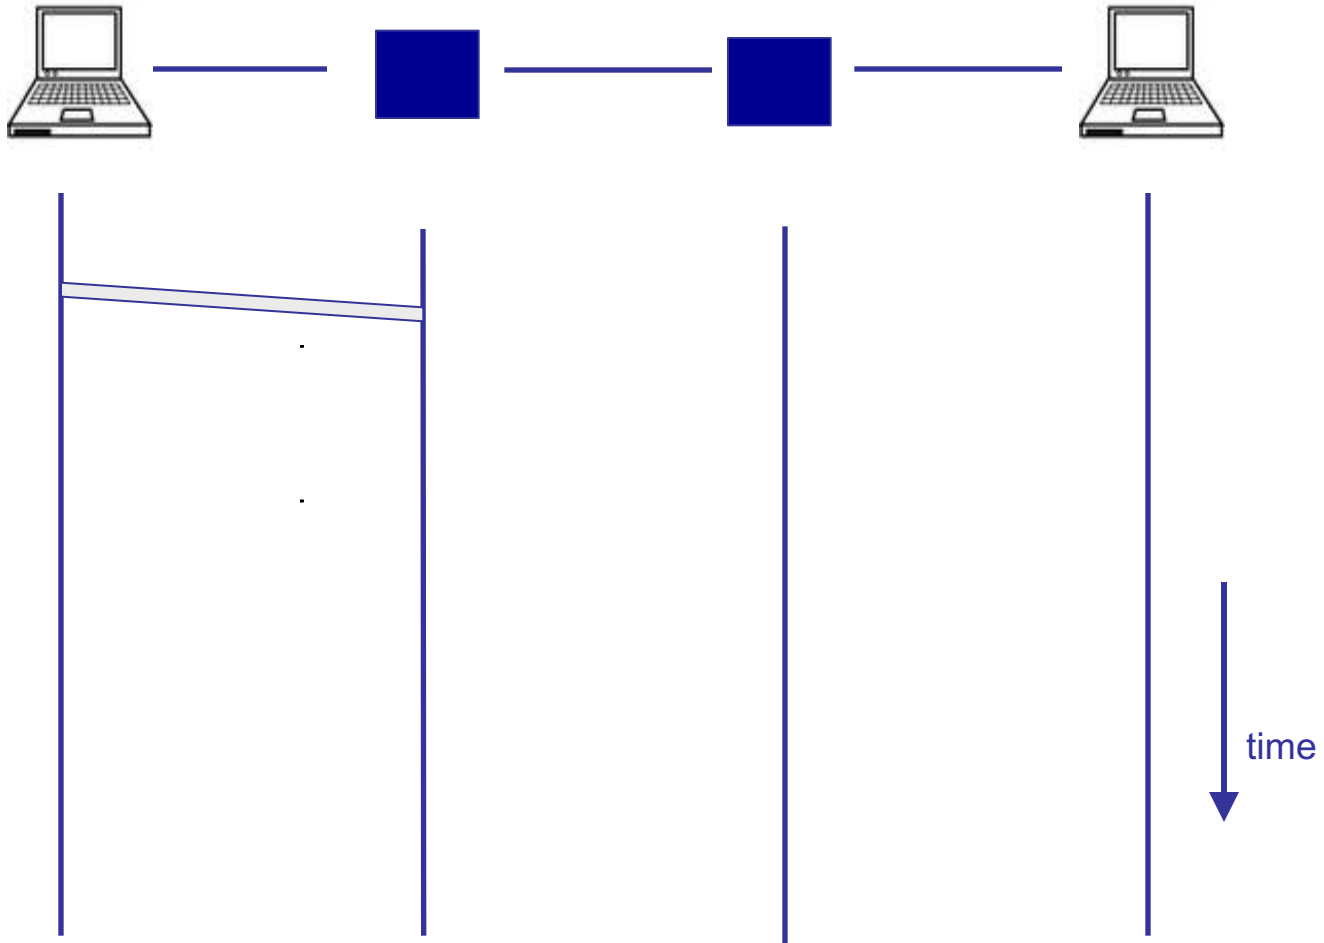


# Timing in circuit switching

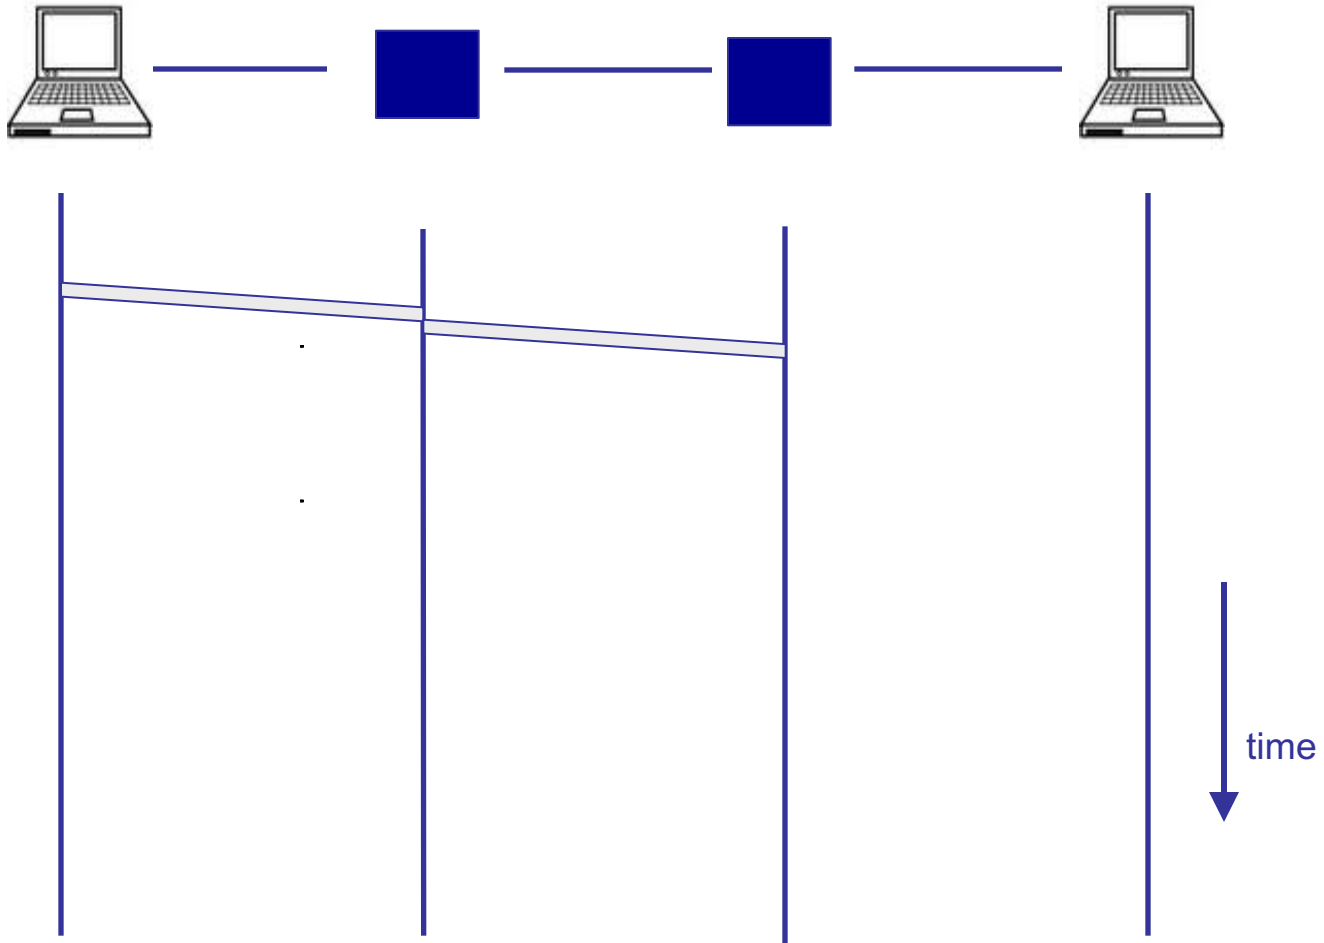
---



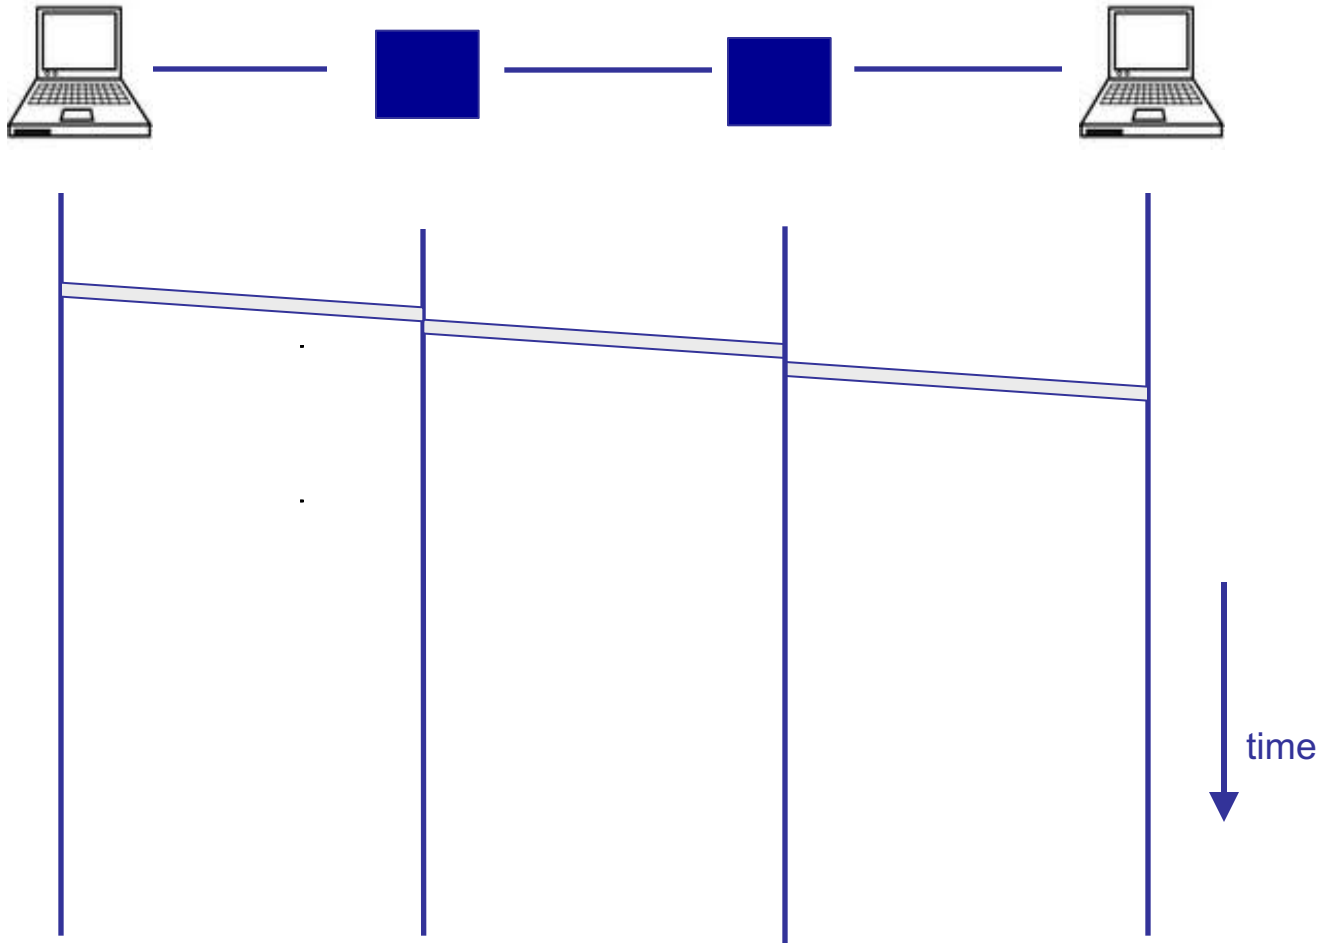
# Timing in circuit switching



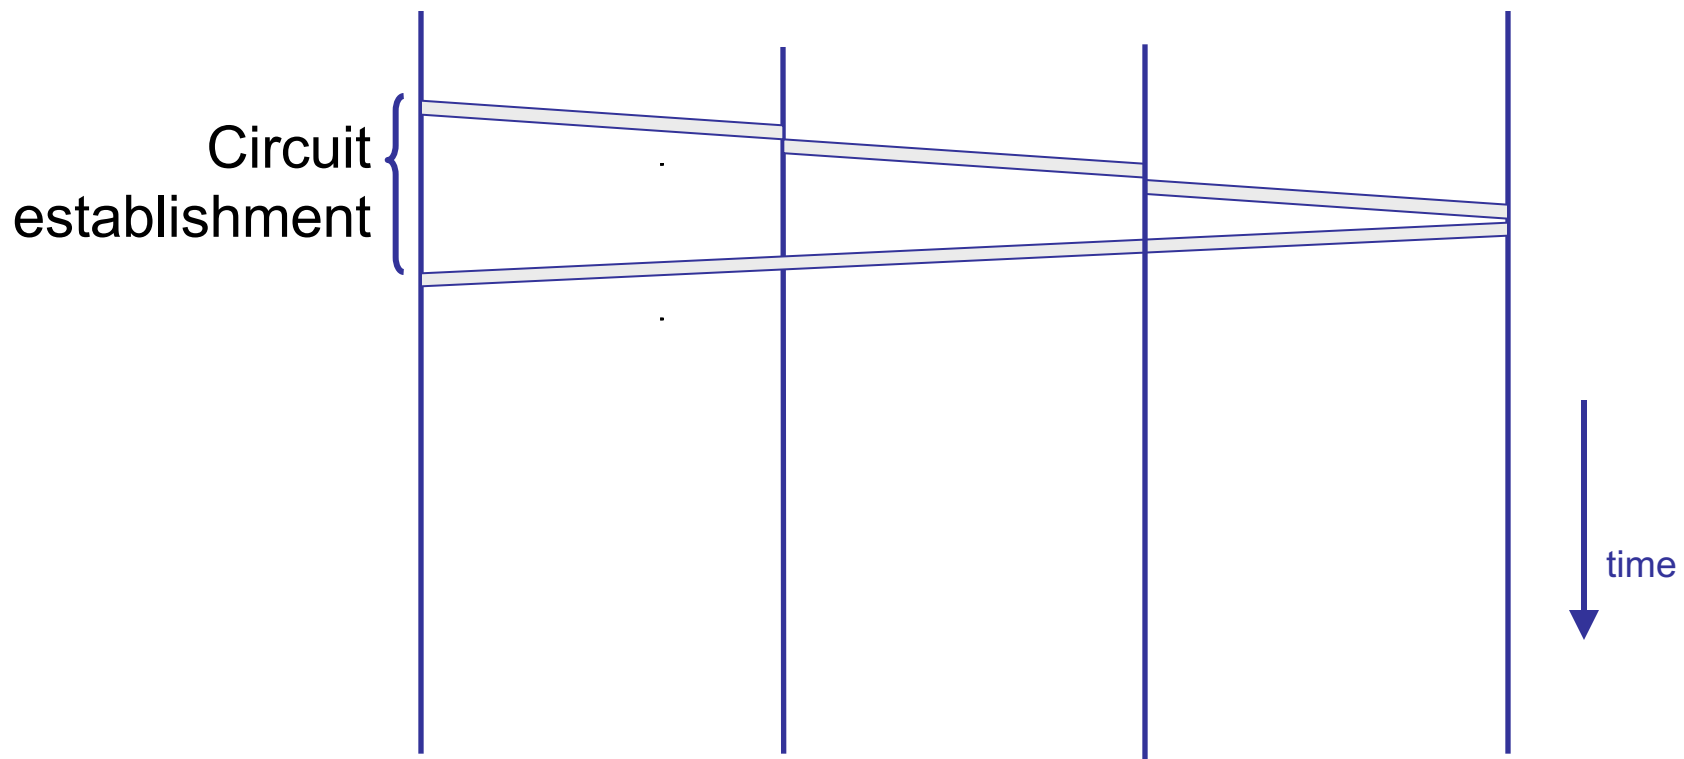
# Timing in circuit switching



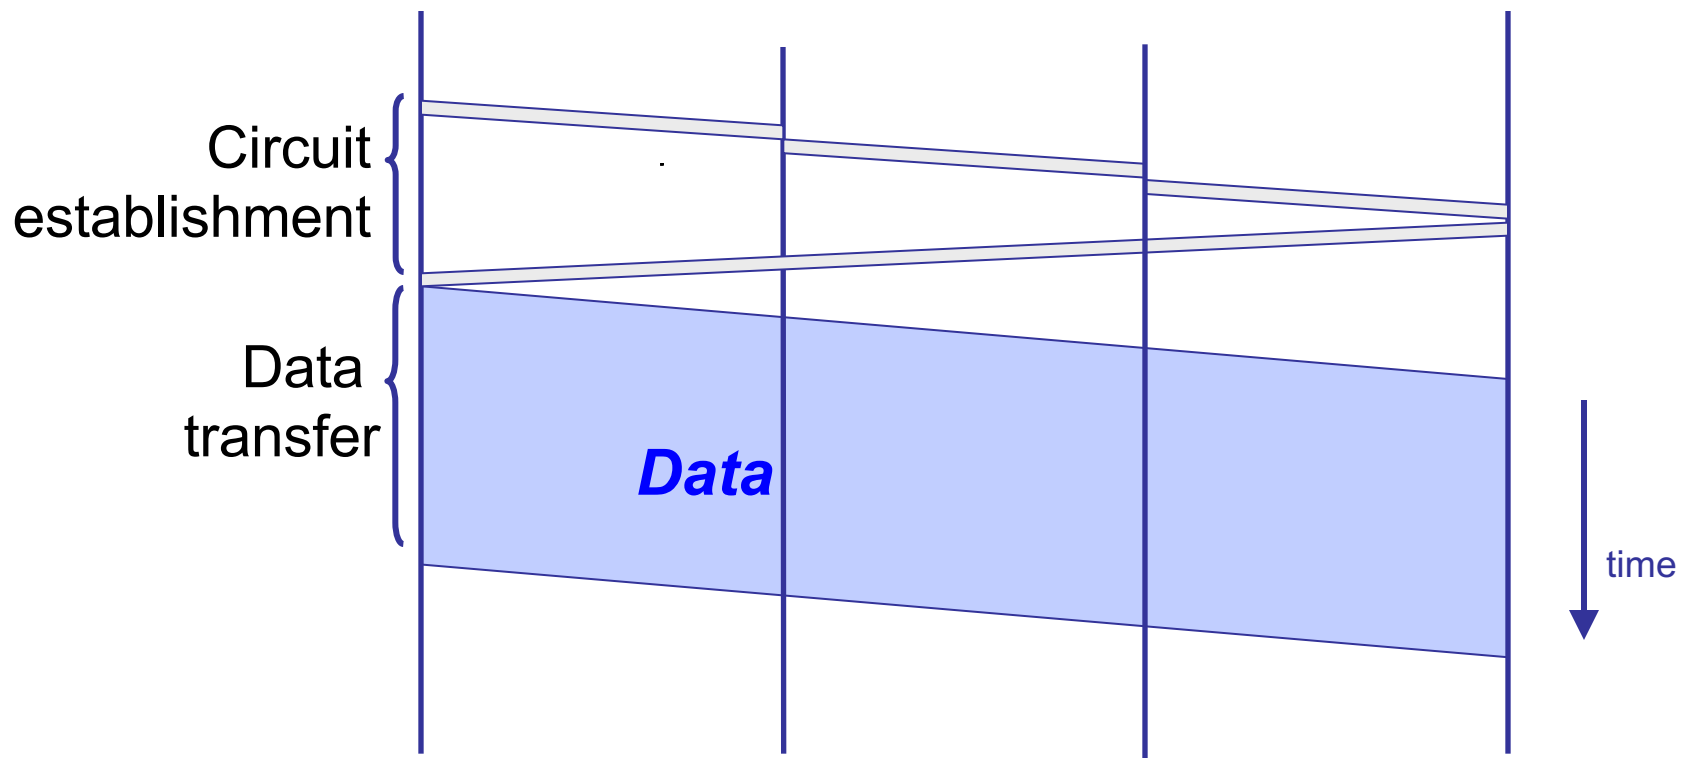
# Timing in circuit switching



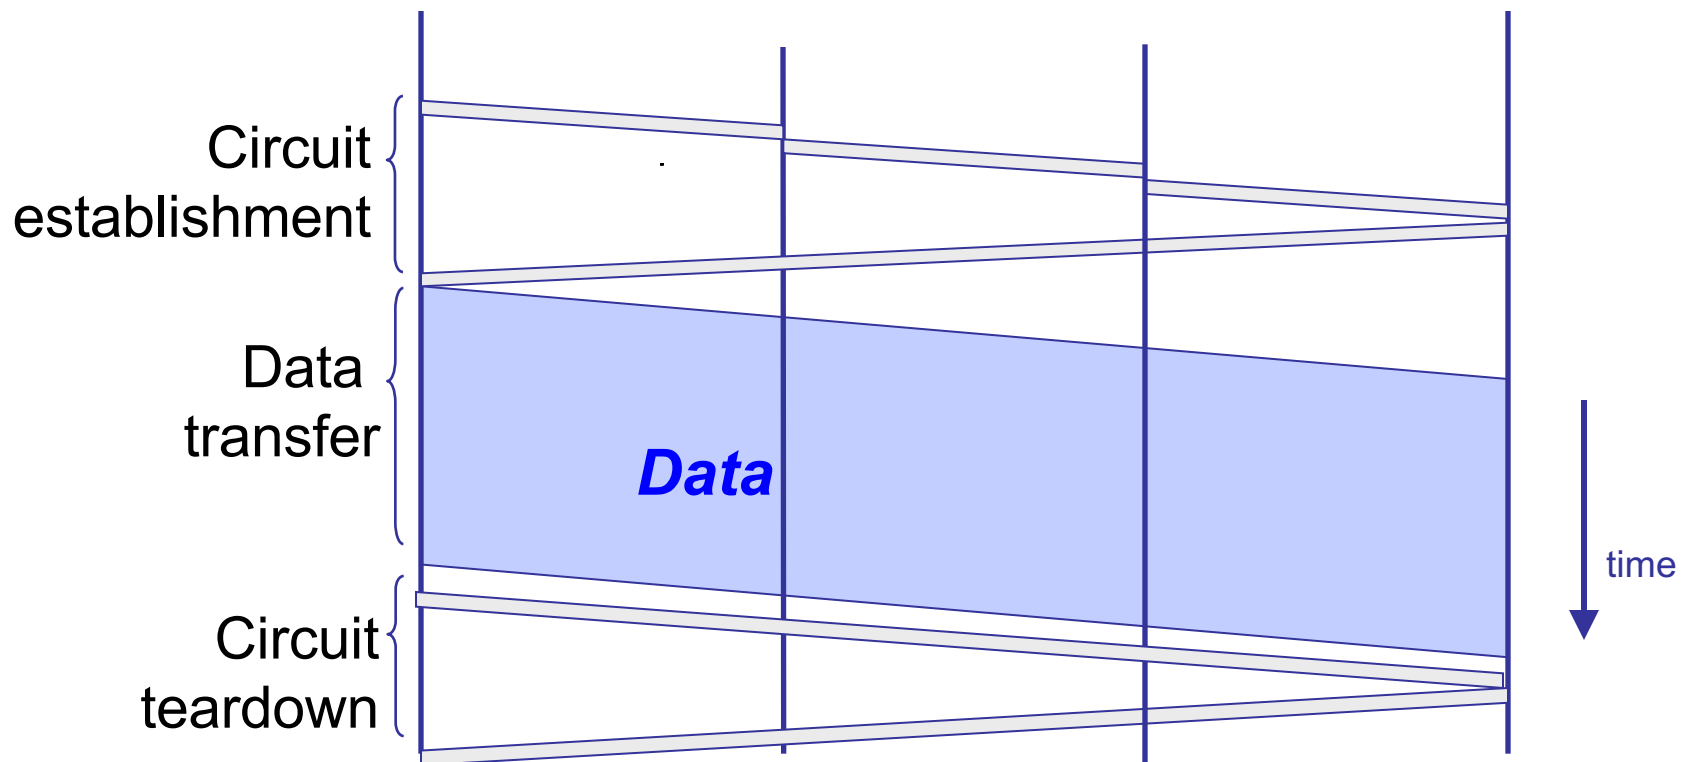
# Timing in circuit switching



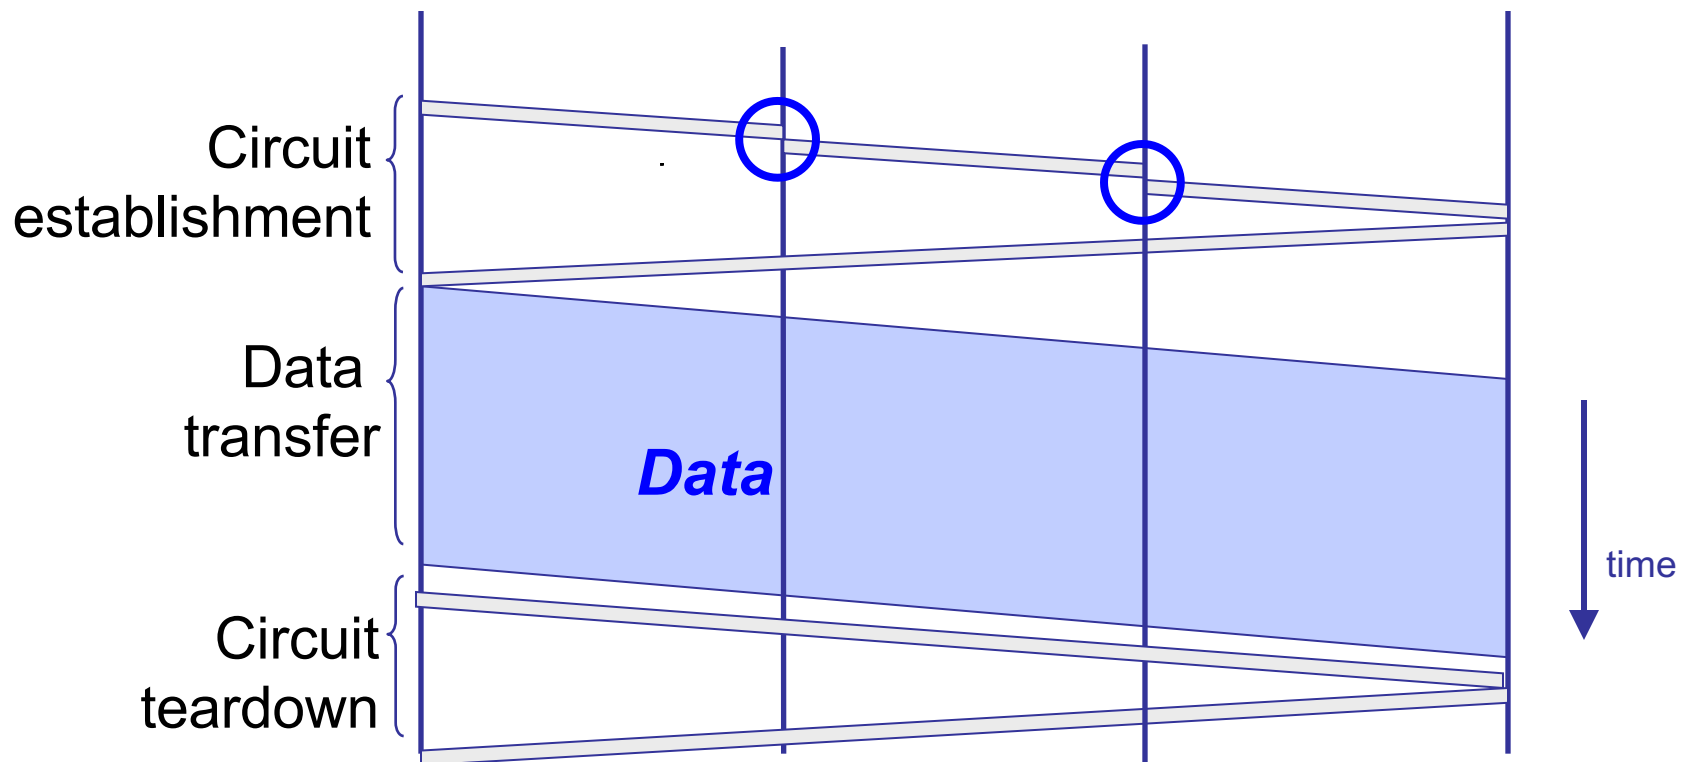
# Timing in circuit switching



# Timing in circuit switching

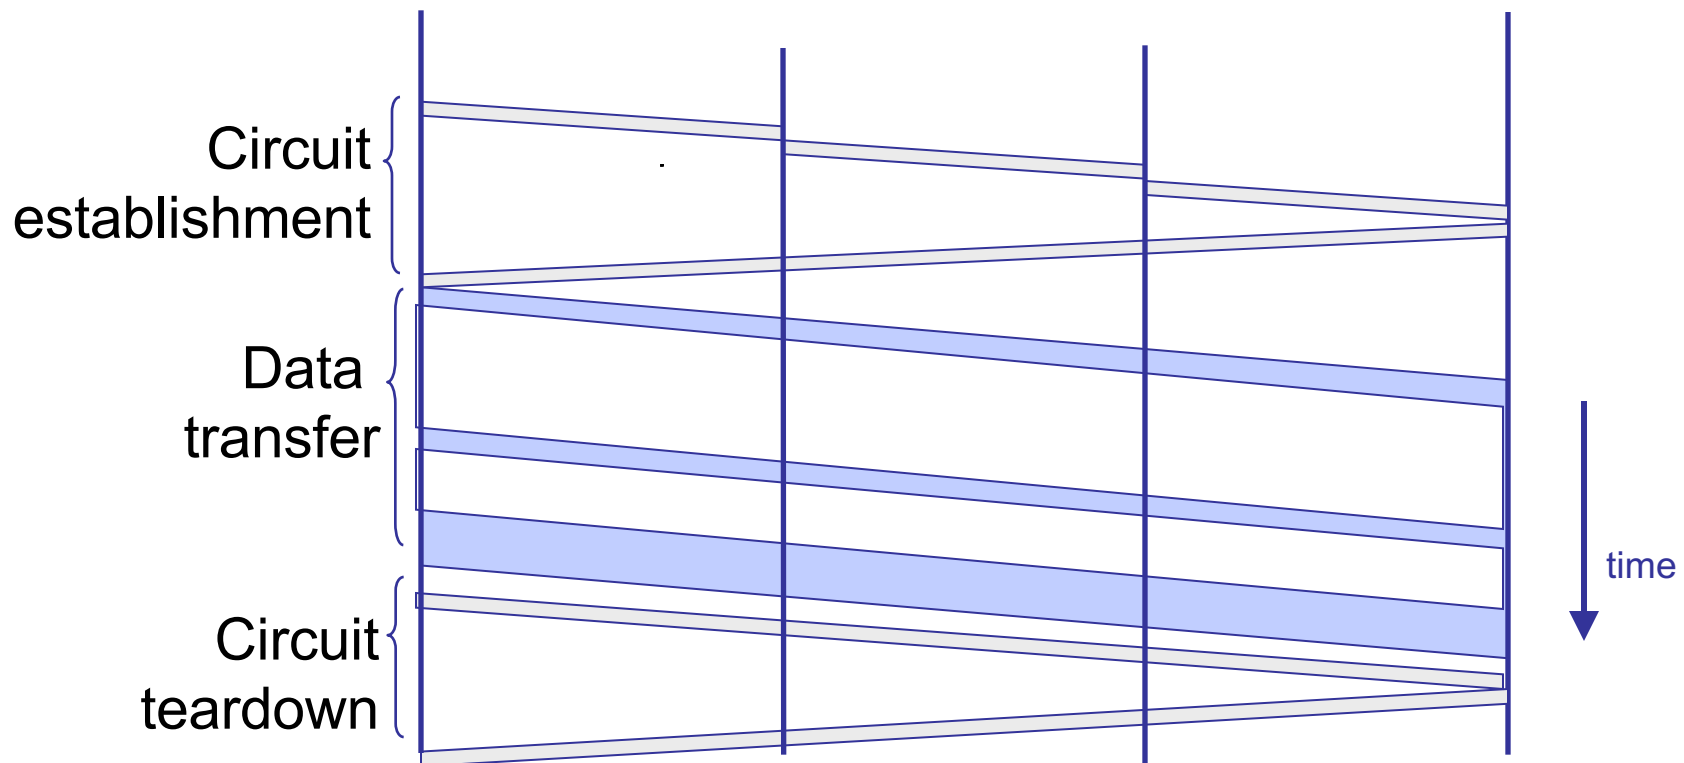


# Why the delays?

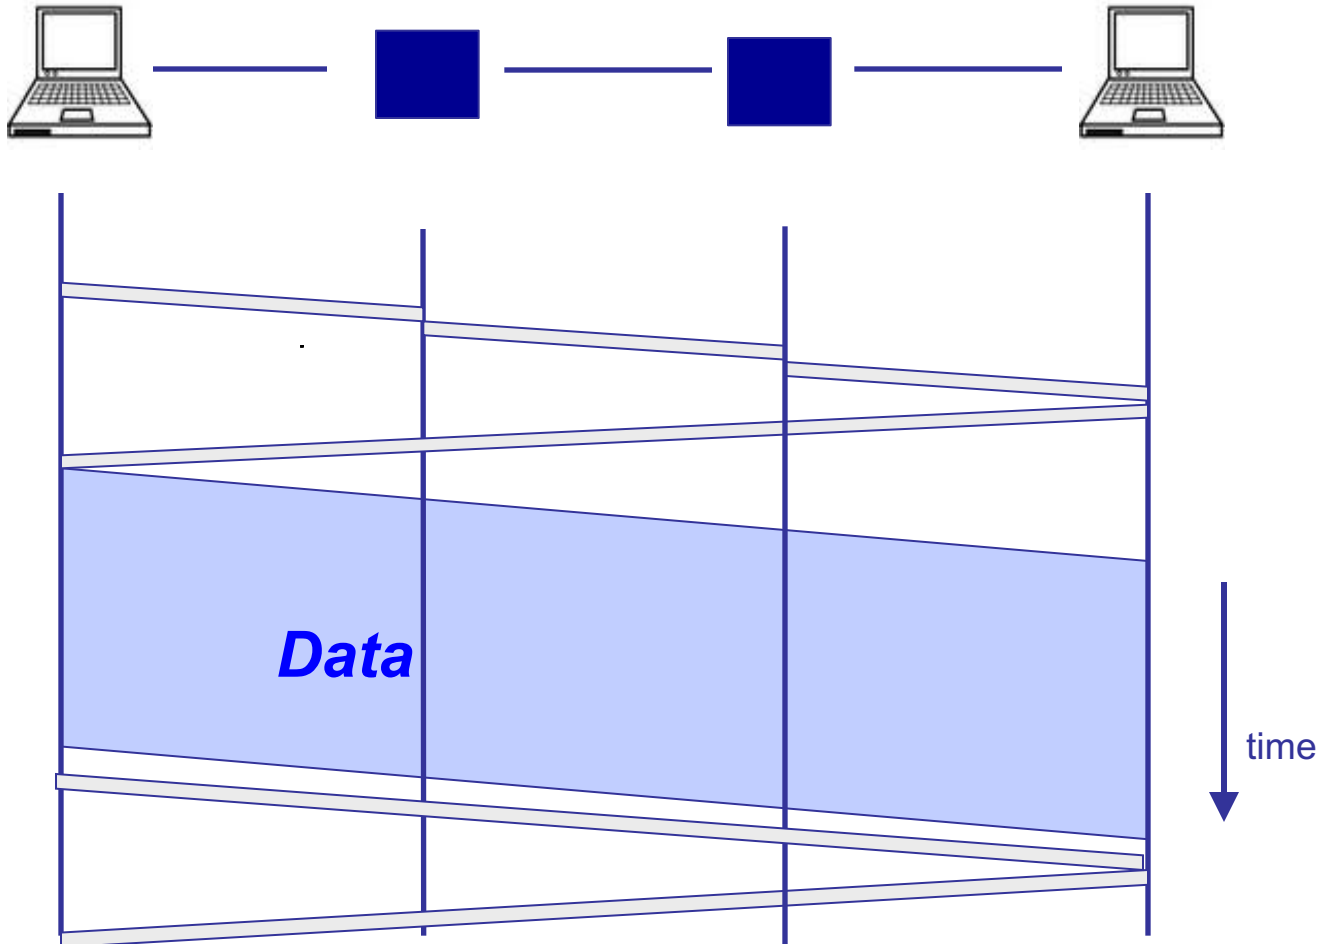




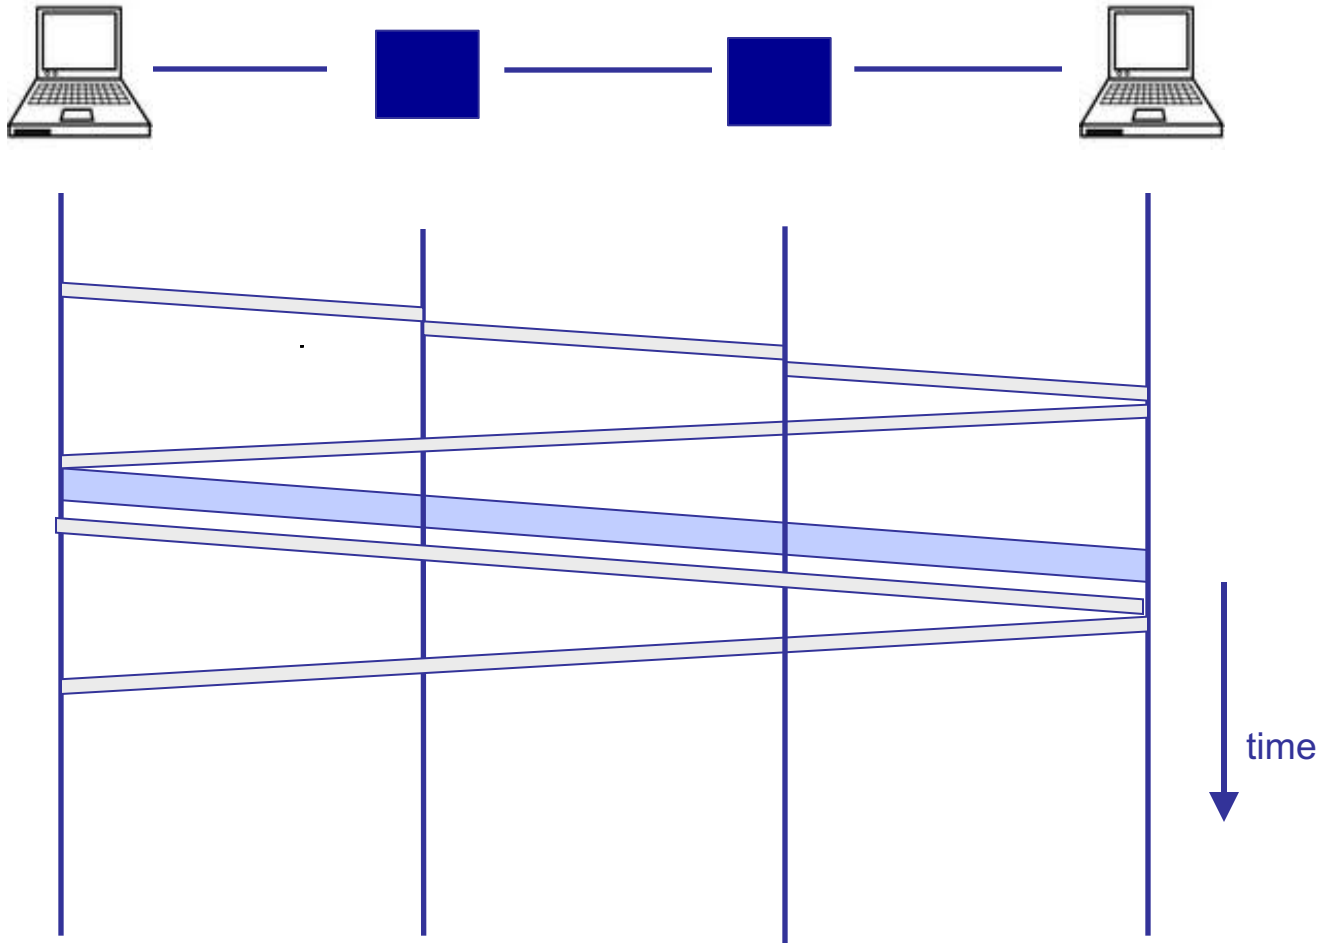
# Timing in circuit switching



# Timing in circuit switching



# Timing in circuit switching





# A network link: BDP

---



- Link bandwidth
  - Number of bits sent/received per unit time (bits/sec or bps)
- Propagation delay
  - Time for one bit to move through the link (seconds)
- Bandwidth-Delay Product (BDP)
  - Number of bits “in flight” at any time
- $\text{BDP} = \text{bandwidth} \times \text{propagation delay}$

# BDP Examples

---

- Same city over a slow link:
  - Bandwidth:  $\sim 100\text{Mbps}$
  - Propagation delay:  $\sim 0.1\text{msec}$
  - BDP:  $10,000\text{bits}$  ( $1.25\text{KBytes}$ )
- Cross-country over fast link:
  - Bandwidth:  $\sim 10\text{Gbps}$
  - Propagation delay:  $\sim 10\text{msec}$
  - BDP:  $10^8\text{bits}$  ( $12.5\text{MBytes}$ )