

EECS 489

Computer Networks

Fall 2020

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Topics

- Basics (lectures 1–2)
- Application layer (lectures 3–5)
 - HTTP, DNS, and CDN
 - Video Streaming

Basic concepts

- You should know:
 - Packet vs. circuit switching
 - Statistical multiplexing
 - Link characteristics
 - Packet delays

Switched networks

- End-systems and networks connected by switches instead of directly connecting them
- Allows us to **scale**
 - For example, directly connecting N nodes to each other would require N^2 links!

How are network resources shared?


- Two approaches
 - Reservations → circuit switching
 - On-demand → packet switching

Two approaches to sharing

- Packet switching
 - Network resources consumed on demand per-packet
 - Admission control: per packet
- Circuit switching
 - Network resources reserved a priori at “connection” initiation
 - Admission control: per connection

Circuit switching

1. **src** sends reservation request to **dst**
2. Switches create circuit *after* admission control
3. **src** sends data
4. **src** sends teardown request



Packet switching

- Data is sent as chunks of formatted bits (Packets)
- Packets consist of a “header” and “payload”
- Switches “forward” packets based on their headers
- Each packet travels **independently**
- No link resources are reserved

Statistical multiplexing

- Allowing more demands than the network can handle
 - Hoping that not all demands are required at the same time
 - Good for bursty traffic (average \ll peak demand)
 - Packet switching exploits statistical multiplexing better than circuit switching

Performance metrics

- Delay
- Loss
- Throughput

Delay

- Consists of four components

- Transmission delay
- Propagation delay
- Queuing delay
- Processing delay

} due to link properties

} due to traffic mix and switch internals

A network link




- Link bandwidth
 - Number of bits sent/received per unit time (bits/sec or bps)
- Propagation delay
 - Time for one bit to move through the link (seconds)

Queueing delay

- How long does a packet have to sit in a buffer before it is processed?
- Characterized with statistical measures
 - Average queuing delay
 - Variance of queuing delay
 - Probability delay exceeds a threshold value

End-to-end delay



What we want


`http://123.xyz`




123.xyz server



(Some of) What happens...



(More of) What happens



What we get



123.xyz server

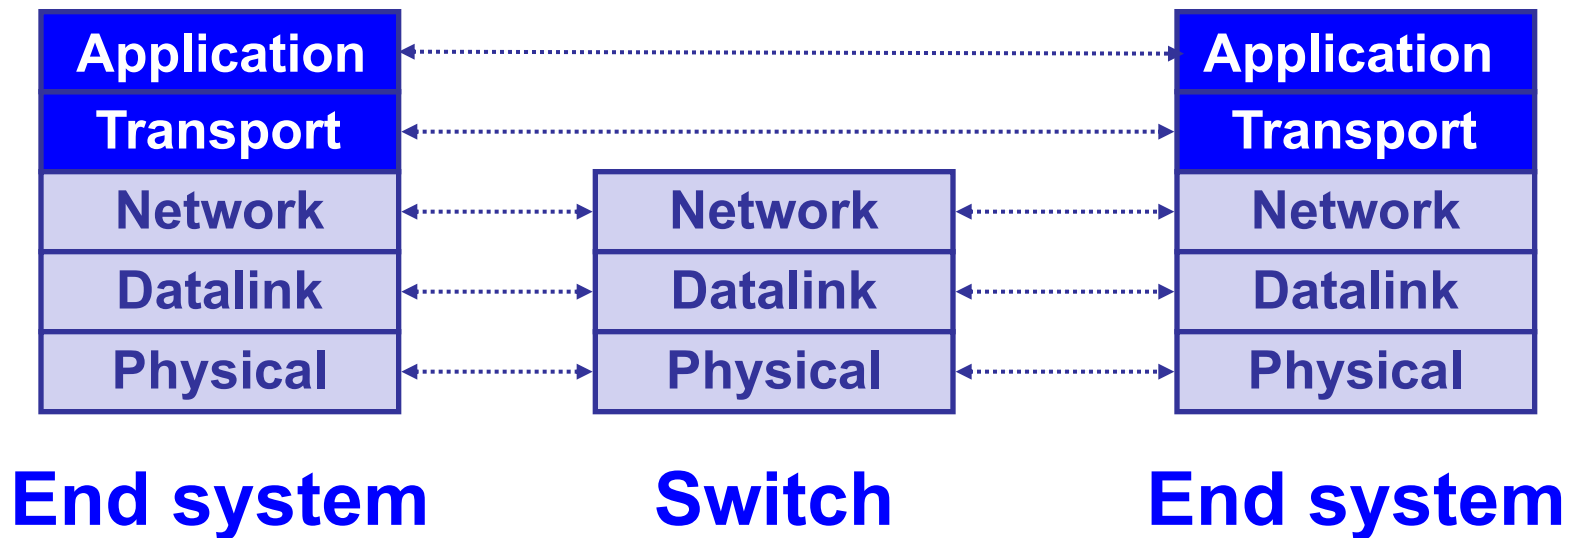


Layers

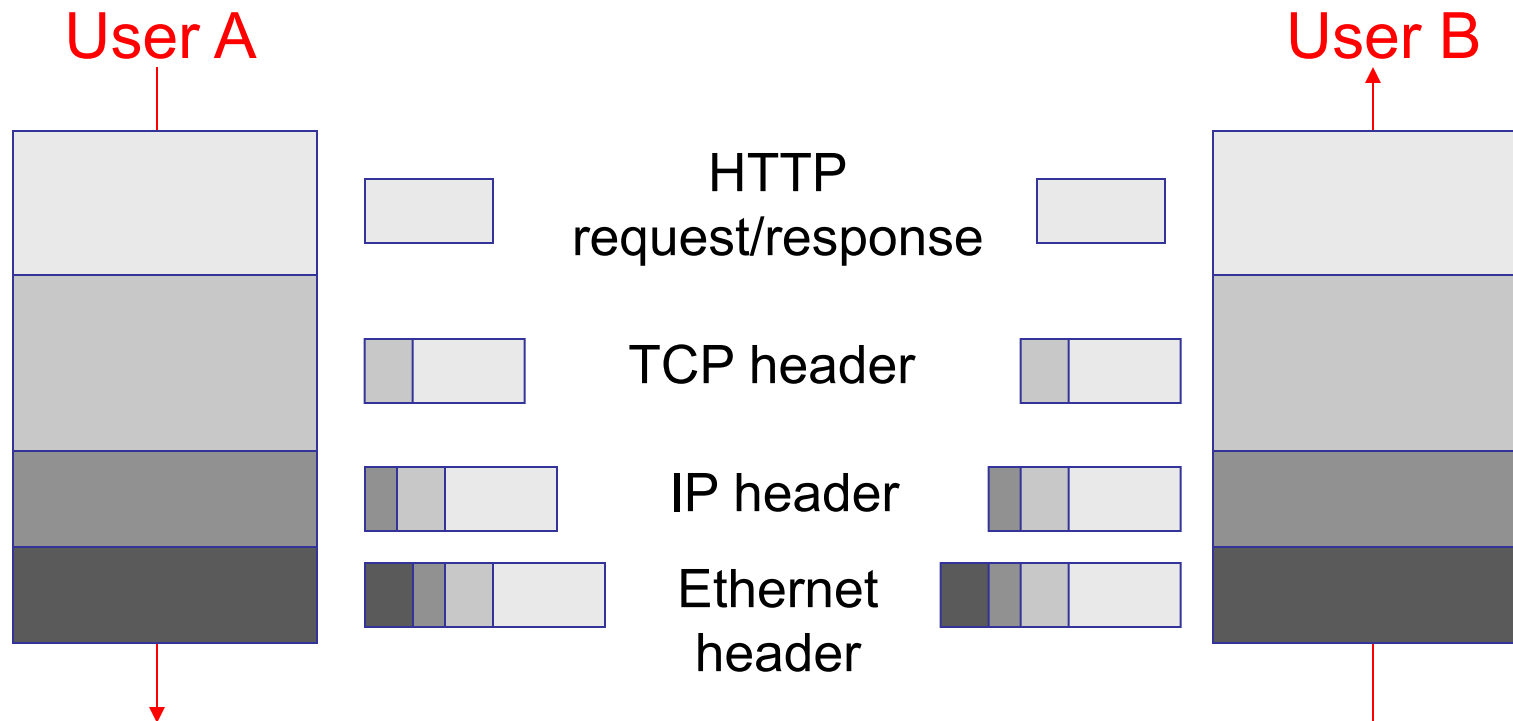
- Layer: a part of a system with well-defined interfaces to other parts
- One layer interacts only with layer above and layer below
- Two layers interact only through the interface between them

Layers in practice


- Lower three layers implemented everywhere
- Top two layers implemented only at hosts



Layer encapsulation: Protocol headers



IP is the narrow waist of the layering hourglass



Placing network functionality

- End-to-end arguments by Saltzer, Reed, and Clark
 - Dumb network and smart end systems
 - Functions that can be *completely* and *correctly* implemented *only* with the knowledge of application end host, should not be pushed into the network
 - Sometimes necessary to break this for performance and policy optimizations
 - **Fate sharing**: fail together or don't fail at all

5-MINUTE BREAK!

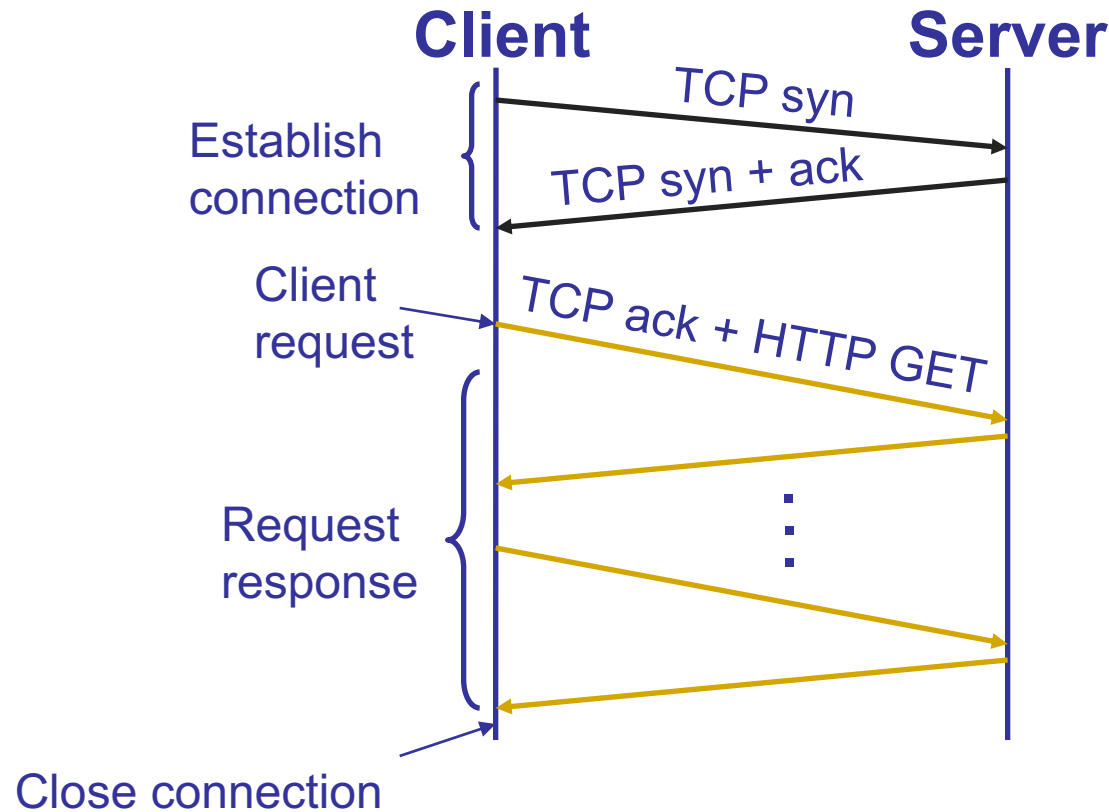
Topics

- Basics (lectures 1–2)
- Application layer (lectures 3–5)
 - HTTP, DNS, and CDN
 - Video Streaming

Hyper Text Transfer Protocol (HTTP)


- Client-server architecture
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- Synchronous request/reply protocol
 - Runs over TCP, Port 80
- Stateless
- ASCII format
 - Before HTTP/2

Steps in HTTP request/response



Object request response time

- RTT (round-trip time)
 - Time for a small packet to travel from client to server and back
- Response time
 - 1 RTT for TCP setup
 - 1 RTT for HTTP request and first few bytes
 - Transmission time
 - **Total** = 2RTT + Transmission Time



Improving HTTP performance

- Optimizing connections using three “P”s
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection
- Caching
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- Replication

Scorecard: Getting n small objects

- Time dominated by latency
- One-at-a-time: $\sim 2n$ RTT
- m concurrent: $\sim 2\lceil n/m \rceil$ RTT
- Persistent: $\sim (n+1)$ RTT
- Pipelined: ~ 2 RTT
- Pipelined and Persistent: ~ 2 RTT first time; RTT later for another n from the same site

Scorecard: Getting n large objects each of size F

- Time dominated by TCP throughput B_C ($\leq B_L$), where link bandwidth is referred by B_L
- One-at-a-time: $\sim nF/B_C$
- m concurrent: $\sim nF/(mB_C)$
 - Assuming each TCP connection gets the same throughput and $mB_C \leq B_L$
- Pipelined and/or persistent: $\sim nF/B_C$
 - The only thing that helps is higher throughput


Content Distribution Networks (CDN)

- Caching and replication as a service
- Combination of caching and replication
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)


Hierarchies in the DNS

- Three intertwined hierarchies
 - Hierarchical namespace
 - » As opposed to flat namespace
 - Hierarchically administered
 - » As opposed to centralized
 - (Distributed) hierarchy of servers
 - » As opposed to centralized storage


Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



DNS caching

- Performing all these queries takes time
 - Up to 1-second latency before starting download
- Caching can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.google.com) visited often
 - Local DNS server often has the information cached
- How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

HTTP streaming

- Video is stored at an HTTP server with a URL
- Clients send a GET request for the URL
- Server sends the video file as a stream
- Client first buffers for a while to minimize interruptions later
- Once the buffer reaches a threshold
 - The video plays in the foreground
 - More frames are downloaded in the background


DASH : Dynamic Adaptive Streaming over HTTP

- Keep multiple resolutions of the same video
 - Stored in a manifest file in the HTTP server
- Client asks for the manifest file first to learn about the options
- Asks for chunks at a time and measures available bandwidth while they are downloaded
 - Low bandwidth \Rightarrow switch to lower bitrate
 - High bandwidth \Rightarrow switch to higher bitrate

Applications

- Common theme: **parallelism**
 - Applications decomposed into tasks
 - Running in parallel on different machines
- Two common paradigms
 - Partition-Aggregate
 - Map-Reduce


Traditional datacenter networks



Challenges

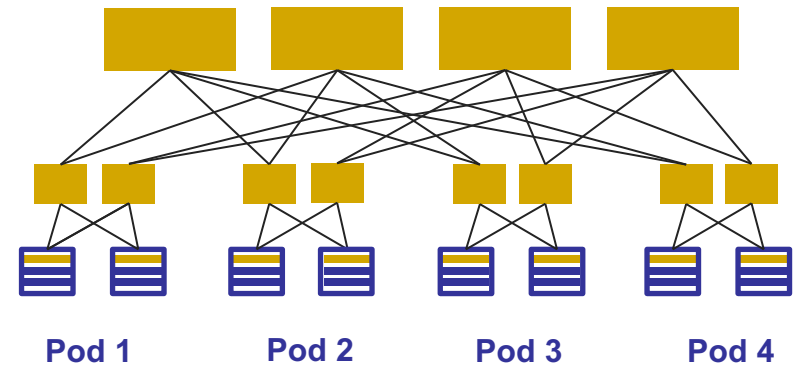
- Not enough bandwidth
 - **Oversubscription**: Less bandwidth in the ToR-Agg links than all the servers' bandwidth in the rack
 - **Oversubscription ratio**: Ratio between bandwidth underneath and bandwidth above
- Not enough paths between server pairs
 - Load balancing issues
 - Failure recovery issues

Modern datacenter networks: More bandwidth, more paths



Clos topology

- Multi-stage network
- k pods, where each pod has two layers of $k/2$ switches
 - $k/2$ ports up and $k/2$ down
- All links have the same b/w
- At most $k^3/4$ machines
- Example
 - $k = 4$
 - 16 machines
- For $k=48$, 27648 machines



Summary

- We will focus on the transport layer from next class