

EECS 489

Computer Networks

Fall 2020

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- From reliable data transfer to TCP
- TCP connection setup and teardown

Recap: Designing a reliable transport protocol

- Stop and wait is correct but inefficient
- Sliding window uses pipelining to increase throughput

Recap: Acknowledgements

- Cumulative
 - Acknowledge many packets at a time
- Selective
 - Acknowledge individual packets

Recap: Sliding window protocols

- Resending packets: two canonical approaches
 - Go-Back-N: Resend all N packets
 - Selective Repeat: Resend only the missing packets
- Many variants that differ in implementation details

TCP: TRANSMISSION CONTROL PROTOCOL

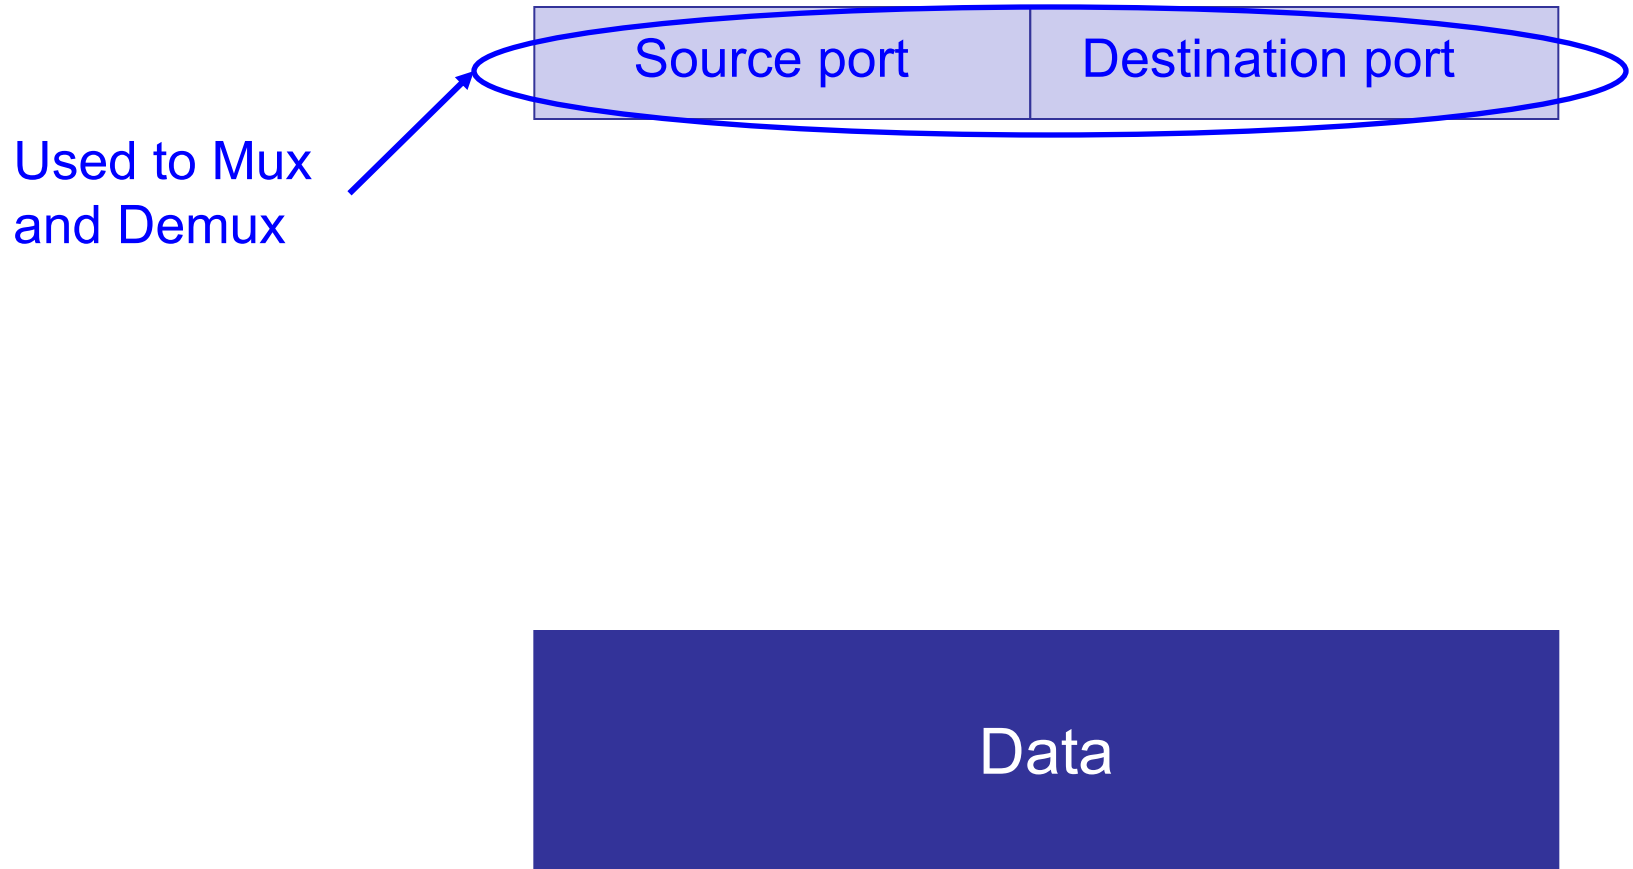
The TCP Abstraction

- TCP delivers a reliable, in-order, byte stream
- **Reliable**: TCP resends lost packets (recursively)
 - Until it gives up and shuts down connection
- **In-order**: TCP only hands consecutive chunks of data to application
- **Byte stream**: TCP assumes there is an incoming stream of data, and attempts to deliver it to app

What does TCP use from what we've seen so far?

- Most of what we've seen
 - Checksums
 - Sequence numbers are byte offsets
 - Sender and receiver maintain a sliding window
 - Receiver sends cumulative acknowledgements (like GBN)
 - » Sender maintains a single retransmission timer
 - Receivers buffer out-of-sequence packets (like SR)
- Few more: fast retransmit, timeout estimation algorithms etc.

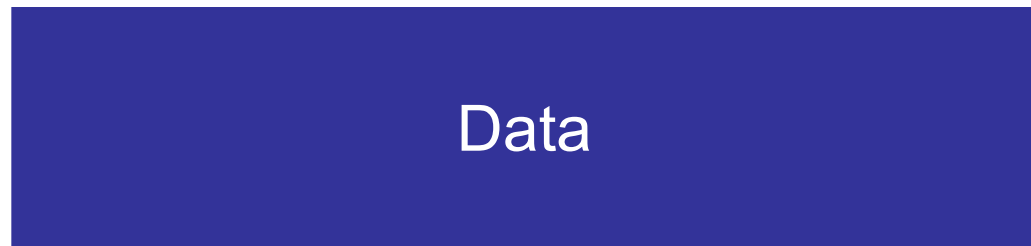
Build the TCP header



Build the TCP header



Computed
over pseudo-header
and data

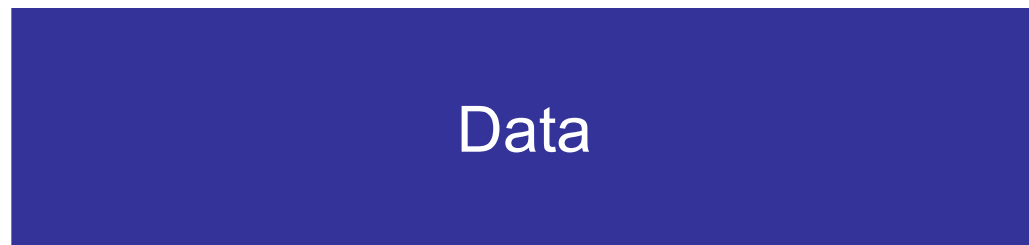
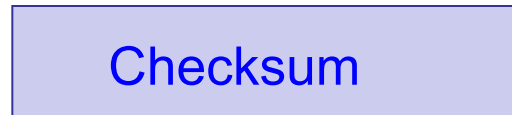
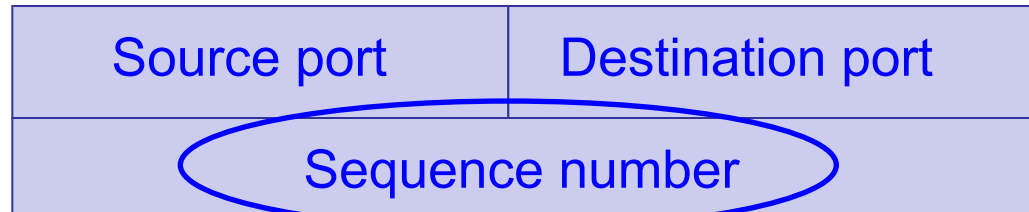


What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets

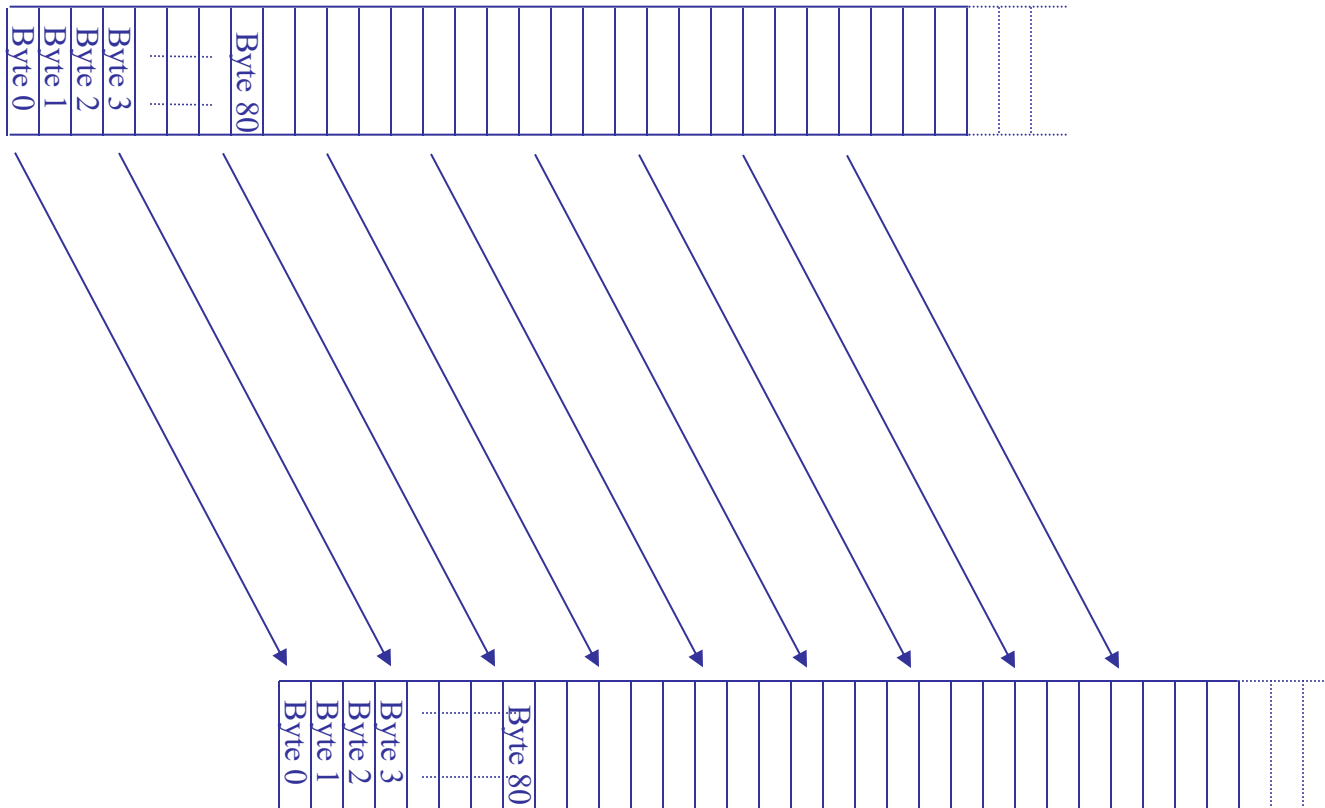
Build the TCP header

Byte offsets
(NOT packet id),
because TCP is a
byte stream



TCP “stream of bytes” service...

Application @ Host A


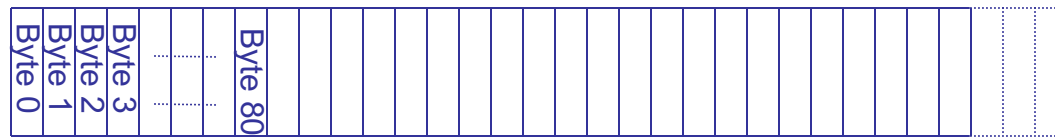


Application @ Host B

... provided using TCP

“segments”


Host A



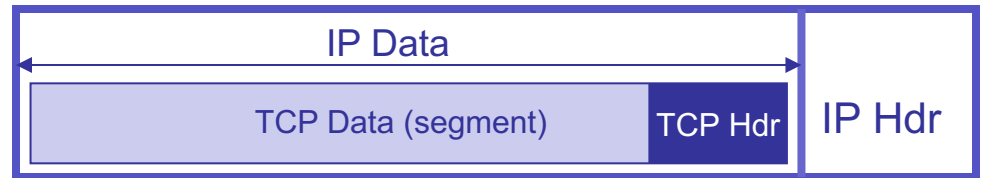
Segment sent when:

1. Segment full (Max Segment Size),
2. Not full, but times out

Host B

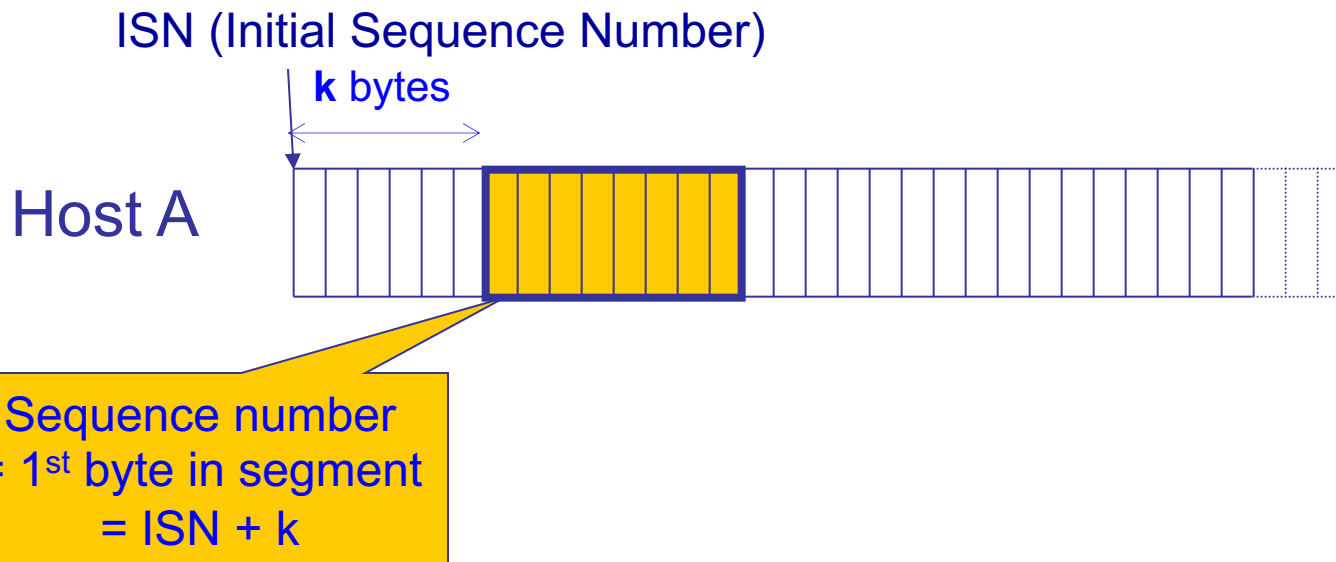


TCP segment

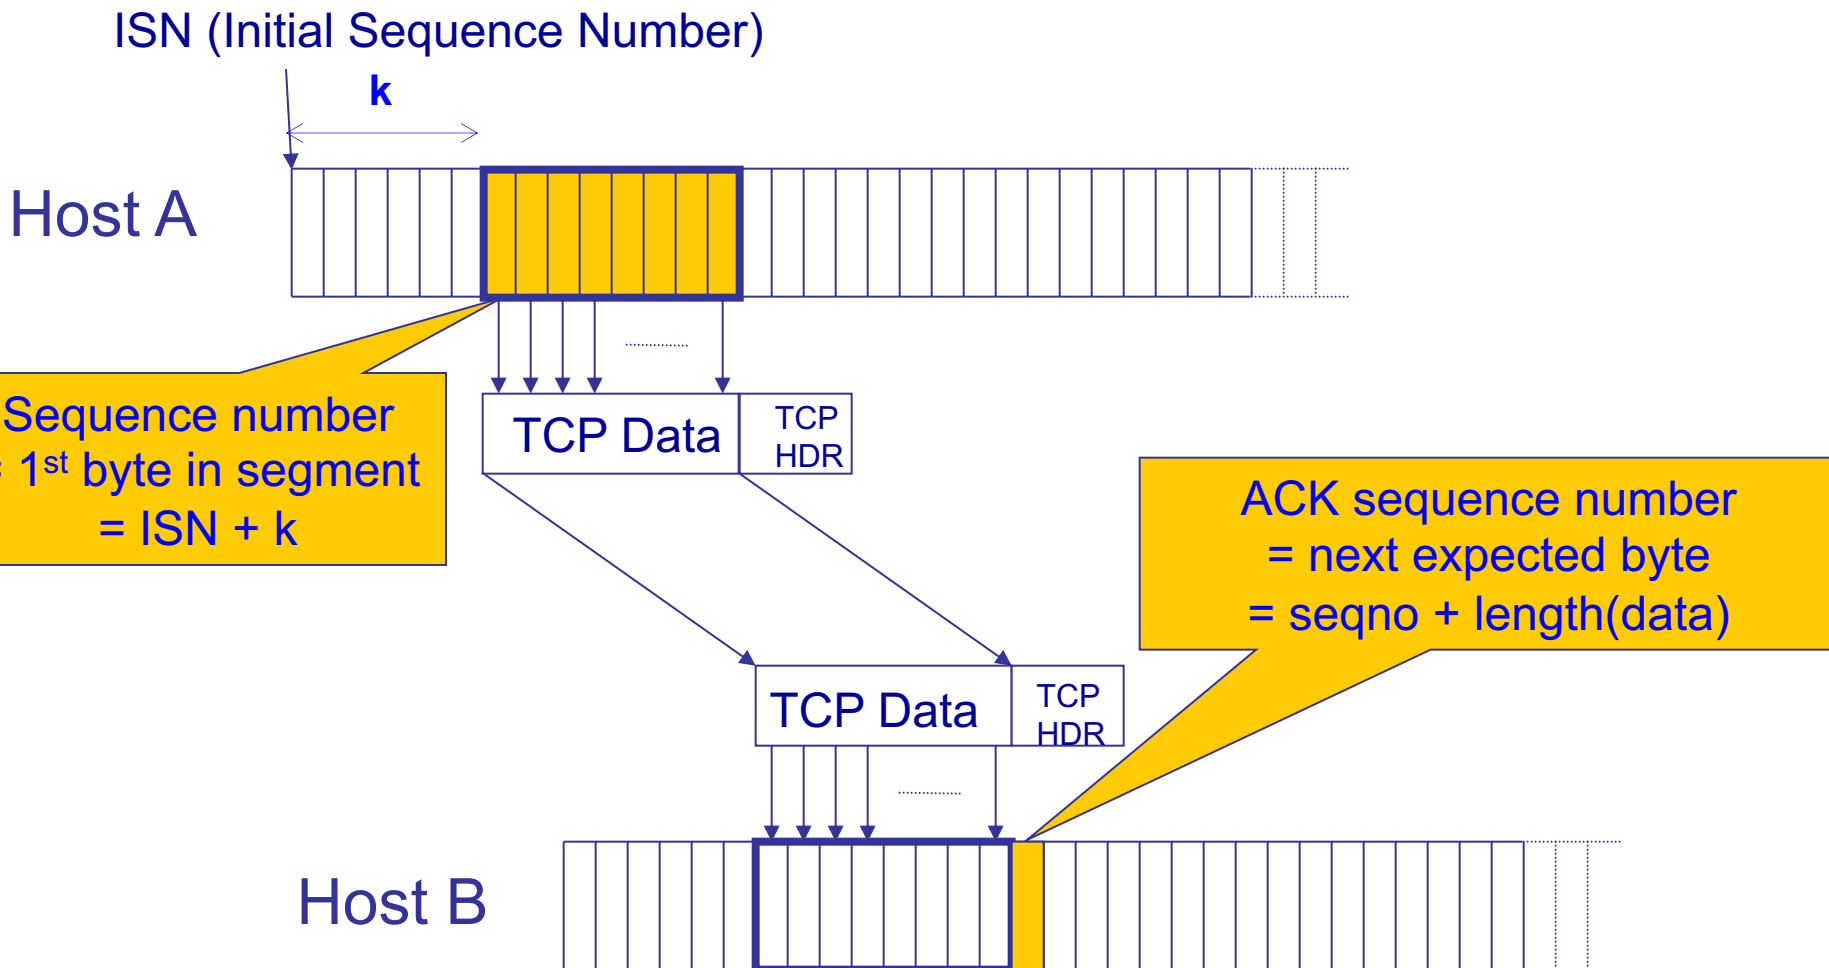


- IP packet
 - No bigger than **Maximum Transmission Unit (MTU)**
 - E.g., up to 1500 bytes with Ethernet
- TCP packet
 - IP packet with a TCP header and data inside
 - TCP header ≥ 20 bytes long
- TCP segment
 - No more than **Maximum Segment Size (MSS)** bytes
 - E.g., up to 1460 consecutive bytes from the stream
 - $MSS = MTU - (IP \text{ header}) - (TCP \text{ header})$

Sequence numbers





Sequence numbers



Build the TCP header

Starting byte
offset of data
carried in this
segment



What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends **cumulative acknowledgements** (like GBN)

ACKs and sequence numbers

- Sender sends packet
 - Data starts with sequence number X
 - Packet contains B bytes $[X, X+1, X+2, \dots, X+B-1]$
- Upon receipt of packet, receiver sends an ACK
 - If all data prior to X already received:
 - » ACK acknowledges $X+B$ (because that is next expected byte)
 - If highest in-order byte received is Y s.t. $(Y+1) < X$
 - » ACK acknowledges $Y+1$
 - » Even if this has been ACKed before

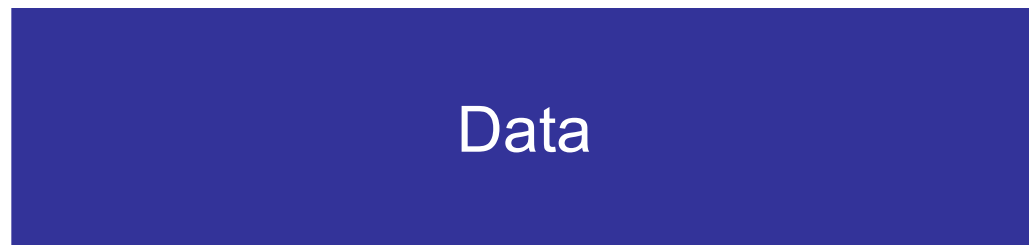
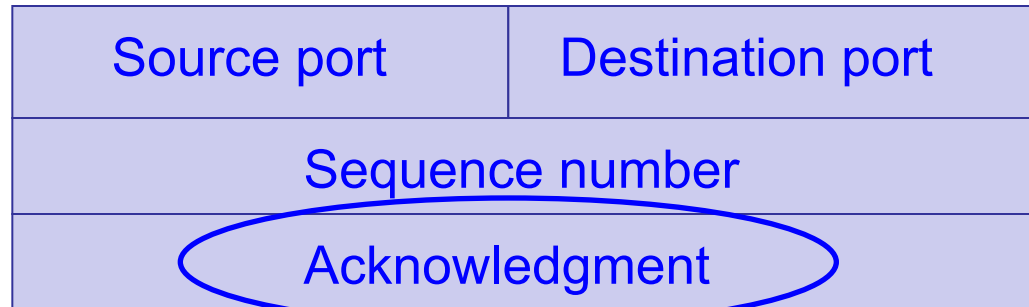
Typical operation

- Sender: $\text{seqno} = X$, $\text{length} = B$
- Receiver: $\text{ACK} = X + B$
- Sender: $\text{seqno} = X + B$, $\text{length} = B$
- Receiver: $\text{ACK} = X + 2B$
- Sender: $\text{seqno} = X + 2B$, $\text{length} = B$

- Seqno of next packet is same as last ACK field

Build the TCP header

Acknowledgment
gives seqno just
beyond highest
seqno received
in order



What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers **can buffer out-of-sequence packets** (like SR)

Loss with cumulative ACKs

- Sender sends packets with 100B and seqnos.:
 - 100, 200, 300, 400, 500, 600, 700, 800, ...
- Assume the fifth packet (seqno 500) is lost, but no others
- Stream of ACKs will be:
 - 200, 300, 400, 500, 500, 500, 500, ...

What does TCP introduce?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers can buffer out-of-sequence packets (like SR)
- Introduces **fast retransmit**: duplicate ACKs trigger early retransmission

Loss with cumulative ACKs

- Duplicate ACKs are a sign of an isolated loss
 - The lack of ACK progress means 500 hasn't been delivered
 - Stream of ACKs means some packets are being delivered
- Trigger retransmission upon receiving k duplicate ACKs
 - » TCP uses $k=3$
 - » Faster than waiting for timeout

Loss with cumulative ACKs

- Two choices after resending
 - Send missing packet and move sliding window by the number of dup ACKs
 - » Speeds up transmission, but might be wrong
 - Send missing packet, and wait for ACK to move sliding window
 - » Is slowed down by single dropped packets
- Which should TCP do?
 - Choose correctness

5-MINUTE BREAK!

Announcements

- Assignment 2 is live!

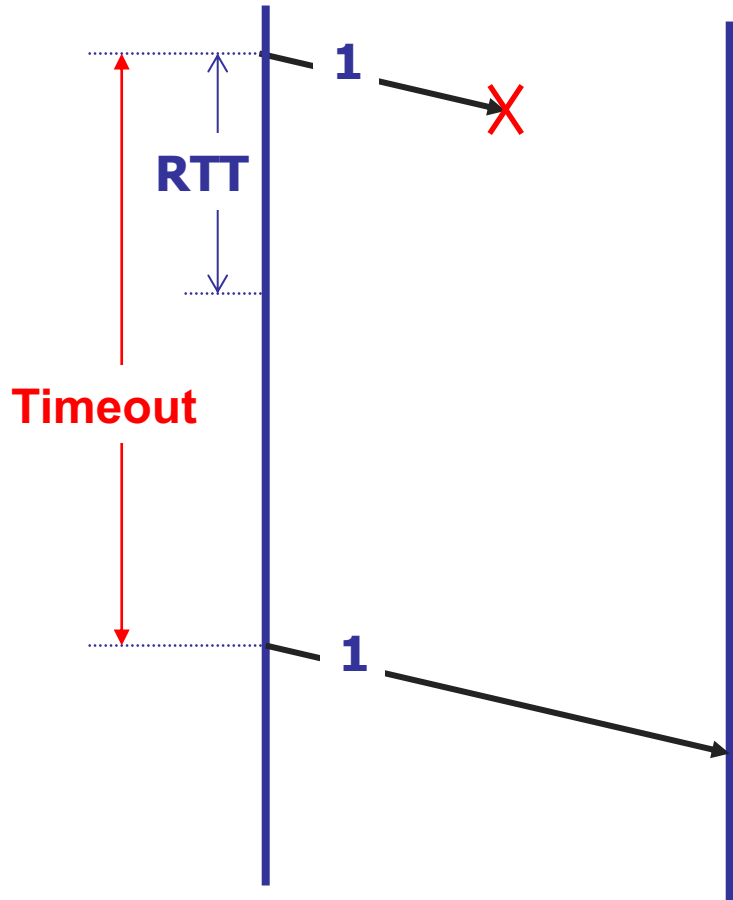
What does TCP introduce?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers buffer out-of-sequence packets (like SR)
- Introduces fast retransmit: duplicate ACKs trigger early retransmission
- Sender maintains a **single retransmission timer** (like GBN) and retransmits on timeout

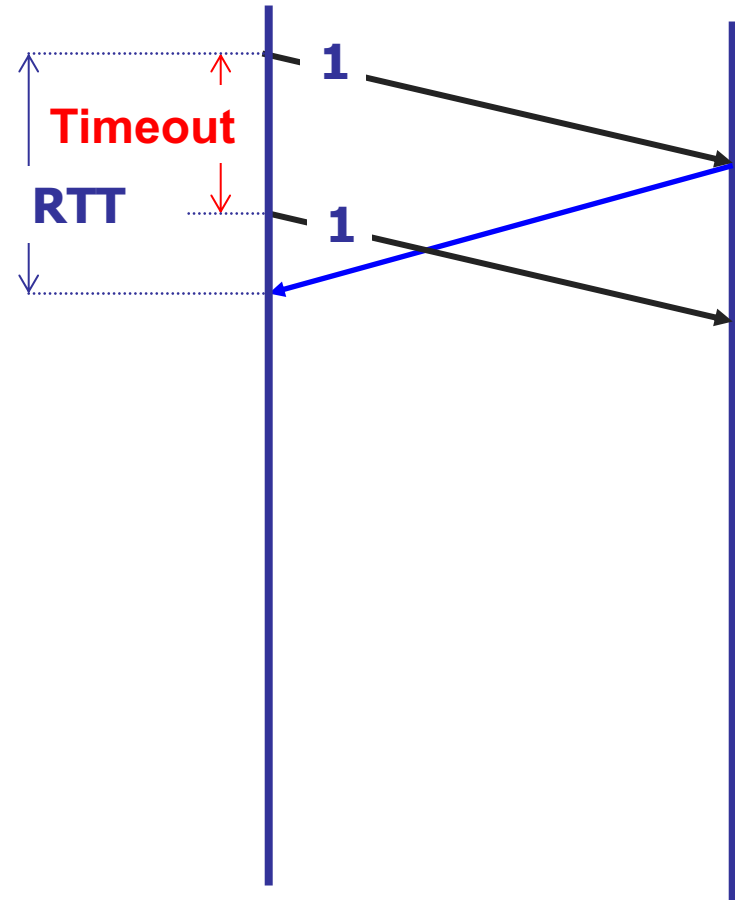
Retransmission timeout

- If the sender hasn't received an ACK by timeout, **retransmit the first packet** in the window
- How do we pick a timeout value?

Timing illustration



Timeout too long → inefficient



Timeout too short →
duplicate packets

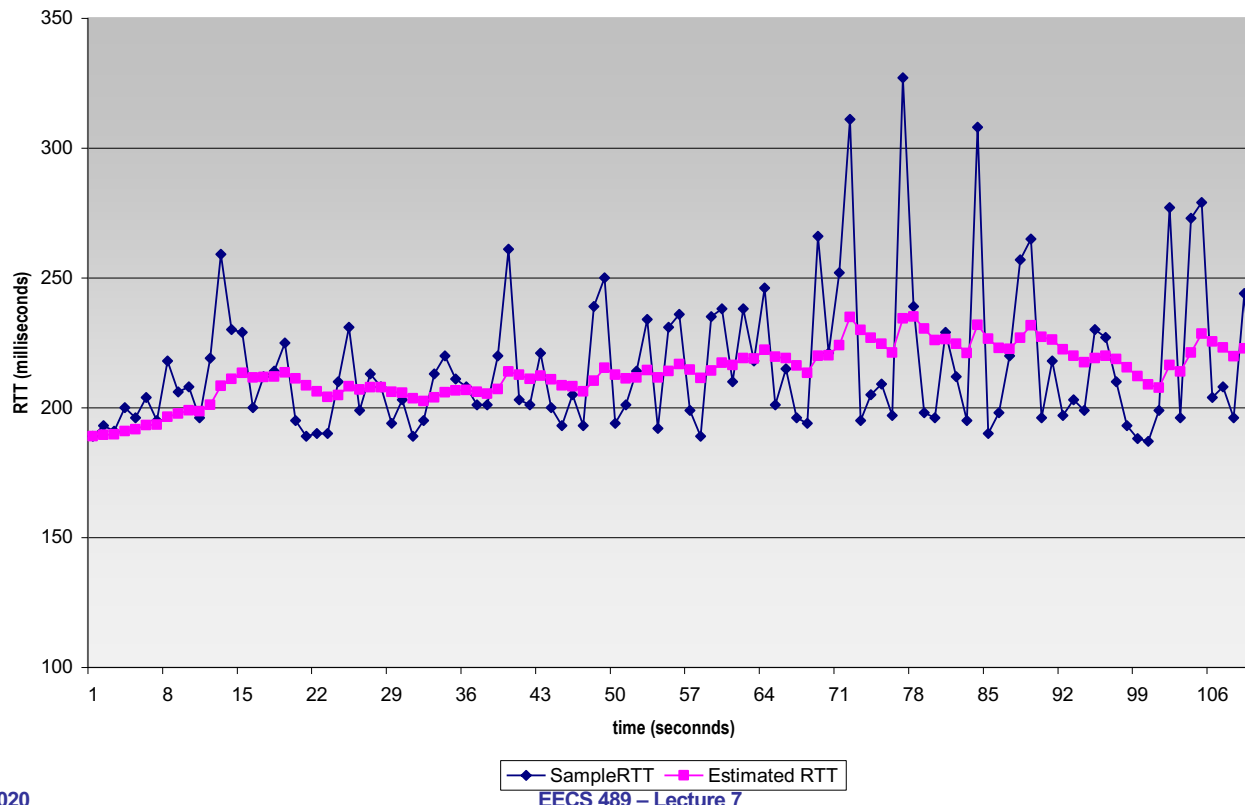
Retransmission timeout

- If the sender hasn't received an ACK by timeout, retransmit the first packet in the window
- How to set timeout?
 - Too long: connection has low throughput
 - Too short: retransmit packet that was just delayed
- Solution: **make timeout proportional to RTT**
 - But how do we measure RTT?

RTT estimation



- Exponential weighted average of RTT samples

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$



Problem: Ambiguous measurements

- How do we differentiate between the real ACK, and ACK of the retransmitted packet?



Karn/Partridge algorithm

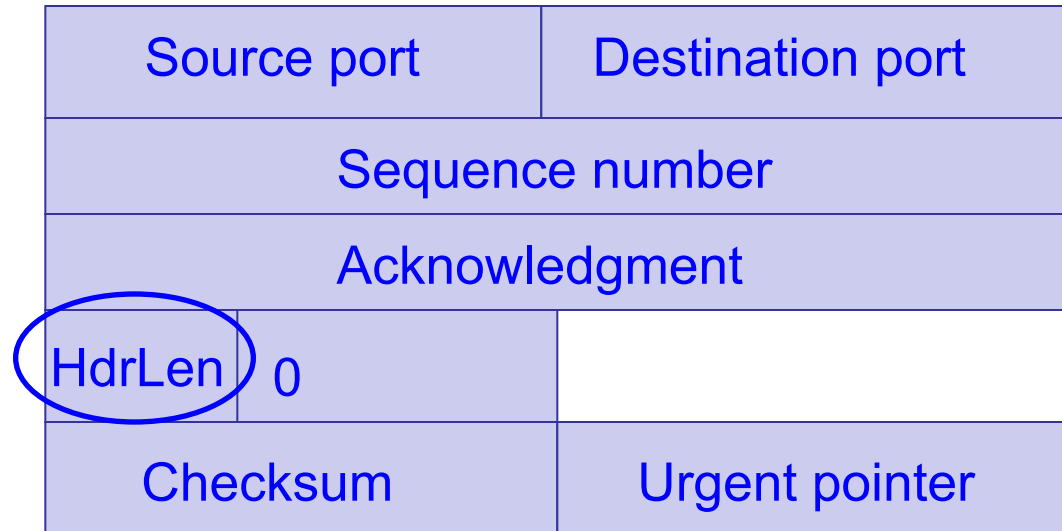
- Don't use SampleRTT from retransmissions
 - Once retransmitted, ignore that segment in the future
- Computes EstimatedRTT using $\alpha = 0.125$
- Timeout value (RTO) = $2 \times \text{EstimatedRTT}$
 - Employs exponential backoff
 - » Every time RTO timer expires, set $\text{RTO} \leftarrow 2 \cdot \text{RTO}$
 - (Up to maximum ≥ 60 sec)
 - » Every time new measurement comes in (= successful original transmission), collapse RTO back to $2 \times \text{EstimatedRTT}$
- Sensitive to RTT variations

Jacobson/Karels algorithm

- **Problem**: need to better capture variability in RTT
 - Directly measure deviation
- Deviation = $| \text{SampleRTT} - \text{EstimatedRTT} |$
- DevRTT: exponential average of Deviation
- $\text{RTO} = \text{EstimatedRTT} + 4 \times \text{DevRTT}$

Build the TCP header

Number of 4-
byte words in the
header;
5: No options



Data

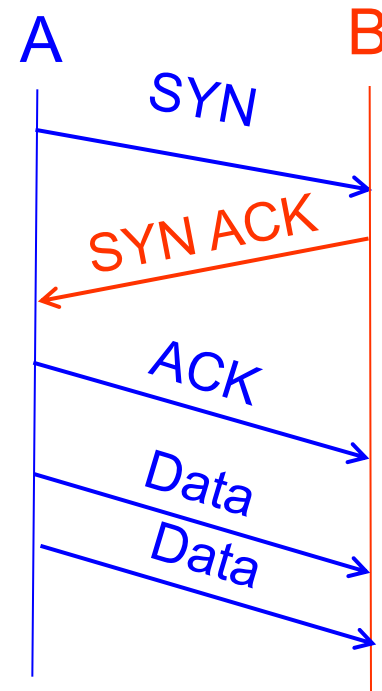
TCP CONNECTION ESTABLISHMENT

Initial Sequence Number (ISN)

- Sequence number for the very first byte
- Why not just use ISN = 0?
 - Practical issue
 - » IP addresses and port #s uniquely identify a connection
 - » Eventually, though, these port #s do get used again; small chance an old packet is still in flight
 - » Also, others might try to spoof your connection
 - Why does using ISN help?
- Hosts exchange ISNs when establishing connection

Establishing a TCP connection

- Three-way handshake to establish connection
 - Host A sends a SYN (open; “synchronize sequence numbers”) to host B
 - Host B returns a SYN acknowledgment (SYN ACK)
 - Host A sends an ACK to acknowledge the SYN ACK



Build the TCP header

Flags:

SYN

ACK

FIN

RST

PSH

URG

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	
Checksum		Urgent pointer	

Data

Step 1: A's initial SYN packet

A tells B to open
a connection

A's port			B's port
A's Initial Sequence Number			
N/A			
5	0	SYN	
Checksum			Urgent pointer

Step 1: B's SYN-ACK packet

B tells it accepts
and is ready to
accept next
packet

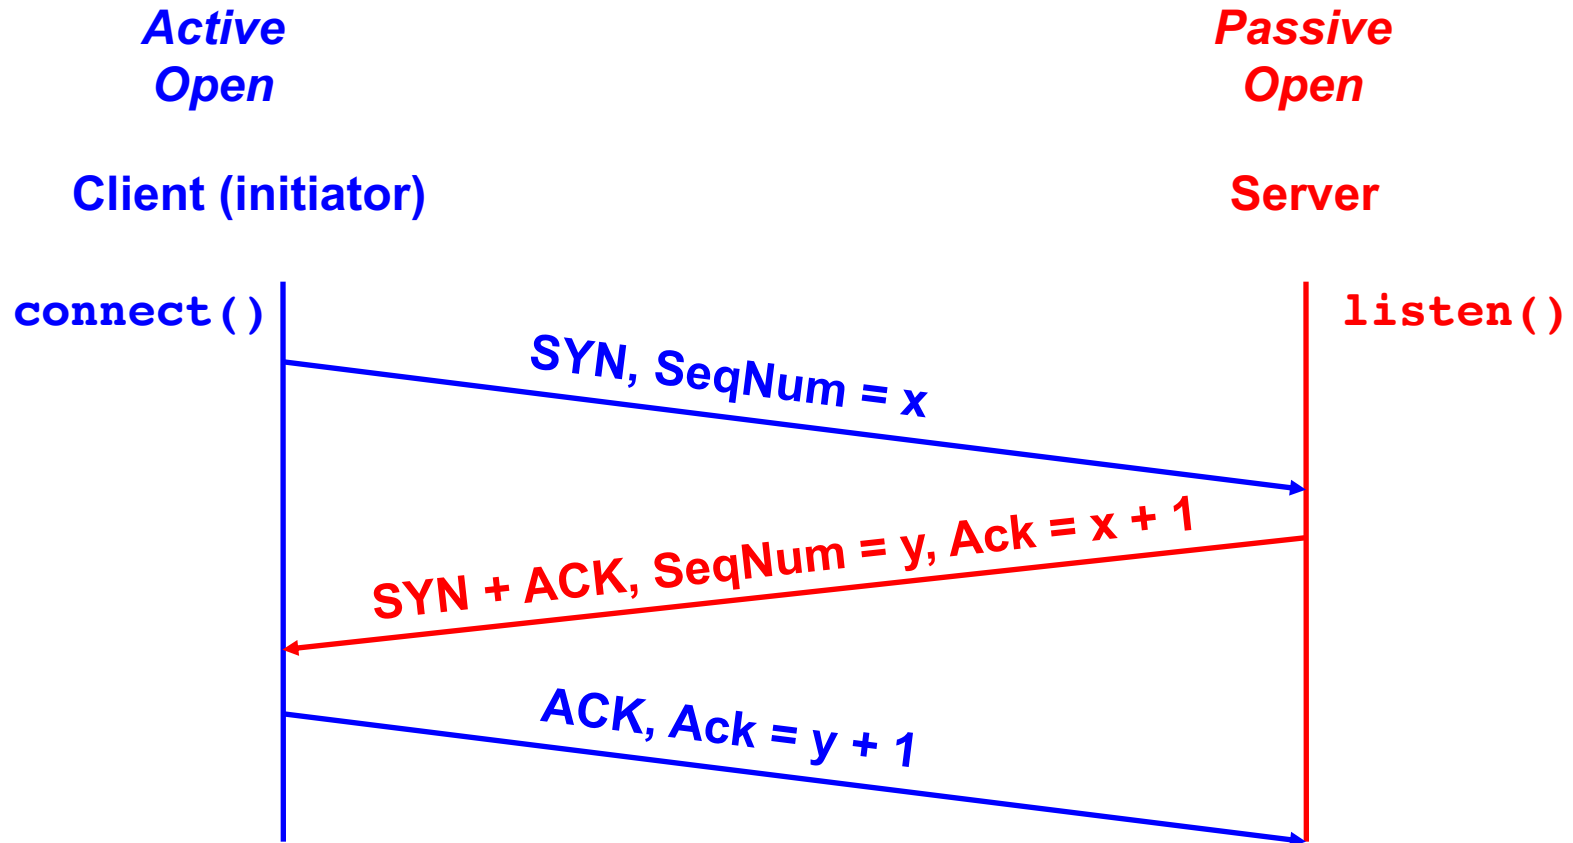
B's port		A's port	
B's Initial Sequence Number			
ACK=A's ISN+1			
5	0	SYN ACK	
Checksum		Urgent pointer	

Step 1: A's ACK to SYN-ACK

A tells B to open
a connection

A's port			B's port
A's Initial Sequence Number + 1			
ACK=B's ISN+1			
5	0	ACK	
Checksum			Urgent pointer

TCP's 3-Way handshaking



What if the SYN Packet Gets Lost?

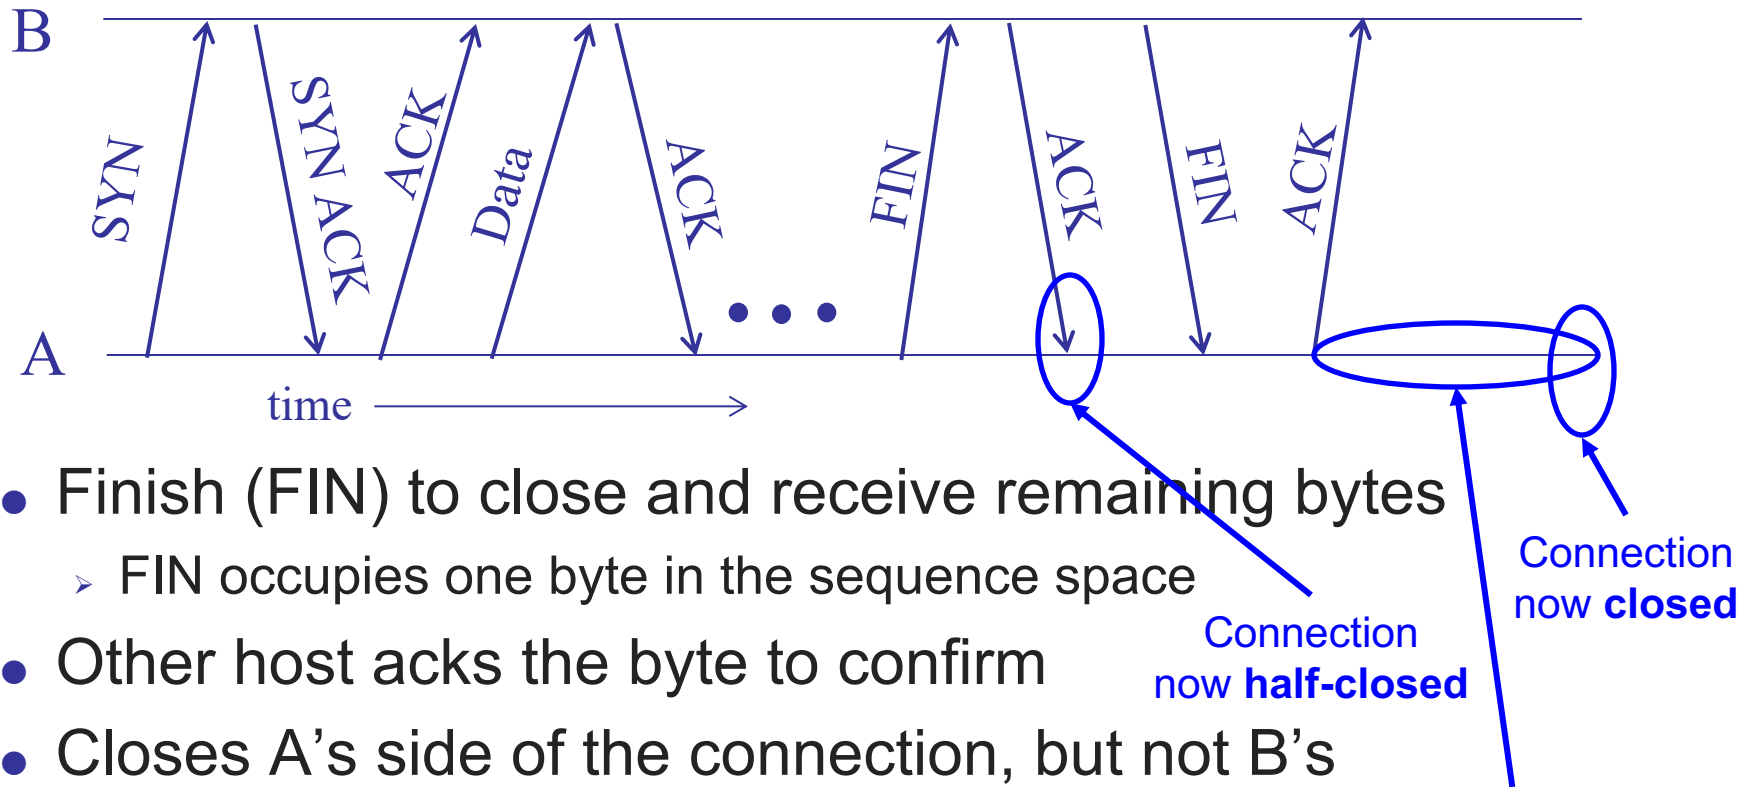
- Suppose the SYN packet gets lost
 - Packet dropped by the network or server is busy
- Eventually, no SYN-ACK arrives
 - Sender retransmits the SYN on timeout
- How should the TCP sender set the timer?
 - Sender has no idea how far away the receiver is
 - Hard to guess a reasonable length of time to wait
 - SHOULD (RFCs 1122 & 2988) use default of 3 seconds
 - » Some implementations instead use 6 seconds

SYN loss and web downloads

- User clicks on a hypertext link
 - Browser creates a socket and does a “connect”
 - The “connect” triggers the OS to transmit a SYN
- If the SYN is lost...
 - 3-6 seconds of delay: can be very long
 - User may become impatient and can retry
- User triggers an “abort” of the “connect”
 - Browser creates a new socket and another “connect”
 - Can be effective in some cases

TCP CONNECTION TEARDOWN

Normal termination, one side at a time



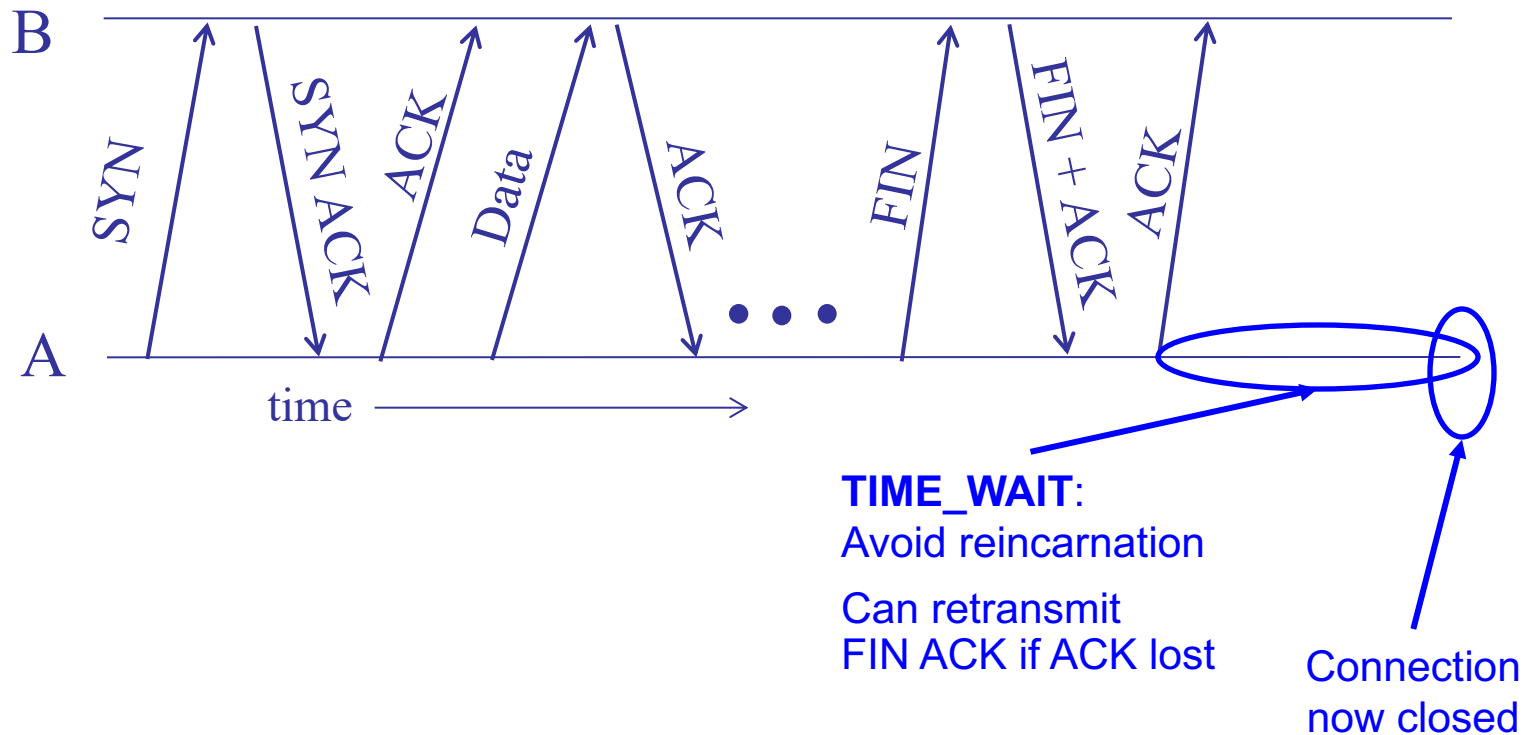
- Finish (FIN) to close and receive remaining bytes
 - FIN occupies one byte in the sequence space
- Other host acks the byte to confirm
- Closes A's side of the connection, but not B's
 - Until B likewise sends a FIN
 - Which A then acks

TIME_WAIT:

Avoid reincarnation


B will retransmit FIN if ACK is lost

Normal termination, both together



- Same as before, but B sets FIN with their ack of A's FIN

Abrupt termination



- A sends a RESET (RST) to B
 - E.g., because application process on A crashed
- That's it
 - B does not ack the RST
 - Thus, RST is not delivered reliably, and any data in flight is lost
 - But: if B sends anything more, will elicit another RST

Summary

- Reliability is not easy!
- Next
 - Flow control
 - LOTS of congestion control