

EECS 489

Computer Networks

Fall 2020

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Xin Jin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- Software-defined networking
- Programmable networks

The field of networking

- CS networking today is largely the study of the Internet
 - Perhaps the only history class many will take in CS

Building an artifact, not a discipline

- Other fields in “systems”: OS, DB, etc.
 - Teach basic principles
 - Are easily managed
 - Continue to evolve
- Networking:
 - Teach big bag of protocols
 - Notoriously difficult to manage
 - Evolves very slowly
- **Networks are much more primitive and less understood than other computer systems**

A tale of two planes

- **Data plane:** forwarding packets
 - Based on local forwarding state
- **Control plane:** computing that forwarding state
 - Involves coordination with rest of system

Original goals for the control plane

- **Basic connectivity:** route packets to destination
 - Local state computed by routing protocols
 - Globally distributed algorithms
- **Inter-domain policy:** find policy-compliant paths
 - Done by globally distributed BGP
- What other goals are there in running a network?

Extended roles of the control plane

- Performs various network management tasks
 - For example,
 - » Where to route?
 - » How much to route?
 - » At what rate to route?
 - » Should we route at all?
 - » ...

Bottom line

- Many different control plane mechanisms
- Each designed from scratch for their intended goal
- Encompassing a wide variety of implementations
 - Distributed, manual, centralized,...
- None of them particularly well designed
- Network control plane is a complicated mess!

“The Power of Abstraction”

- “Modularity based on abstraction is the way things get done”
 - Barbara Liskov
- Abstractions → Interfaces → Modularity

Analogy: Mainframe to PC evolution

Vertical integration, closed

- Specialized application
- Specialized operating system
- Specialized hardware

Open interfaces

- Arbitrary applications
- Commodity operating systems
- Microprocessor

We want the same for networking!

Many control plane mechanisms

- Variety of goals, no modularity
 - **Routing**: distributed routing algorithms
 - **Isolation**: ACLs, Firewalls,...
 - **Traffic engineering**: adjusting weights,...
- Control Plane: mechanism without abstraction
 - Too many mechanisms, not enough functionality

Task: Compute forwarding state

- Consistent with low-level hardware/software
 - Which might depend on vendor
- Based on entire network topology
 - Because many control decisions depend on topology
- For all routers/switches in network
 - Every router/switch needs forwarding state

Separate concerns with abstractions

- Be compatible with low-level hardware/software
 - Forwarding abstraction
- Make decisions based on entire network
 - Network state abstraction
- Compute configuration of each physical device
 - Specification abstraction

#1: Forwarding abstraction

- Express intent independent of implementation
 - Don't want to deal with proprietary HW and SW
- Design details concern exact nature of:
 - Header matching
 - Allowed actions

#2: Network state abstraction


- Abstraction: **global network view**
 - Annotated network graph provided through an API
- Creates a logically centralized view of the network (Network Operating System)
 - Runs on replicated servers in network (“controllers”)
- Information flows both ways
 - Information from routers/switches to form “view”
 - Configurations to routers/switches to control forwarding

#3: Specification abstraction

- Control mechanism expresses desired behavior
 - Whether it be isolation, access control, or QoS
- It should not be responsible for implementing that behavior on physical network infrastructure
 - Requires configuring the forwarding tables in each switch
- **Abstract view** of network
 - Models only enough detail to specify goals
 - Will depend on task semantics

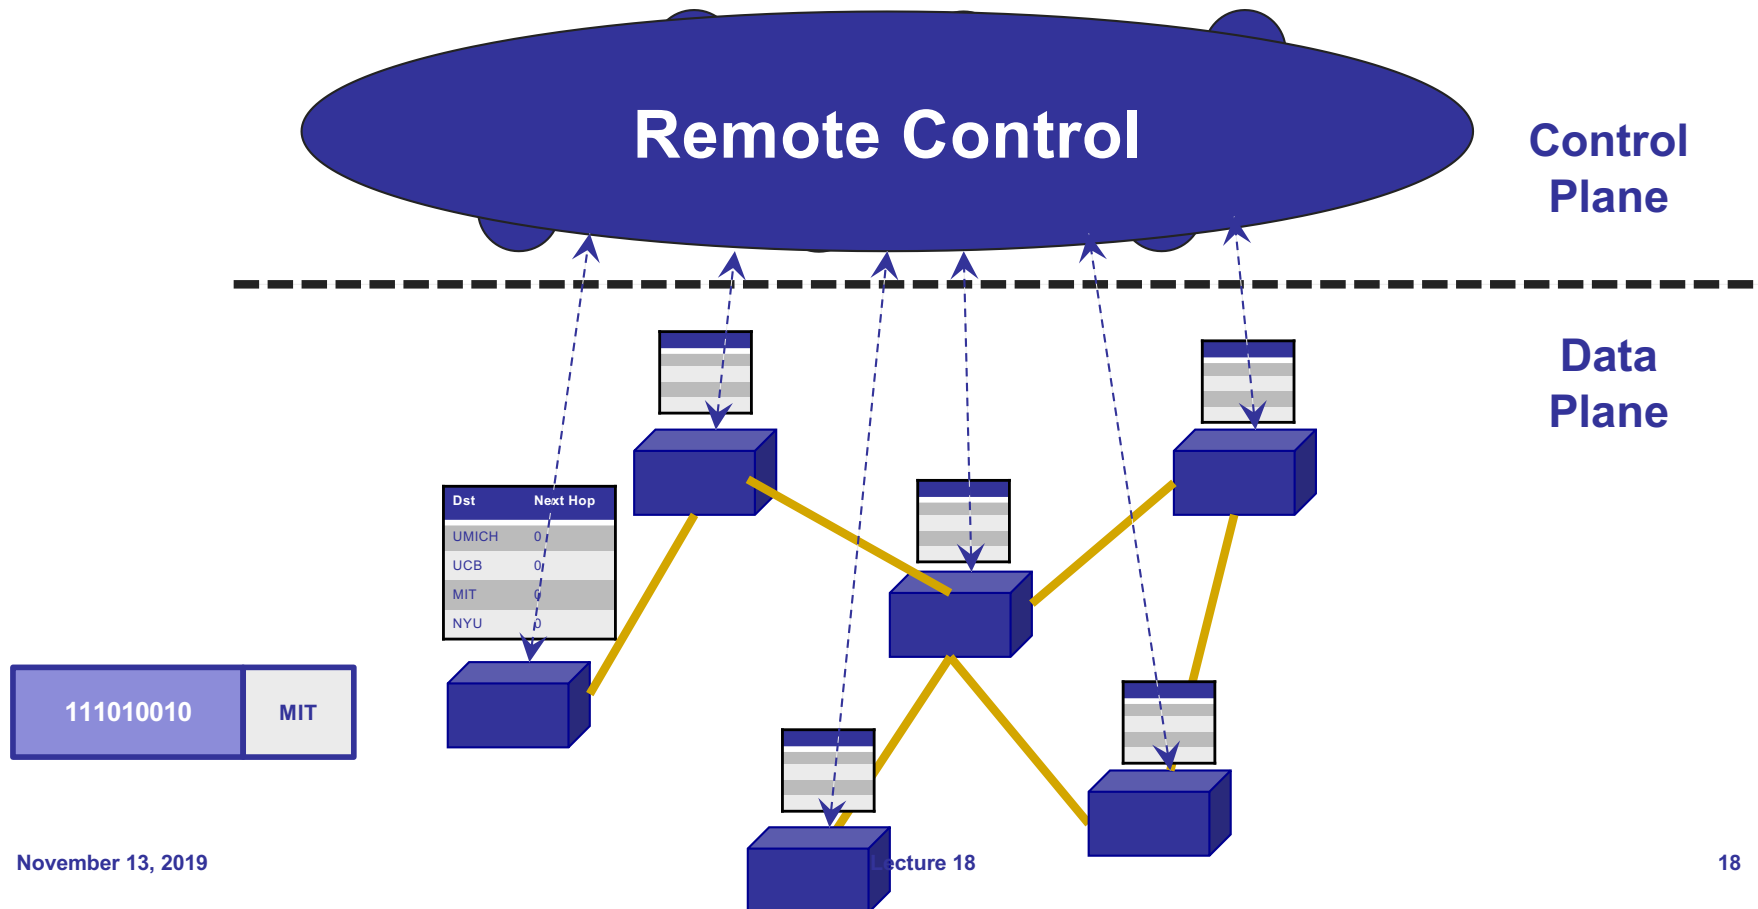
Traditional fully decentralized control plane

- Individual routing algorithm components in every router interact in the control plane



Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



Each goal is an app via specification abstraction

- What if an operator wants X?
- What if a customer wants to do weighted traffic splitting?
- ...
- There is an app for it!
 - Write your own routing protocol, load balancing algorithm, access control policies

Reason about each app via network state abstraction

- Now that the network is not distributed anymore and is a simple graph, we can **verify** whatever the correctness of whatever we specified

Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)
- Each router contains a **flow table**
- Each entry of the flow table defines a **match-action** rule
- Entries of the flow table is computed and distributed by the (logically) centralized controller

SDN: Many challenges remain

- Hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
 - Robustness to failures: leverage strong theory of reliable distributed system for control plane
 - Dependability, security: “baked in” from day one?
- Networks, protocols meeting mission-specific requirements
 - E.g., real-time, ultra-reliable, ultra-secure
- Internet-scaling

Some progress in the wide-area network (WAN)

- Google and Microsoft use SDN to manage traffic between datacenters
- One centralized controller to rule the entire world (well, their world)

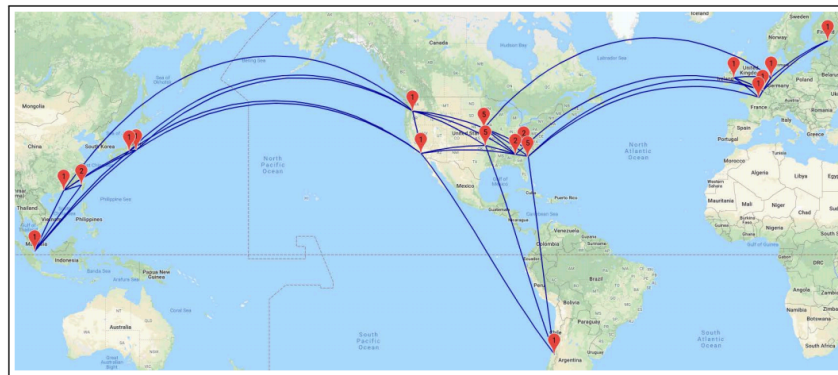


Figure 1: B4 global topology. Each marker indicates a site or multiple sites located in close geographical proximity. B4 consists of 33 sites as of January, 2018.

5-MINUTE BREAK!

Announcements

- Prof. Jennifer Rexford will be giving a distinguished lecture on Nov 13 2:45-3:45PM
 - Topic: **Networks Capable of Change**
 - <https://eecs.engin.umich.edu/event/networks-capable-of-change/>

A tale of two planes

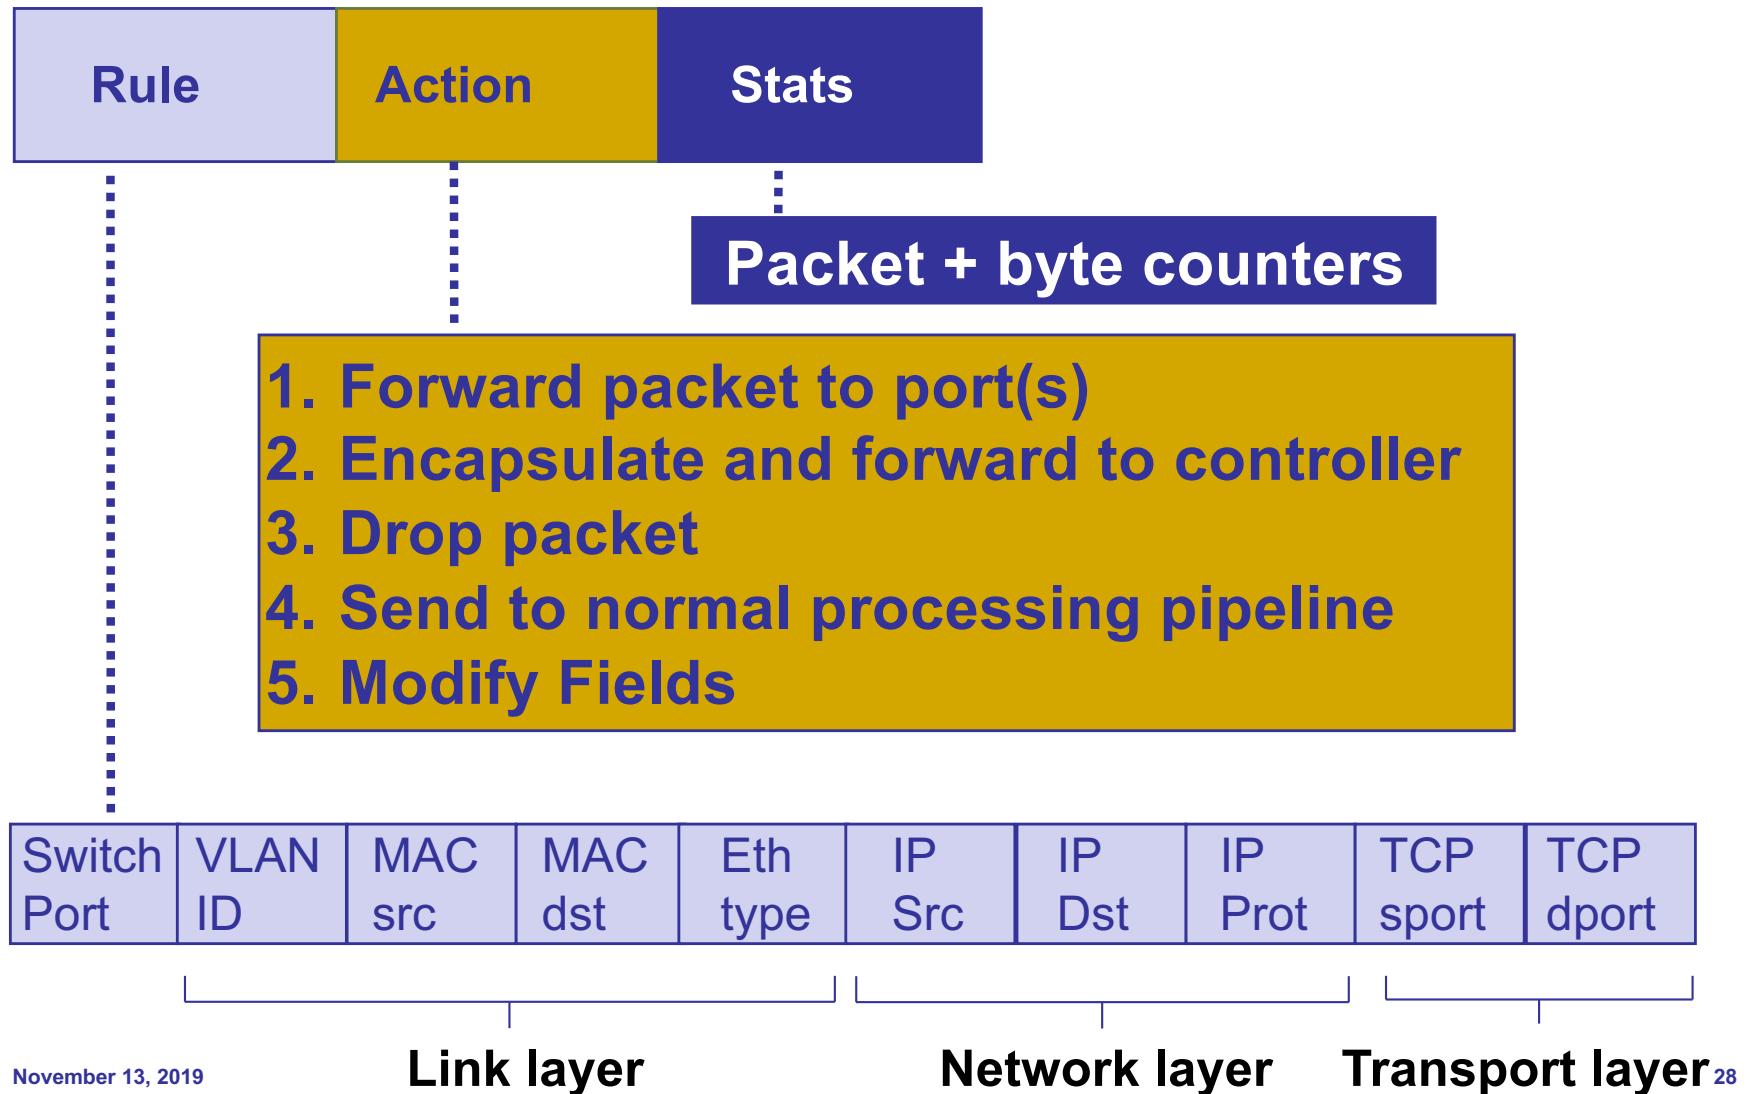
- **Data plane:** forwarding packets
 - Based on local forwarding state
- **Control plane:** computing that forwarding state
 - Involves coordination with rest of system

OpenFlow data plane abstraction

- Flow is defined by header fields
- Generalized forwarding: simple packet-handling rules
 - **Pattern**: match values in packet header fields
 - **Actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
 - **Priority**: disambiguate overlapping patterns
 - **Counters**: #bytes and #packets

1. src=1.2.*.*, dest=3.4.5.* → drop
2. src = *.*.*, dest=3.4.*.* → forward(2)
3. src=10.1.2.3, dest=*.*.* → send to controller

OpenFlow: Flow table entries




Forwarding abstraction

Match + Action: unifies different kinds of devices

- Router
 - Match: longest destination IP prefix
 - Action: forward out a link
- Switch
 - Match: destination MAC address
 - Action: forward or flood
- Firewall
 - Match: IP addresses and TCP/UDP port numbers
 - Action: permit or deny
- NAT
 - Match: IP address and port
 - Action: rewrite address and port

OpenFlow example

match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)




Example: datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2

match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)

OpenFlow protocol

OpenFlow Controller




- Operates between controller, switch
- TCP used to exchange messages
 - Optional encryption
- Three classes of OpenFlow messages:
 - Controller-to-switch
 - Asynchronous (switch to controller)
 - Symmetric (misc.)

Fixed-function data plane

- Traditional switches are fixed-function
 - They can do whatever they can do at birth, but they cannot change!
 - Bottom-up design
- Even OpenFlow was designed to be a fixed protocol
 - With a fixed table format
 - Capable of doing limited things

Takes forever to get a feature



Programmable data plane

- What if we could tell switches exactly what we want?
 - What table to keep?
 - What rules to use?
 - What data to keep track of?
 - ...

Top-down approach

- **Precisely specify** what you want to do and how you want a packet to be processed

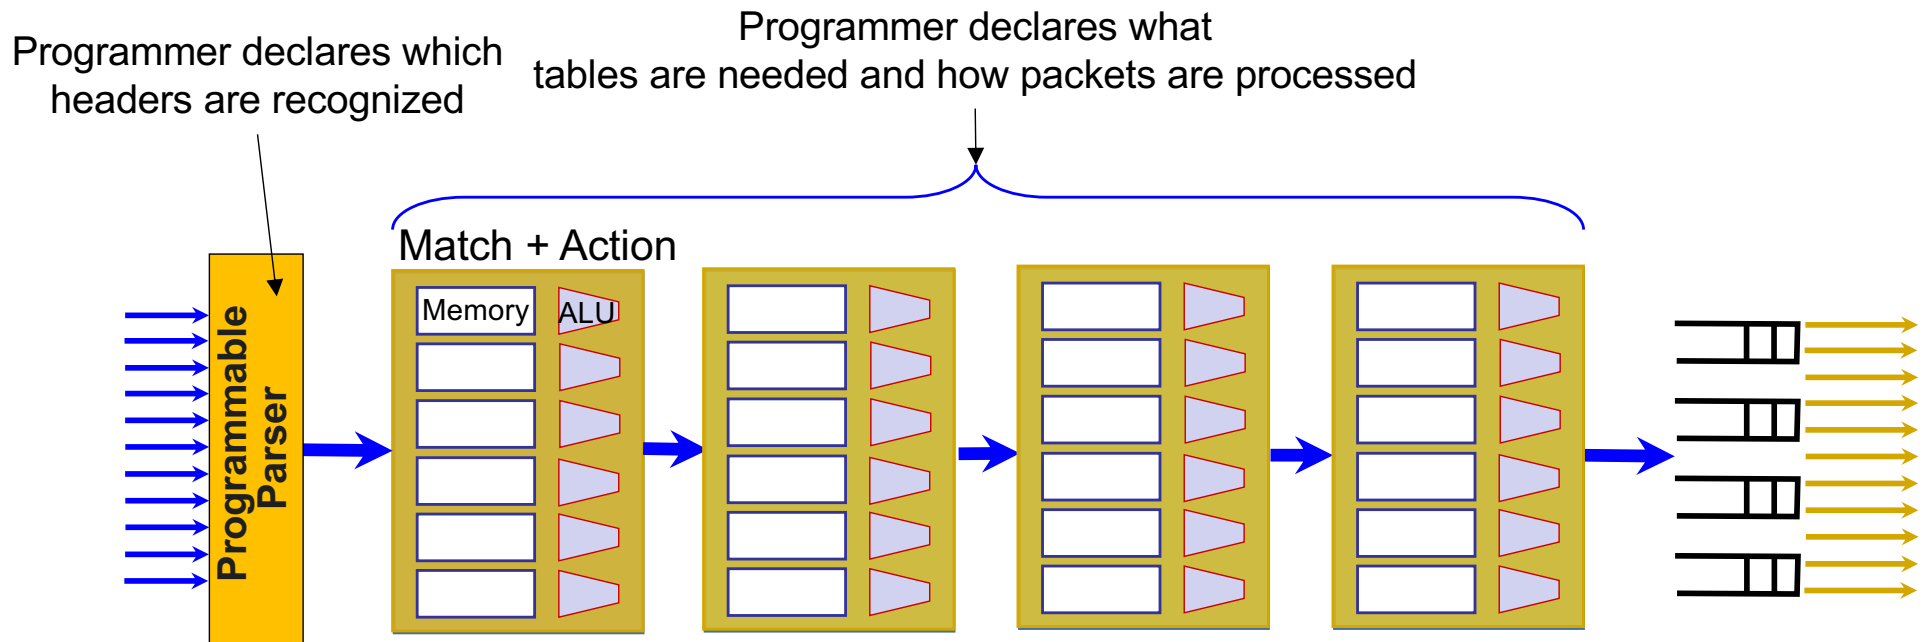
```
table int_table {  
    reads {  
        ip.protocol  
    }  
    actions {  
        export_queue_latency  
    }  
}
```

```
action export_queue_latency (sw_id) {  
    add_header(int_header);  
    modify_field(int_header.kind, TCP_OPTION_INT);  
    modify_field(int_header.len, TCP_OPTION_INT_LEN);  
    modify_field(int_header.sw_id, sw_id);  
    modify_field(int_header.q_latency,  
                intrinsic_metadata.deq_timedelta);  
    add_to_field(tcp.dataOffset, 2);  
    add_to_field(ipv4.totalLen, 8);  
    subtract_from_field(ingress_metadata.tcpLength, 12);  
}
```

What's left?

- Compile it down to be something runnable on a programmable switch
 - Similar to other high-level languages we use to run code on hardware like CPU, GPU, FPGA etc.
 - P4 for programmable switches
- But which switch?

PISA: Protocol Independent Switch Architecture



All stages are identical – makes PISA a good “compiler target”

How's programmability used today?

- Remove features to reduce complexity
- Add proprietary features
- Silicon independence or avoid vendor lock-in
- Telemetry and measurements

Example: In-band network telemetry (INT)

- “Which path did my packet take?”
- “Which rules did my packet follow?”
- “How long did it queue at each switch?”
- “Who did it share the queues with?”

Why now?

- One of the earlier incarnation of programmable networks was in mid 90s
 - Active networks
- What's changed after two+ decades?
 - **Hardware:** We can now make programmable switches as fast as fixed ones
 - **Software:** We have found a (so far) reasonable balance between programmability, performance, and security

Summary


- Abstractions beget modularity
 - Modularity is (almost always) good
- Programmability is powerful
 - Finding the right balance is hard
- Next lecture: Layer 2



Traffic engineering

- Want to avoid persistent overloads on links
- Choose routes to spread traffic load across links

Traffic engineering: Difficult




Q: What if network operator wants u-to-z traffic to flow along uvwz, x-to-z traffic to flow xwyz?

A: Need to define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!

Link weights are only control “knobs”


Traffic engineering: Difficult



Q: What if network operator wants to split u-to-z traffic along uvwz and uxyz (load balancing)?

A: Can't do it (or need a new routing algorithm)

Traffic engineering: Difficult



Q: What if w wants to route the two flows differently?

A: Can't do it (with LS or DV)

OpenFlow: Controller-to-switch messages

- Key controller-to-switch messages
 - **Features**: controller queries switch features, switch replies
 - **Configure**: controller queries/sets switch configuration parameters
 - **Modify-state**: add, delete, modify flow entries in the OpenFlow tables
 - **Packet-out**: controller can send this packet out of specific switch port

OpenFlow: Switch-to-controller messages

- Key switch-to-controller messages
 - **Packet-in**: transfer packet (and its control) to controller. See packet-out message from controller
 - **Flow-removed**: flow table entry deleted at switch
 - **Port status**: inform controller of a change on a port
- Network operators do not “program” switches by creating/sending OpenFlow messages directly.
 - Instead, they use higher-level abstraction at controller