

# **EECS 489**

# **Computer Networks**

**Fall 2020**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

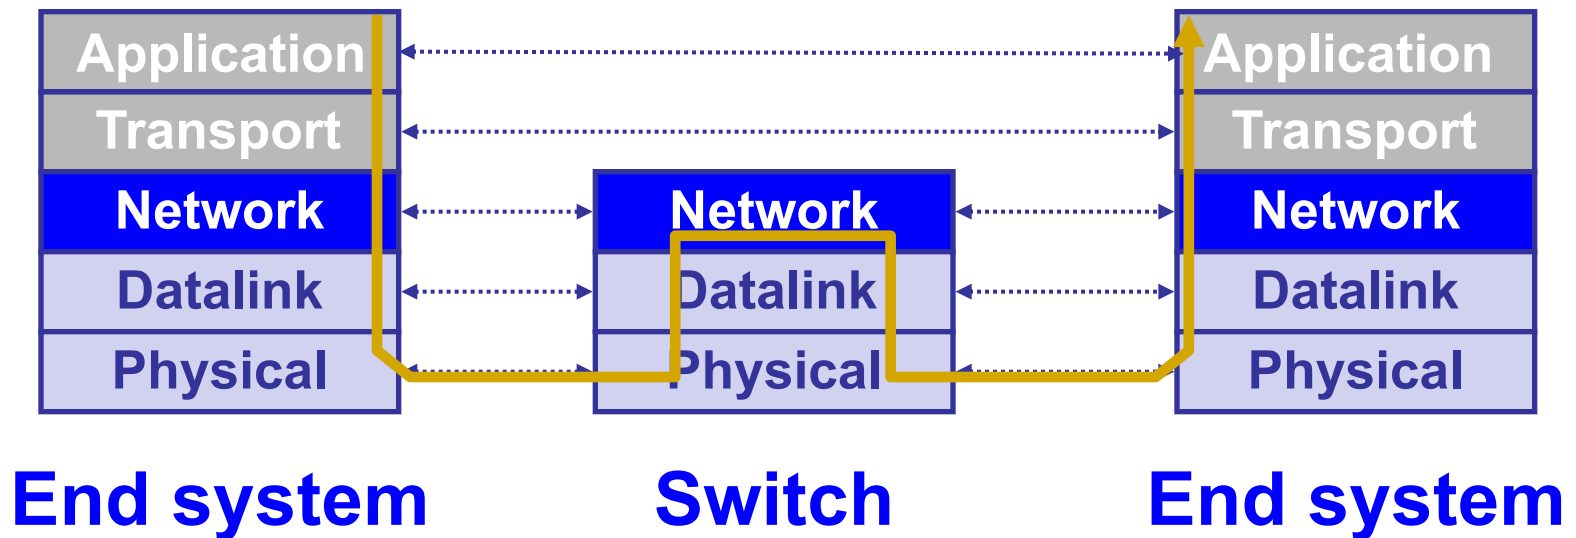
# Topics

---

- Network layer (lectures 12–16)
  - Intra-domain routing
  - Inter-domain routing
  - SDN

# Network layer

- Present everywhere
- Performs **addressing**, **forwarding**, and **routing**, among other tasks

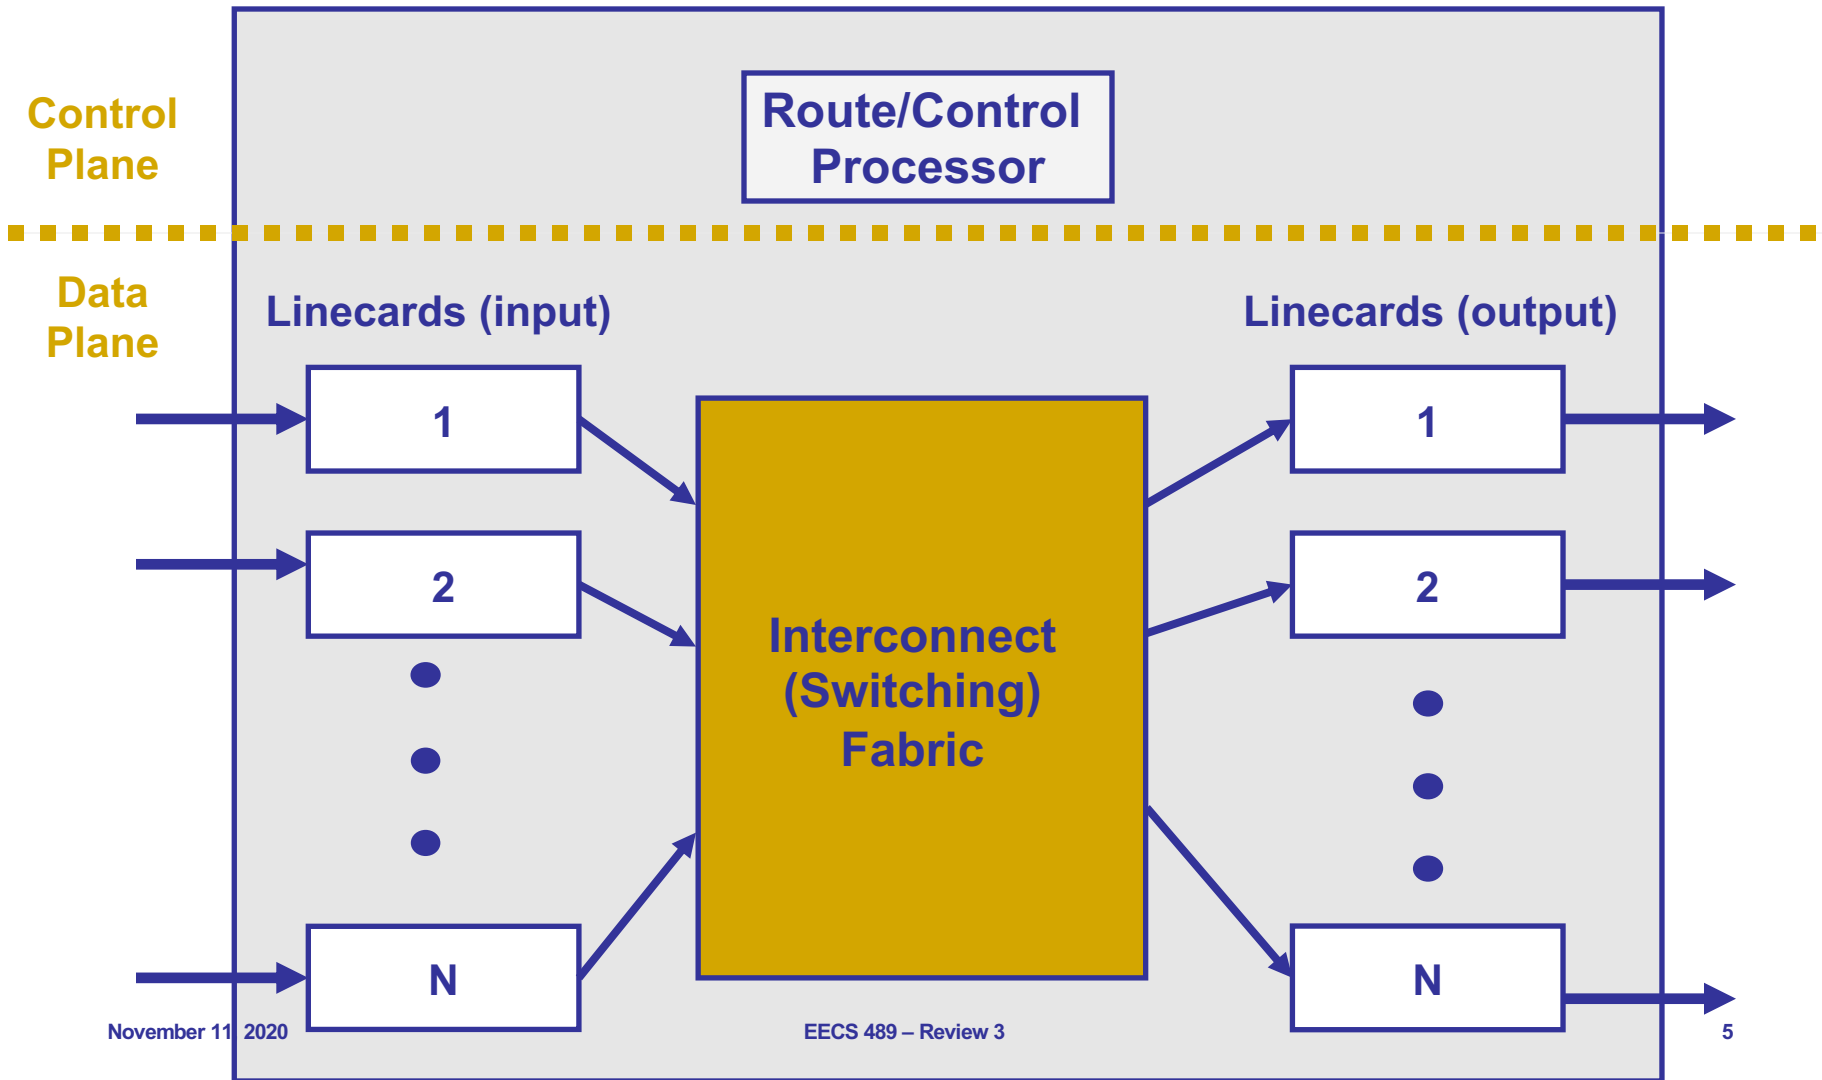


# Forwarding vs. routing

---

- Forwarding: “data plane”
  - Directing one data packet
  - Each router using local routing state
- Routing: “control plane”
  - Computing the forwarding tables that guide packets
  - Jointly computed by routers using a distributed algorithm
- Very different timescales!

# What's inside a router?



# Routing: Local vs. global view

---

- *Local* routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated
  - It must be evaluated in terms of the global context
- *Global* state refers to the collection of forwarding tables in each of the routers
  - Global state determines which paths packets take

# “Valid” routing state

---

- Global state is “valid” if it produces forwarding decisions that always deliver packets to their destinations
- Goal of routing protocols: compute valid state
  - How can we tell if routing state is valid?

# Necessary and sufficient condition

---

- Global routing state is valid *if and only if*:
  - There are no dead ends (other than destination)
  - There are no loops
- A **dead end** is when there is no outgoing link (next-hop)
  - A packet arrives, but the forwarding decision does not yield any outgoing link
- A **loop** is when a packet cycles around the same set of nodes forever



# Least-cost routes

---

- Least-cost routes provide an easy way to avoid loops
  - No reasonable cost metric is minimized by traversing a loop
- Least-cost paths form a spanning tree for each destination rooted at that destination

# Intra-domain routing

---

- Link-state (LS) routing protocol
  - Dijkstra's algorithm
  - Broadcast neighbors' info to everyone
- Distance vector (DV) routing protocol
  - Bellman-Ford algorithm
  - Gossip to neighbors about everyone

# Link-state routing

---

- Every router knows its local “link state”
  - Router u: “(u,v) with cost=2; (u,x) with cost=1”
- Each router floods its **local link state to all other routers** in the network
  - Does so periodically or when its link state changes
- Every router learns the entire network graph
  - Each runs Dijkstra’s Shortest-Path First (SPF) algorithm locally to compute forwarding table

# Distance-vector protocol

---

- Link-state routing protocol
  - Each node **broadcasts** its **local** information
- Distance-vector routing protocol
  - The opposite (sort of)
  - Each node **tells its neighbors** about its **global** view
- Use Bellman-Ford equation

# Similarities between LS and DV routing

---

- Both are shortest-path based routing
  - Minimizing cost metric (link weights) a common optimization goal
    - » Routers share a common view as to what makes a path “good” and how to measure the “goodness” of a path
- Due to shared goal, commonly used inside an organization
  - RIP and OSPF are mostly used for **intra**-domain routing

# Comparison of LS and DV routing

---

## Messaging complexity

- LS: with  $N$  nodes,  $E$  links,  $O(NE)$  messages sent
- DV: exchange between neighbors only

## Speed of convergence

- LS: relatively fast
- DV: convergence time varies
  - **Count-to-infinity** problem

## Robustness: what happens if router malfunctions?

- LS:
  - Node can advertise incorrect **link** cost
  - Each node computes its *own* table
- DV:
  - Node can advertise incorrect **path** cost
  - Each node's table used by others (**errors propagate**)

---

# **INTER-DOMAIN ROUTING**

# Autonomous systems (AS)

---

- An AS is a network under a single administrative control
  - Currently over 70,000 ASes
  - Updated daily at <http://www.cidr-report.org/as2.0/>
- ASes are sometimes called “domains”
- Each AS is assigned a unique identifier (ASN)
  - E.g., University of Michigan owns ASNs 177 to 180



# Addressing is key to scalable inter-domain routing

---

- Ability to aggregate addresses is crucial for
  - State: Small forwarding tables at routers
    - » Much less than the number of hosts
  - Churn: Limited rate of change in routing tables

# Classful addressing

---

- Three classes
  - 8-bit network prefix (Class A),
  - 16-bit network prefix (Class B), or
  - 24-bit network prefix (Class C)
- Example: an organization needs 500 addresses.
  - A single class C address is not enough ( $<500$  hosts)
  - Instead, a class B address is allocated ( $\sim 65K$  hosts)
    - » Huge waste!

# CIDR: Classless inter-domain routing

---

- Flexible division between network and host addresses
- Offers a better tradeoff between size of the routing table and efficient use of the IP address space


# Allocation done hierarchically

---

- Internet Corporation for Assigned Names and Numbers (ICANN) gives large blocks to...
- Regional Internet Registries, such as the American Registry for Internet Names (ARIN), which give blocks to...
- Large institutions (ISPs), which give addresses to...
- Individuals and smaller institutions
- FAKE Example:
  - ICANN → ARIN → AT&T → UMICH → EECS

# Hierarchy in IP addressing

- 32 bits are partitioned into a prefix and suffix components
- Prefix is the **network** component; suffix is the **host** component



- Inter-domain routing operates on network prefix

---

**5-MINUTE BREAK!**

# Announcements

---

- Prof. Jennifer Rexford will be giving a distinguished lecture on Nov 13 2:45-3:45PM
  - Topic: **Networks Capable of Change**
  - <https://eecs.engin.umich.edu/event/networks-capable-of-change/>


# Administrative structure shapes Inter-domain routing

---

- ASes want freedom to pick routes based on policy
- ASes want autonomy
- ASes want privacy



# Business relationships



## *Relations between ASes*


provider  $\longleftrightarrow$  customer

peer  $\bullet\text{---}\bullet$  peer

## *Business implications*

- Customers pay provider
- Peers don't pay each other

# Routing follows the money!



- ASes provide “transit” between their customers
- Peers do not provide transit between other peers


# BGP inspired by Distance-Vector with four differences

---


- Shortest-path routes may not be picked to enforce policy
- Path-Vector routing to avoid loops
- Selective route advertisement may affect reachability
- Routes may be aggregated for scalability

# BGP: Basic idea

---



# Policy dictates how routes are “selected” and “exported”



- **Selection:** Which path to use?
  - Controls whether/how traffic leaves the network
- **Export:** Which path to advertise?
  - Controls whether/how traffic enters the network

# Typical export policy

Destination prefix advertised by...	Export route to...
Customer	Everyone (providers, peers, other customers)
Peer	Customers
Provider	Customers


We'll refer to these as the “Gao-Rexford” rules  
(capture common – **but not required!** – practice)

# Selection using attributes

- Rules for route selection in priority order

Priority	Rule	Remarks
1	LOCAL PREF	Pick highest LOCAL PREF
2	ASPATH	Pick shortest ASPATH length
3	MED	Lowest MED preferred
4	eBGP > iBGP	Did AS learn route via eBGP (preferred) or iBGP?
5	iBGP path	Lowest IGP cost to next hop (egress router)
6	Router ID	Smallest next-hop router's IP address as tie-breaker

# Who speaks BGP?



Border routers in an Autonomous System



# eBGP, iBGP, and IGP

---

- **eBGP**: BGP sessions between border routers in different ASes
  - Learn routes to external destinations
- **iBGP**: BGP sessions between border routers and other routers within the same AS
  - Distribute externally learned routes internally
- **IGP**: “Interior Gateway Protocol” = Intra-domain routing protocol
  - Provide internal reachability via shortest path
  - E.g., OSPF, RIP

---

# **SOFTWARE-DEFINED AND PROGRAMMABLE NETWORKS**

# “The Power of Abstraction”

---

- “Modularity based on abstraction is the way things get done”
  - Barbara Liskov
- Abstractions → Interfaces → Modularity


# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Need an abstraction for general forwarding model
- Make decisions based on entire network
  - Need an abstraction for network state
- Compute configuration of each physical device
  - Need an abstraction that simplifies configuration

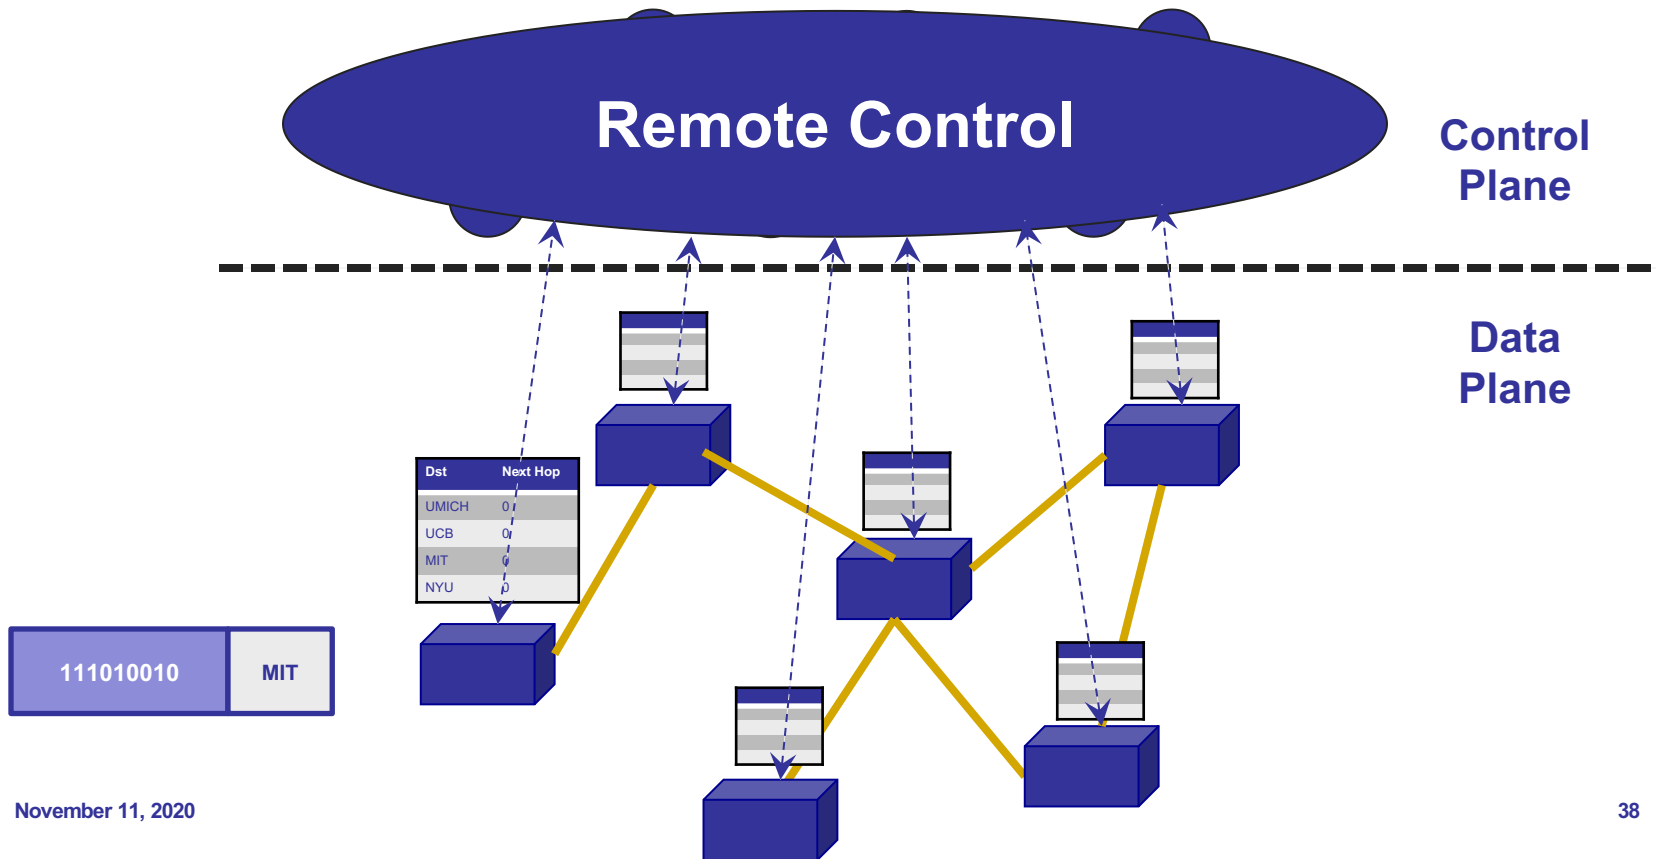
# Traditional fully decentralized control plane

- Individual routing algorithm components in every router interact in the control plane



# Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



# SDN: Many challenges remain

---

- Hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
  - **Robustness to failures**: leverage strong theory of reliable distributed system for control plane
  - **Security**: “baked in” from day one?
- Networks, protocols meeting mission-specific requirements
  - E.g., real-time, ultra-reliable, ultra-secure
- **Internet-scaling**

# OpenFlow data plane abstraction

---

- Flow is defined by header fields
- Generalized forwarding: simple packet-handling rules
  - **Pattern**: match values in packet header fields
  - **Actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority**: disambiguate overlapping patterns
  - **Counters**: #bytes and #packets

1. `src=1.2.*.*, dest=3.4.5.*` → drop
2. `src = *.*.*.*, dest=3.4.*.*` → forward(2)
3. `src=10.1.2.3, dest=*.*.*.*` → send to controller



# Fixed-function data plane

---

- Traditional switches are fixed-function
  - They can do whatever they can do at birth, but they cannot change!
  - Bottom-up design
- Even OpenFlow was designed to be a fixed protocol
  - With a fixed table format
  - Capable of doing limited things

# Programmable data plane

---

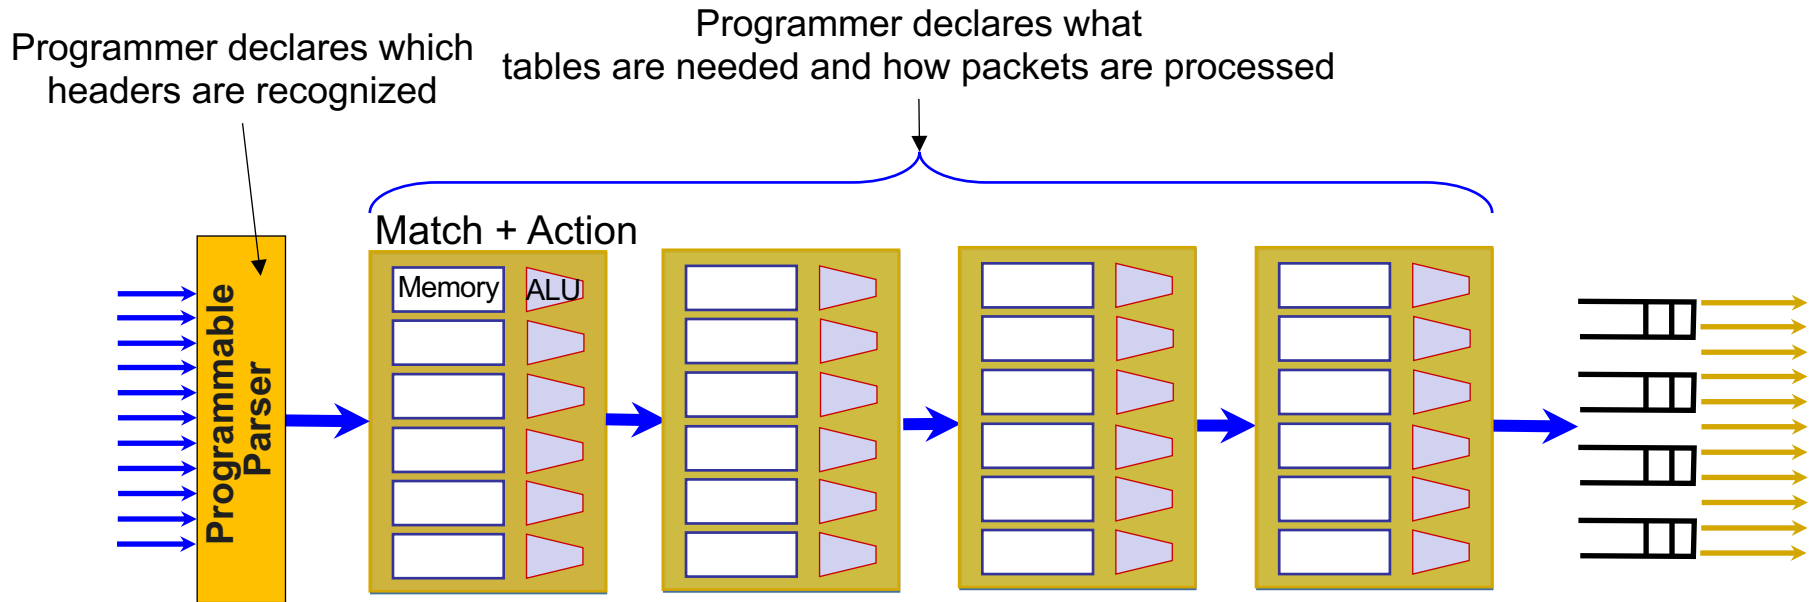
- What if we could tell switches exactly what we want?
  - What table to keep?
  - What rules to use?
  - What data to keep track of?
  - ...

# Top-down workflow

---

- Precisely specify using a well-defined language
- Compile it down to run on a standardized hardware (e.g., using P4)
- Run at line speed

# PISA: Protocol Independent Switch Architecture



**All stages are identical – makes PISA a good “compiler target”**

# Summary

---

- Intra-domain routing minimizes a cost metric
- Inter-domain routing is more complex due to policies
- Programmable networks are more flexible than fixed-function ones
- Next week: Layer 2
- Join us on Friday in welcoming Prof. Jen Rexford back to Michigan for a DLS talk