

EECS 489

Computer Networks

Fall 2020

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- Ethernet wrap-up
- Putting everything together

Recap: Switched Ethernet

- Ethernet started as a broadcast medium
 - Faced broadcast storm in larger setups due to flooding
- Constraints of switched Ethernet (for backward compatibility)
 - No changes to end-hosts
 - Maintain plug-n-play aspect

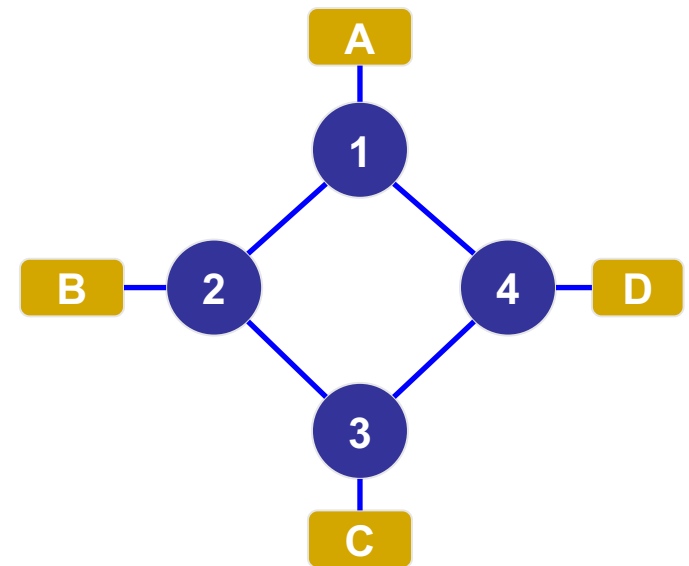
Ethernet topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery

Flooding (still) leads to loops

- Example: A wants to broadcast a message

- A sends packet to 1
- 1 Floods to 2 and 4
- 2 Floods to B and 3
- 4 Floods to D and 3
- 3 Floods packet from 2 to C and 4
- 3 Floods packet from 4 to C and 2
- 4 Floods packet from 3 to D and 1
- 2 Floods packet from 3 to B and 1
- 1 Floods packet from 2 to A and 4
- 1 Floods packet from 4 to B and 2
-



- Broadcast storm still happens in a switched network if it contains a cycle of switches

Spanning tree approach

- Take arbitrary topology
- Pick subset of links that form a spanning tree

Algorithm has two aspects

- Pick a root
 - Destination to which shortest paths go
 - Pick the one with the smallest identifier (MAC addr.)
- Compute shortest paths to the root
 - No shortest path can have a cycle
 - Only keep the links on shortest-paths
 - Break ties in some way (so we only keep one shortest path from each node)
- Ethernet's spanning tree construction does both with a single algorithm

Breaking ties





- When there are multiple shortest paths to the root, choose the path that uses the neighbor switch with the lower ID
 - One could use any tiebreaking system, but this is an easy one to remember and implement

Constructing a spanning tree

- Messages (Y, d, X)
 - From node X
 - Proposing Y as the root
 - Advertising a distance d to Y
- Switches elect the node with smallest identifier (MAC address) as root
- Each switch determines if a link is on its shortest path to the root; excludes it from the tree if not

Steps in the spanning tree algorithm

- Initially, each switch proposes itself as the root
 - Switch X announces $(X, 0, X)$ to its neighbors
- Switches update their view of the root
 - Upon receiving (Y, d, Z) from Z, check Y's id
 - If Y's id < current root: set root = Y
- Switches compute their distance from the root
 - Add 1 to the shortest distance received from a neighbor
- If root or shortest distance to it **changed**, send neighbors updated message $(Y, d+1, X)$



Robust spanning tree algorithm

- Algorithm must react to failures
 - Failure of the root node
 - Failure of other switches and links
- Root switch sends periodic root announcement messages
 - Other switches continue forwarding messages
- Detecting failures through **timeout**
 - If no word from root, time out and claim to be the root!

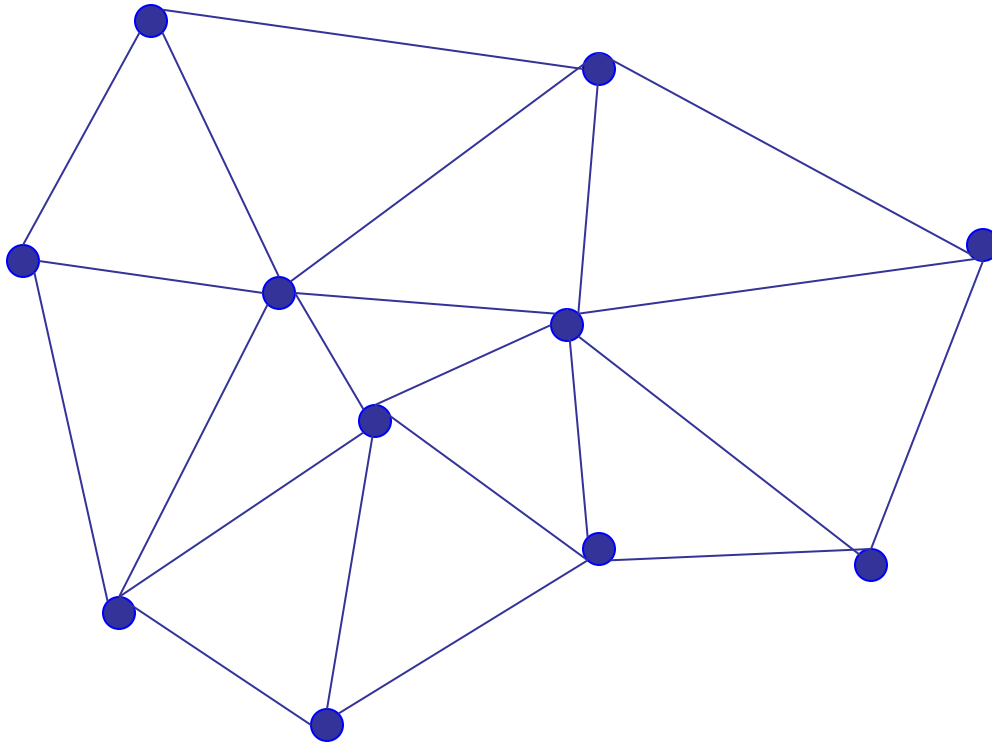
Ethernet topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery

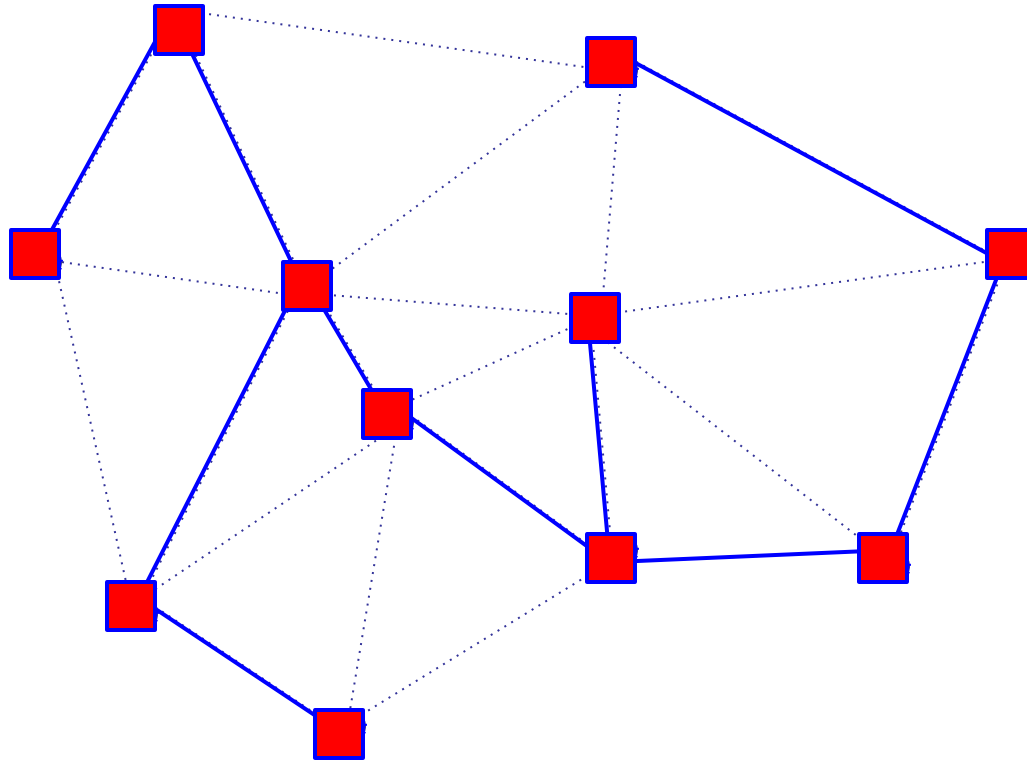
Flooding on a spanning tree

- Switches flood using the following rule:
 - (Ignore all ports not on spanning tree!)
 - Originating switch sends packet out all ports
 - When a packet arrives on one incoming port, send it out all ports other than the incoming port

Flooding on spanning tree



Flooding on spanning tree



But isn't flooding wasteful?

- Yes, but we can use it to bootstrap more efficient forwarding
- **Idea**: watch the packets going by, and learn from them
 - If node A sees a packet from node B come in on a particular port, it knows what port to use to reach B!
 - **Works because there is only one path to B**

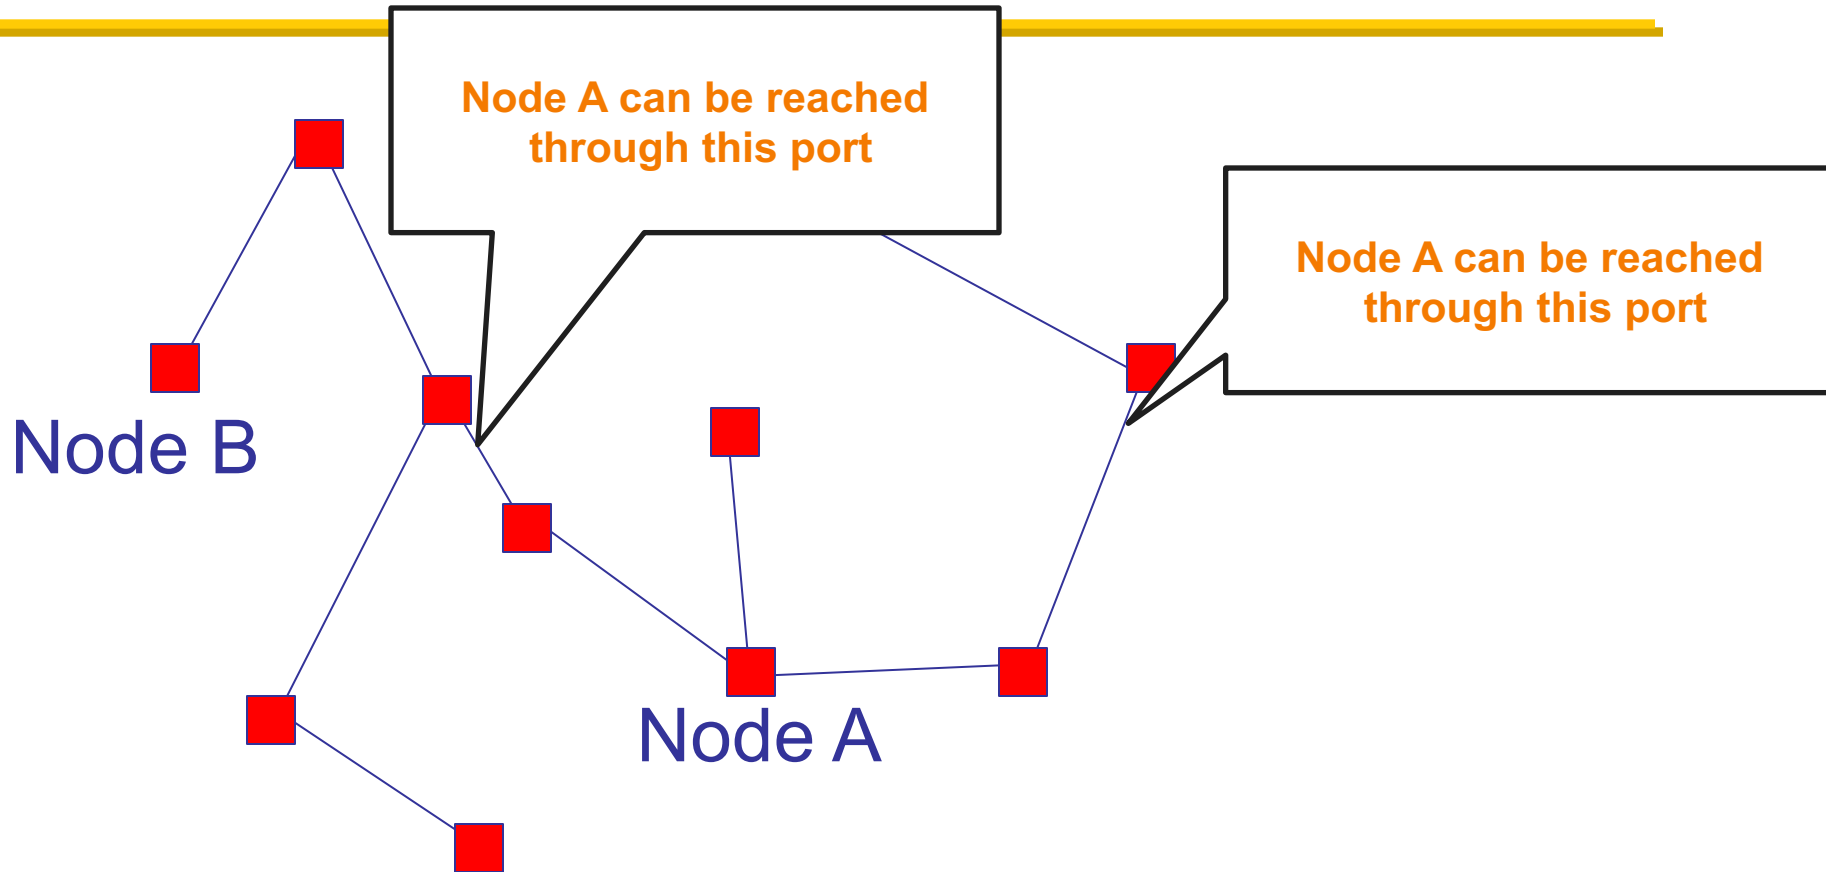
Nodes can “learn” routes

- Switch learns how to reach nodes by remembering where flooding packets came from
 - If flood packet from Node A entered switch on port 4, then switch uses port 4 to send to Node A

General approach

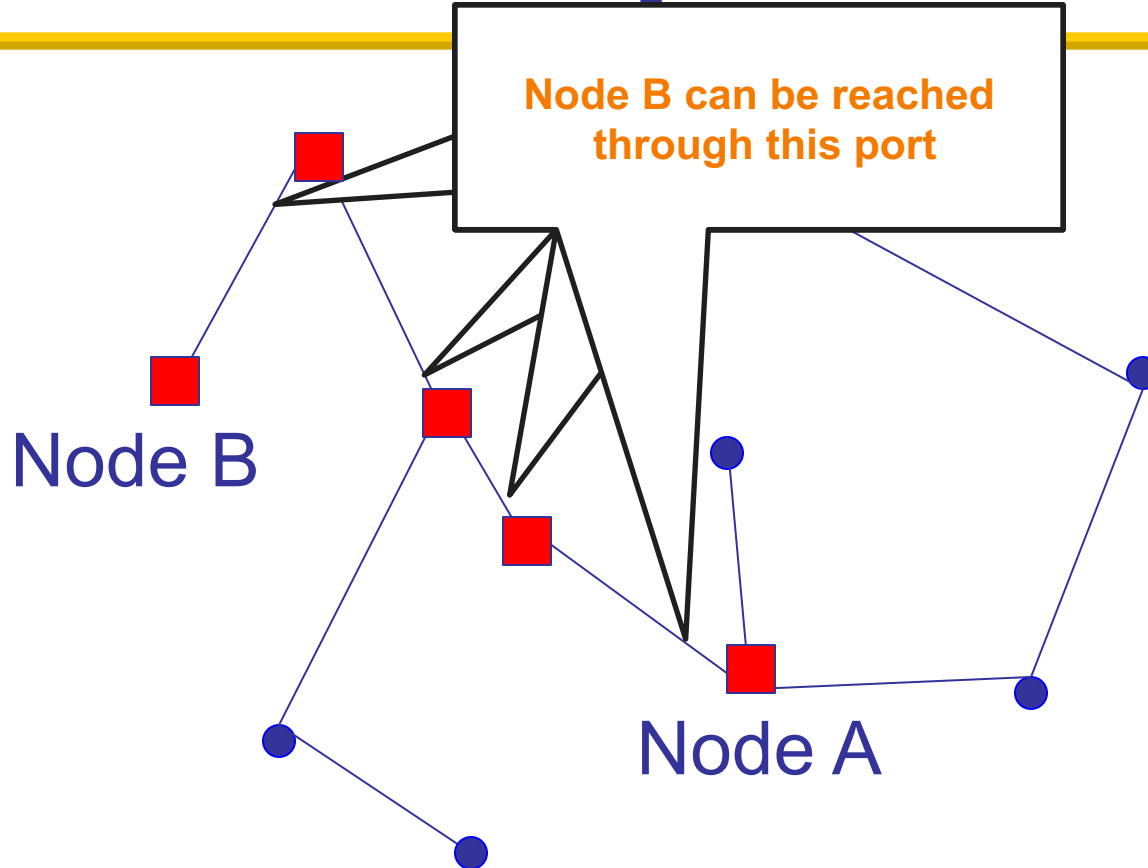
- Flood first packet to node you are trying to reach
- All switches learn where you are
- When destination responds, some switches learn where it is...
 - Only some switches, because packet to you follows direct path, and is not flooded

Learning from flood packets



Once a node has sent a flood message, all other switches know how to reach it....

Node B responds



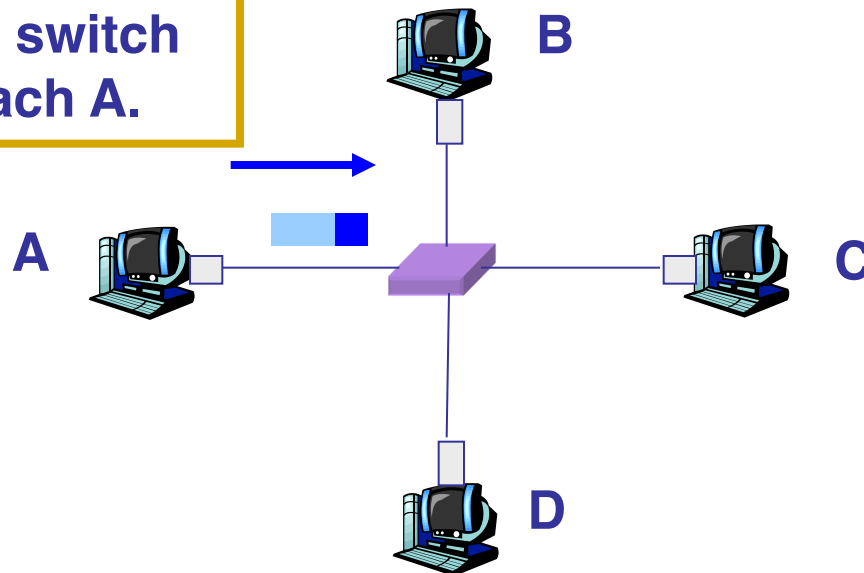
When a node responds, some of the switches learn where it is

November 18, 2020 21

Ethernet switches are “self learning”

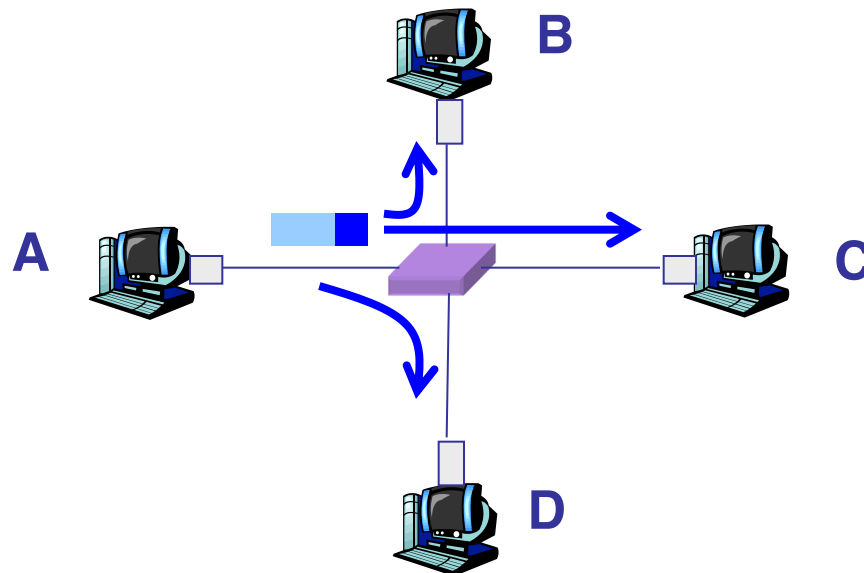
- When a packet arrives:
 - Inspect source MAC address, associate with incoming port
 - Store mapping in the switch table
 - Use **time-to-live** field to eventually forget mapping

Packet tells switch
how to reach A.



Self learning: Handling misses

- When packet arrives with unfamiliar destination
- Forward packet out all other ports
- Response may teach switch about that destination



Summary of learning approach

- Avoids loop by restricting to spanning tree
 - This makes flooding possible
- Flooding allows packet to reach destination
- And in the process switches learn how to reach source of flood
- No route “computation”
- Forwarding entries a consequence of traffic pattern

Contrast

IP

- Packets forwarded on all available links
- Addresses can be aggregated
- Routing protocol computes loop-free paths
- Forwarding table computed by routing protocol

Ethernet

- Packets forwarded on subset of links (spanning tree)
- Flat addresses
- “Routing” protocol computes loop-free topology
- Forwarding table derived from data packets(+ spanning tree for floods)

Strengths of Ethernet's approach

- Plug-n-Play: zero-configuration / self-*
- Simple
- Cheap

Weaknesses of Ethernet's approach

- Much of the network bandwidth goes unused
 - Forwarding is only over the spanning tree
- Delay in reestablishing spanning tree
 - Network is “down” until spanning tree rebuilt
 - Rebuilt spanning tree may be quite different
- Slow to react to host movement
 - Entries must time out
- Poor predictability
 - Location of root and traffic pattern determines forwarding efficiency

5-MINUTE BREAK!

Link layer topics


- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery and bootstrapping

Discovery

- A host is “born” knowing only its MAC address
- Must **discover** lots of information before it can communicate with a remote host B
 - What is my IP address?
 - What is B’s IP address? (remote)
 - What is B’s MAC address? (if B is local)
 - What is my first-hop router’s address? (if B is not local)
 - ...

ARP and DHCP

- Link layer discovery protocols
 - ARP → Address Resolution Protocol
 - DHCP → Dynamic Host Configuration Protocol
 - Confined to a single local-area network (LAN)
 - Rely on broadcast capability



ARP and DHCP

- Link layer discovery protocols
- Serve two functions
 - Discovery of local end-hosts
 - » For communication between hosts on the same LAN
 - Bootstrap communication with remote hosts
 - » What's my IP address?
 - » Who/where is my local DNS server?
 - » Who/where is my first hop router?

DHCP

- Dynamic Host Configuration Protocol
 - Defined in RFC 2131
- A host uses DHCP to discover
 - Its own IP address
 - Its netmask
 - IP address(es) for its local DNS name server(s)
 - IP address(es) for its first-hop “default” router(s)

DHCP: Operation

- One or more local DHCP servers maintain required information
 - IP address pool, netmask, DNS servers, etc.
 - Application that listens on UDP port 67

DHCP: Operation

- One or more local DHCP servers maintain required information
- Client broadcasts a DHCP discovery message
 - L2 broadcast, to MAC address FF:FF:FF:FF:FF:FF

DHCP: Operation

- One or more local DHCP servers maintain required information
- Client broadcasts a DHCP discovery message
- One or more DHCP servers responds with a DHCP “offer” message
 - Proposed IP address for client, lease time
 - Other parameters

DHCP: Operation

- One or more local DHCP servers maintain required information
- Client broadcasts a DHCP **discovery** message
- One or more DHCP servers responds with a DHCP “offer” message
- Client broadcasts a DHCP **request** message
 - Specifies which offer it wants
 - Echoes accepted parameters
 - Other DHCP servers learn they were not chosen

DHCP: Operation

- One or more local DHCP servers maintain required information
- Client broadcasts a DHCP discovery message
- One or more DHCP servers responds with a DHCP “offer” message
- Client broadcasts a DHCP request message
- Selected DHCP server responds with an ACK

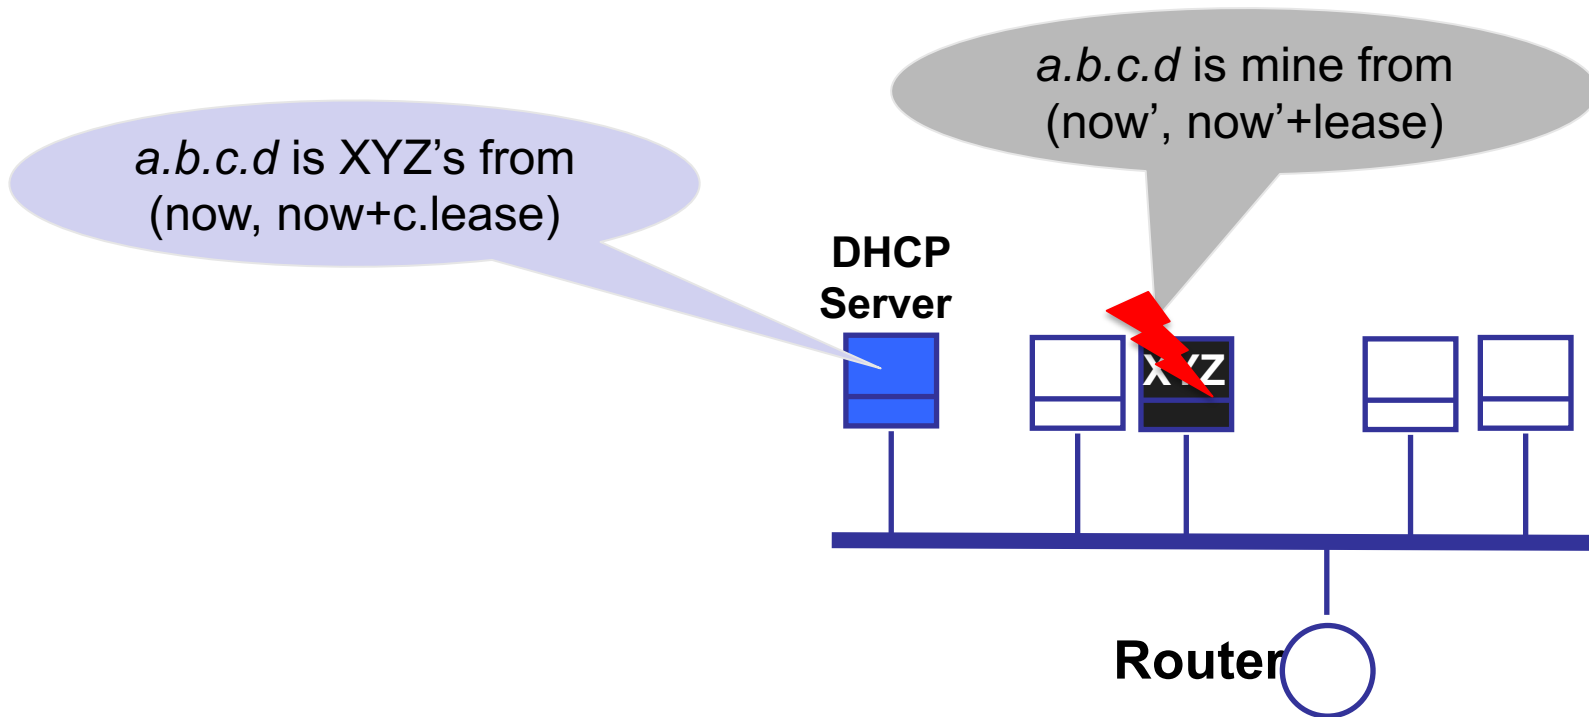
DHCP: Operation

- One or more local DHCP servers maintain required information
- Client broadcasts a DHCP discovery message
- One or more DHCP servers responds with a DHCP “offer” message
- Client broadcasts a DHCP request message
- Selected DHCP server responds with an ACK
- DHCP “relay agents” used when the DHCP server is not on the same broadcast domain

DHCP uses “soft state”

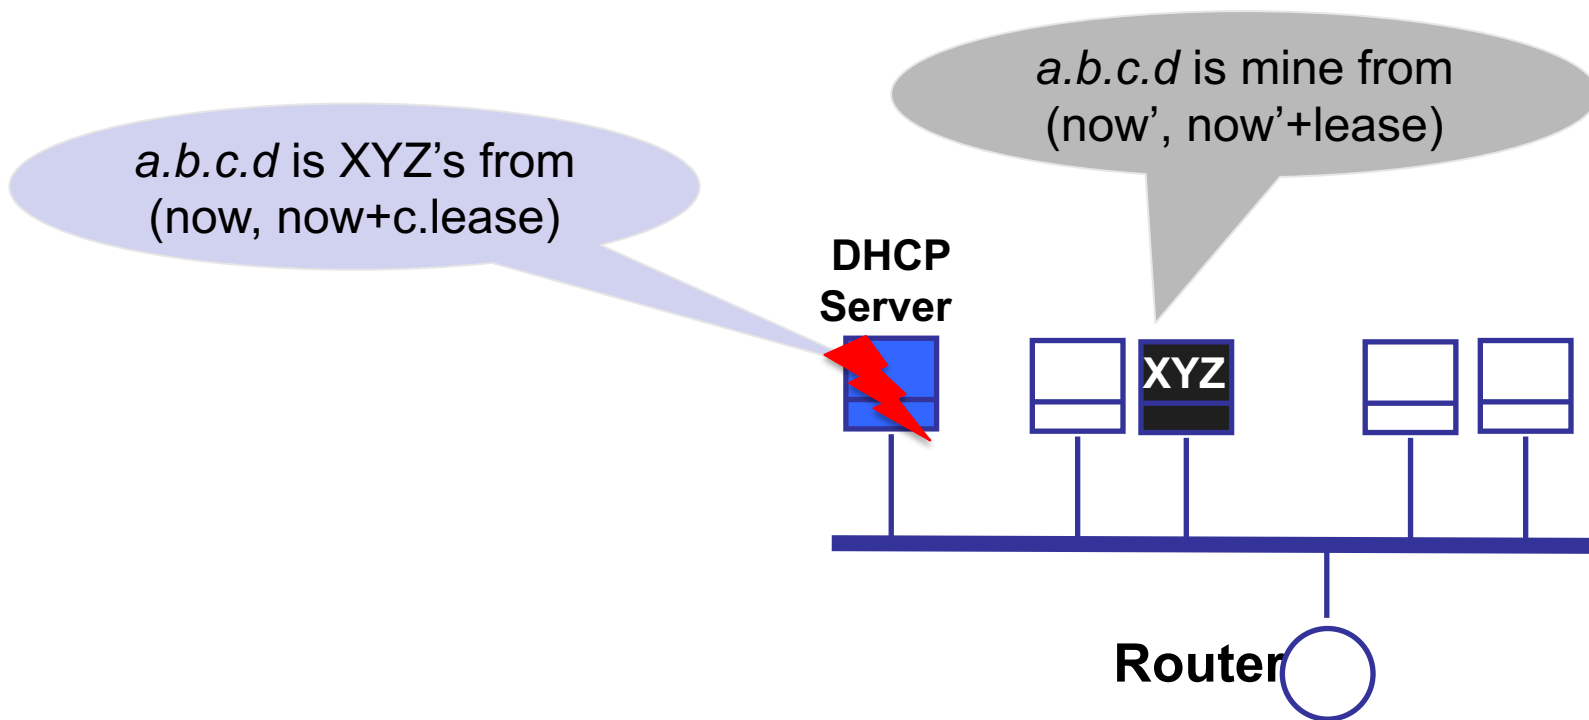
- Soft state: if not refreshed, state is forgotten
 - Hard state: allocation/revocation is deliberate
- Implementation:
 - Address allocations have a **lease** period
 - Server sets a timer for each allocation
 - Client must request a refresh before lease expires
 - Server resets timer when a refresh arrives and ACKs
 - »OR reclaims allocated address when timer expires
- Simple, yet robust under failure

Soft state under failure




- What happens when host XYZ fails?
 - Refreshes from XYZ stop
 - Server reclaims *a.b.c.d* after $O(\text{lease period})$

Soft state under failure




- What happens when server fails?
 - ACKs from server stop
 - XYZ releases address after $O(\text{lease period})$; send new request
 - A new DHCP server can come up from a 'cold start' and we are back on track in $\sim \text{lease time}$

Soft state under failure



- What happens if the network fails?
 - Refreshes and ACKs don't get through
 - XYZ release address; DHCP server reclaims it

Are we there yet?



What I learnt from DHCP


my IP: 1.2.3.48

netmask: 1.2.3.0/24
(255.255.255.0)

Local DNS: 1.2.3.156

router: 1.2.3.19

Sending packets over link Layer




- Link layer only understands MAC addresses
 - Translate the destination IP address to MAC address
 - Encapsulate the IP packet in a link-level (Ethernet) frame

ARP: Address Resolution Protocol

- Every host maintains an ARP table
 - List of (IP address → MAC address) pairs
- Consult the table when sending a packet
 - Map dest. IP address to dest. MAC address
 - Encapsulate (IP) data packet with MAC header; xmit
- What if IP address not in the table?
 - Sender broadcasts: Who has IP address 1.2.3.156?
 - Receiver replies: MAC address 58-23-D7-FA-20-B0
 - Sender caches result in its ARP table

What if the destination is remote?

- Look up the MAC address of the first hop router
 - 1.2.3.48 uses ARP to find MAC address for first-hop router 1.2.3.19 rather than ultimate destination IP address
- How does the red host know the destination is not local?
 - Uses netmask (discovered via DHCP)
- How does the red host know about 1.2.3.19?
 - Also DHCP



Key ideas in both ARP and DHCP

- **Broadcasting**: Can use broadcast to make contact
 - Scalable because of limited size
- **Caching**: remember the past for a while
 - Store the information you learn to reduce overhead
- **Soft state**: eventually forget the past
 - Associate a time-to-live field with the information
 - ... and either refresh or discard the information
 - Key for robustness in the face of unpredictable change

ID resolution in the networking stack

| Layer | Examples | Structure | Configuration | Resolution Service |
|---------------|-------------------|--------------------------|---------------|--------------------|
| App. Layer | cse.umich.edu | Organizational hierarchy | ~ manual | ↕ DNS |
| Network Layer | 123.45.6.78 | topological hierarchy | DHCP | |
| Link layer | 45-CC-4E-12-F0-97 | vendor (flat) | hard-coded | ↕ ARP |

Discovery mechanisms

- We have seen two approaches
 - Broadcast (ARP, DHCP)
 - » Flooding does not scale
 - » No centralized point of failure
 - » Zero configuration
 - Directory service (DNS)
 - » No flooding = scalable
 - » Root of the directory is vulnerable (caching is key)
 - » Needs configuration to bootstrap (local, root servers, etc.)

Summary

- Spanning tree enables Ethernet to efficiently flood a network to learn routes while forwarding packets
- DHCP and ARP form the discovery backplane of networking and make everything work together
- **After the break:** the end-to-end picture and specialized networks (e.g., datacenter, mobile)